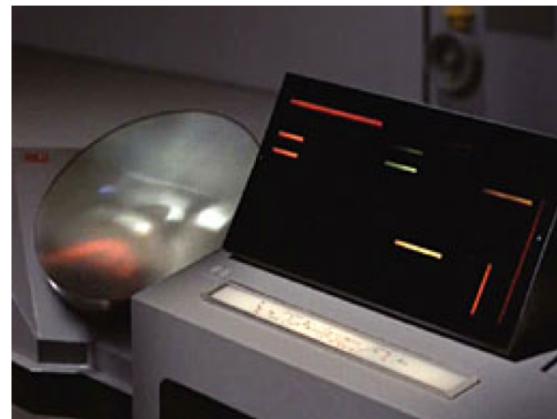
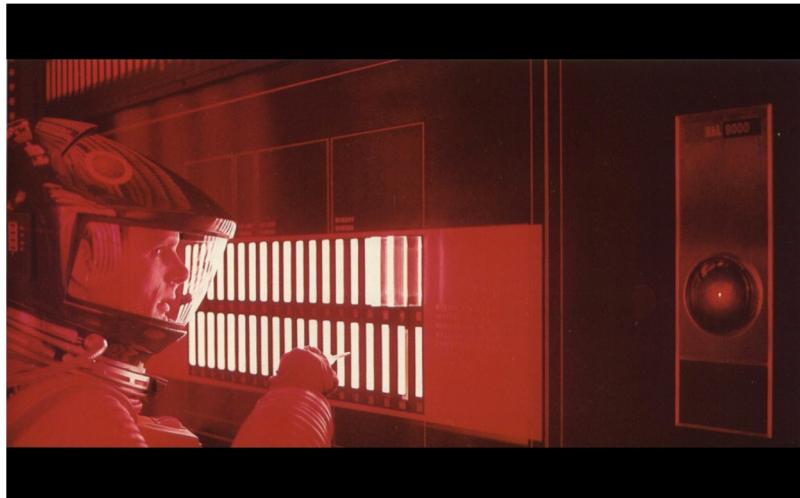
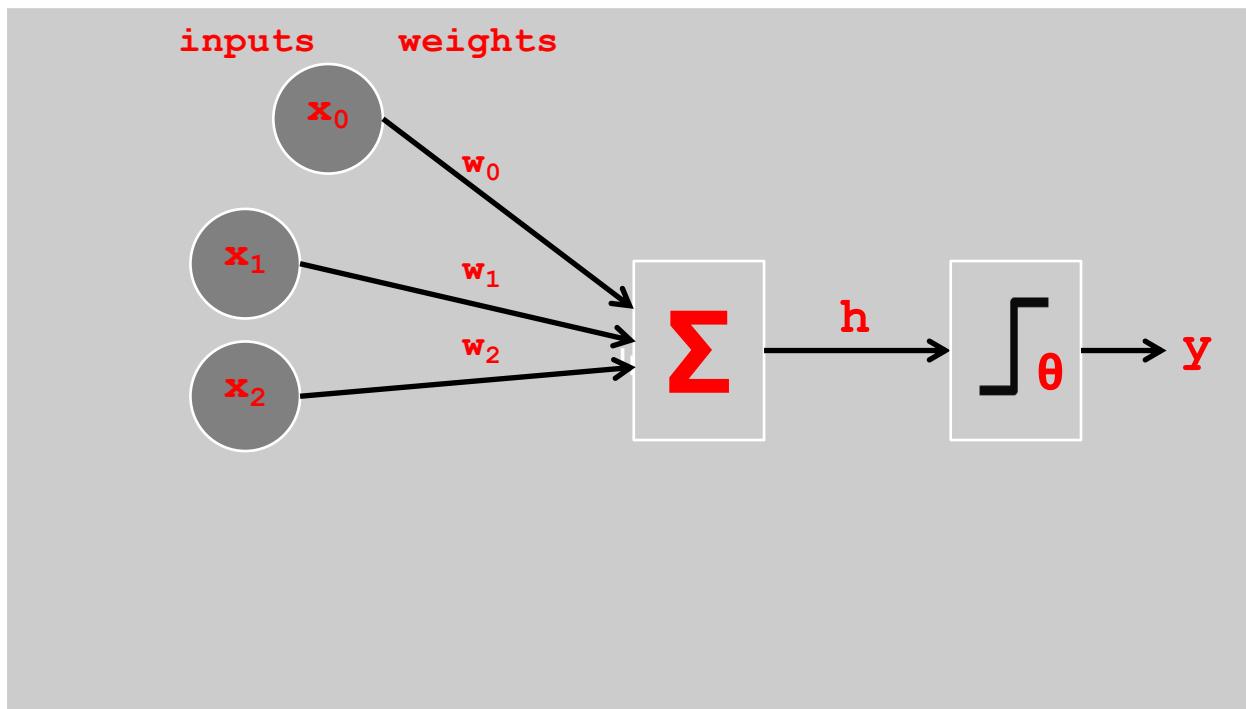


# Perceptrons are The Future!

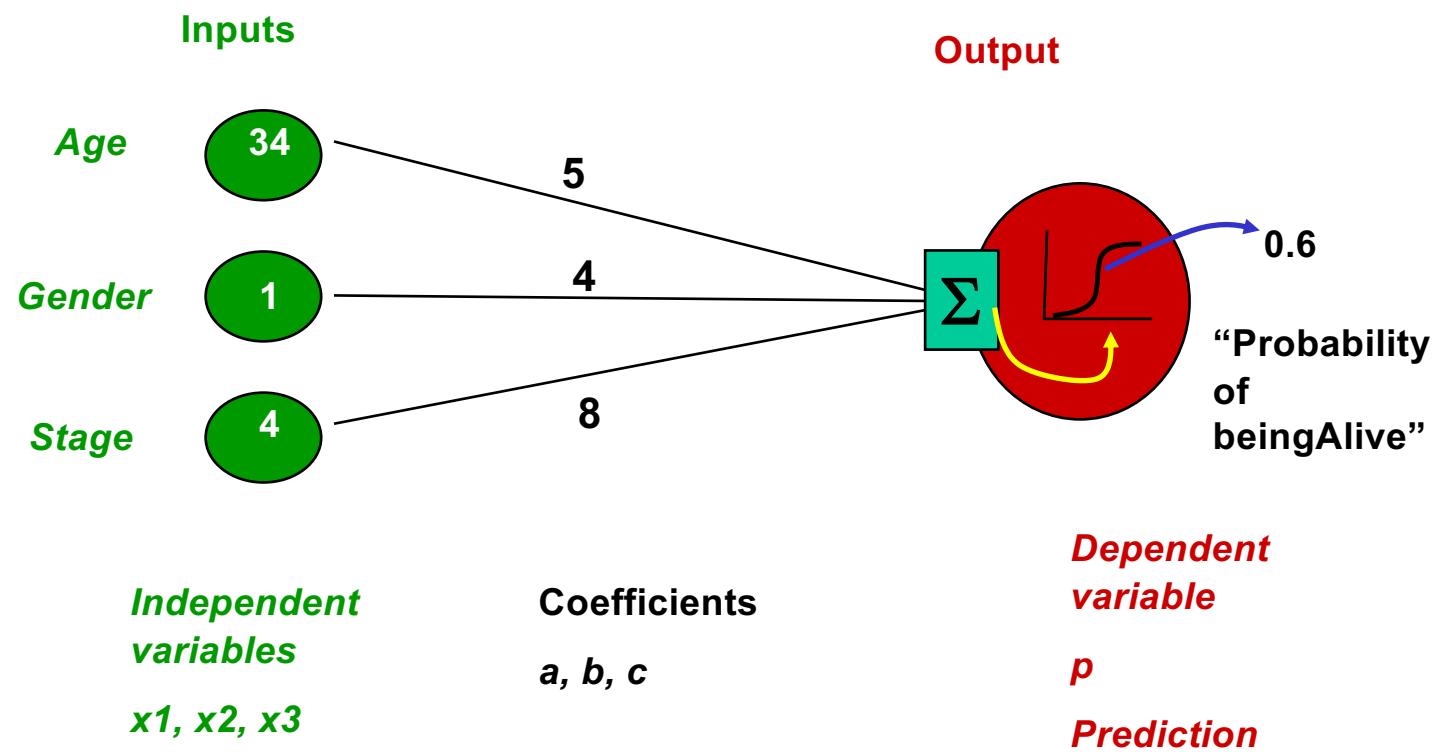
---



# Perceptron Networks

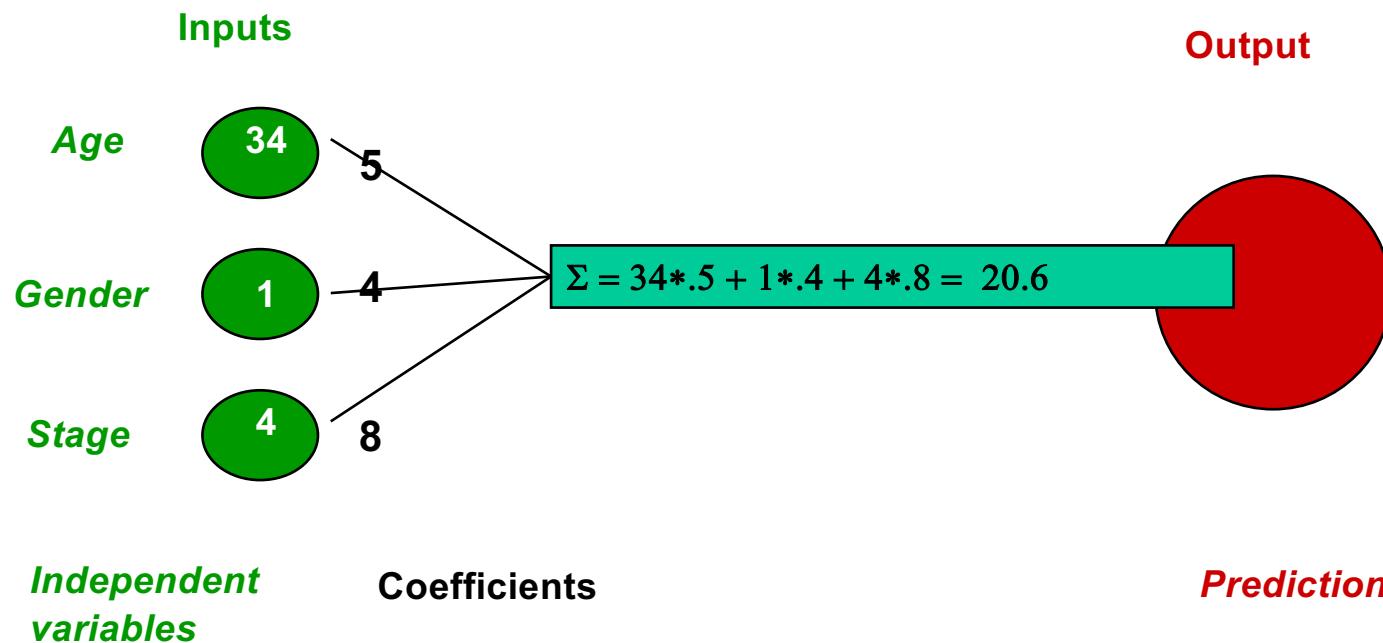


# Logistic Regression Model

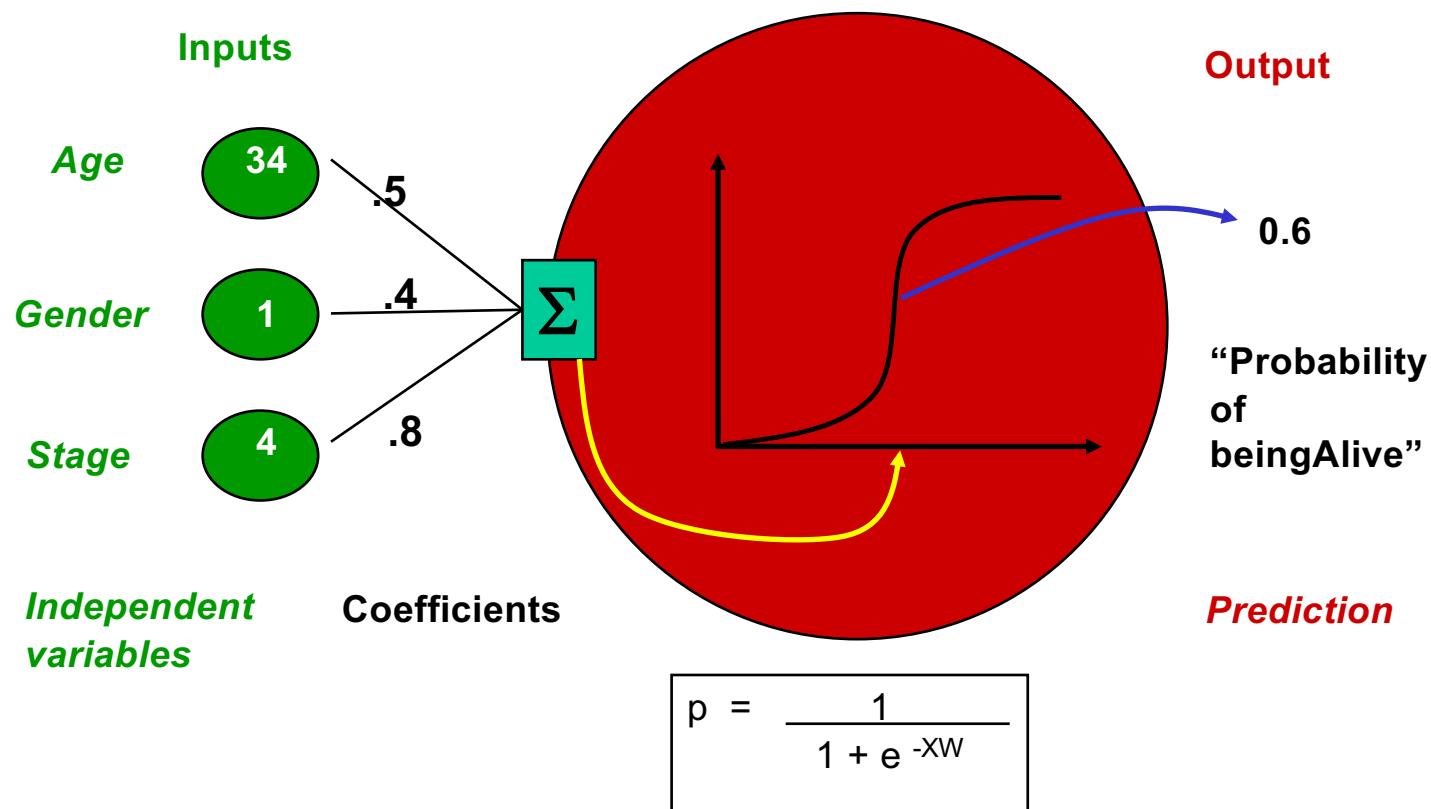


---

$\Sigma$  is the sum of inputs \* weights



# Logistic function

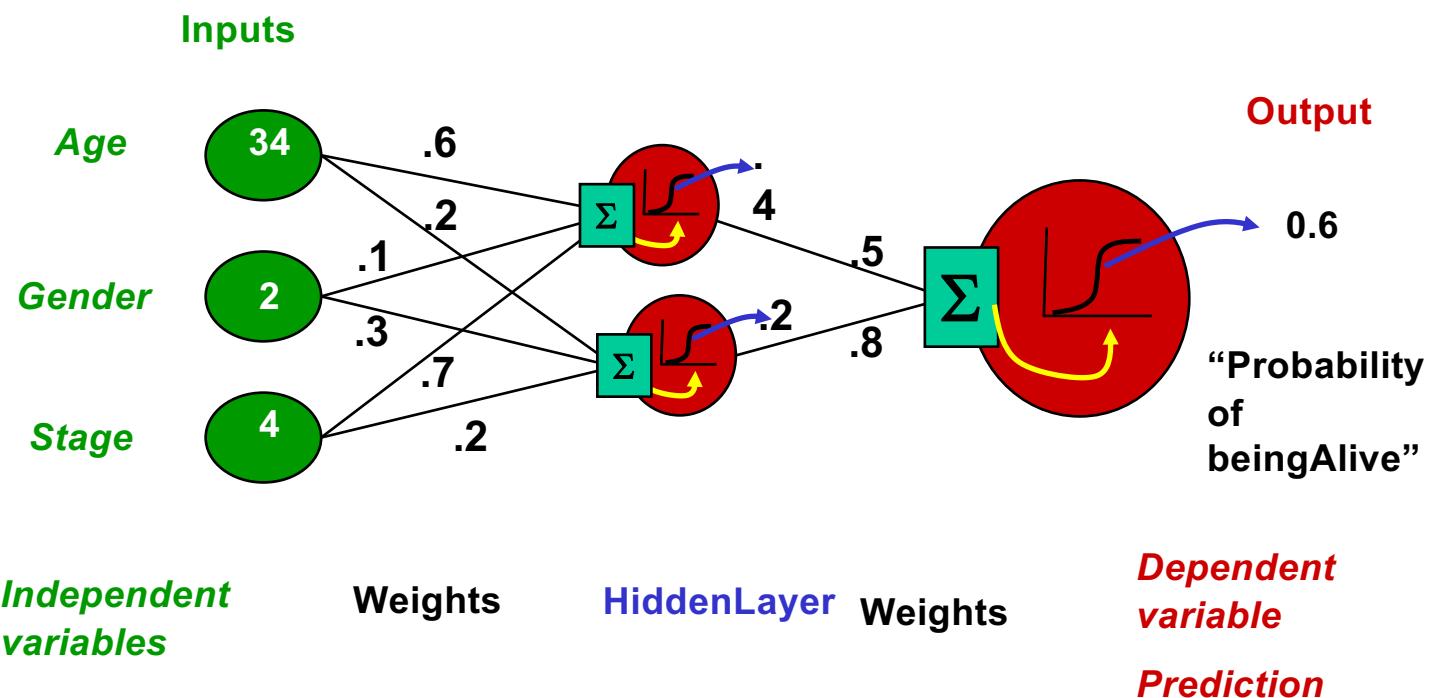


# Activation Functions...

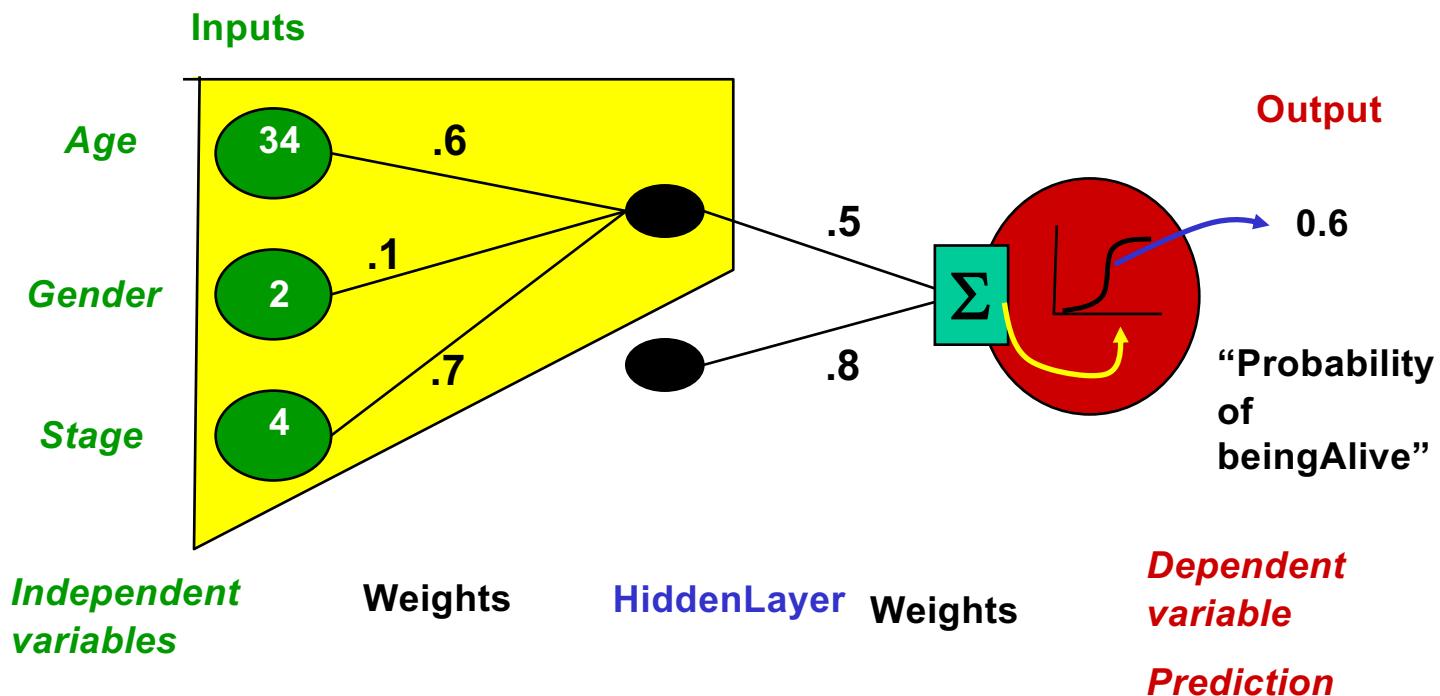
---

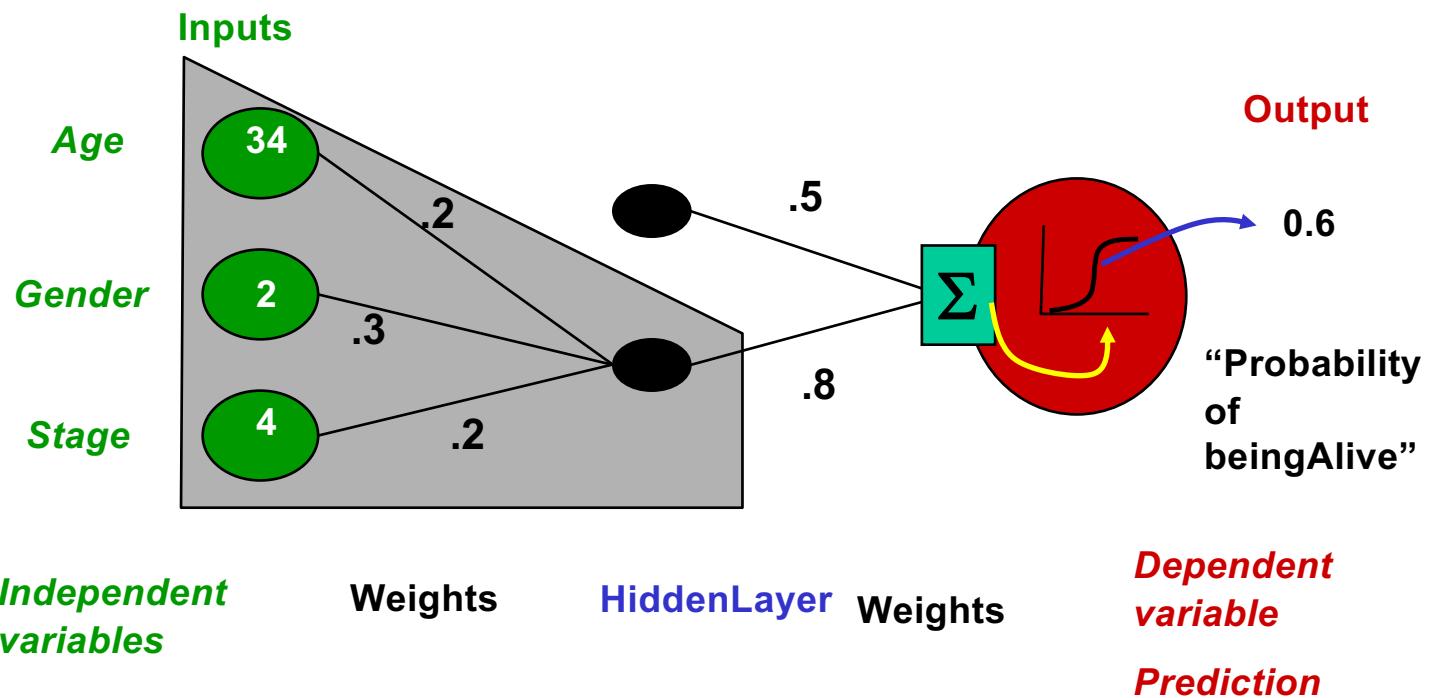
- **Linear**
- **Threshold or step function**
- **Logistic, sigmoid, “squash”**
- **Hyperbolic tangent**

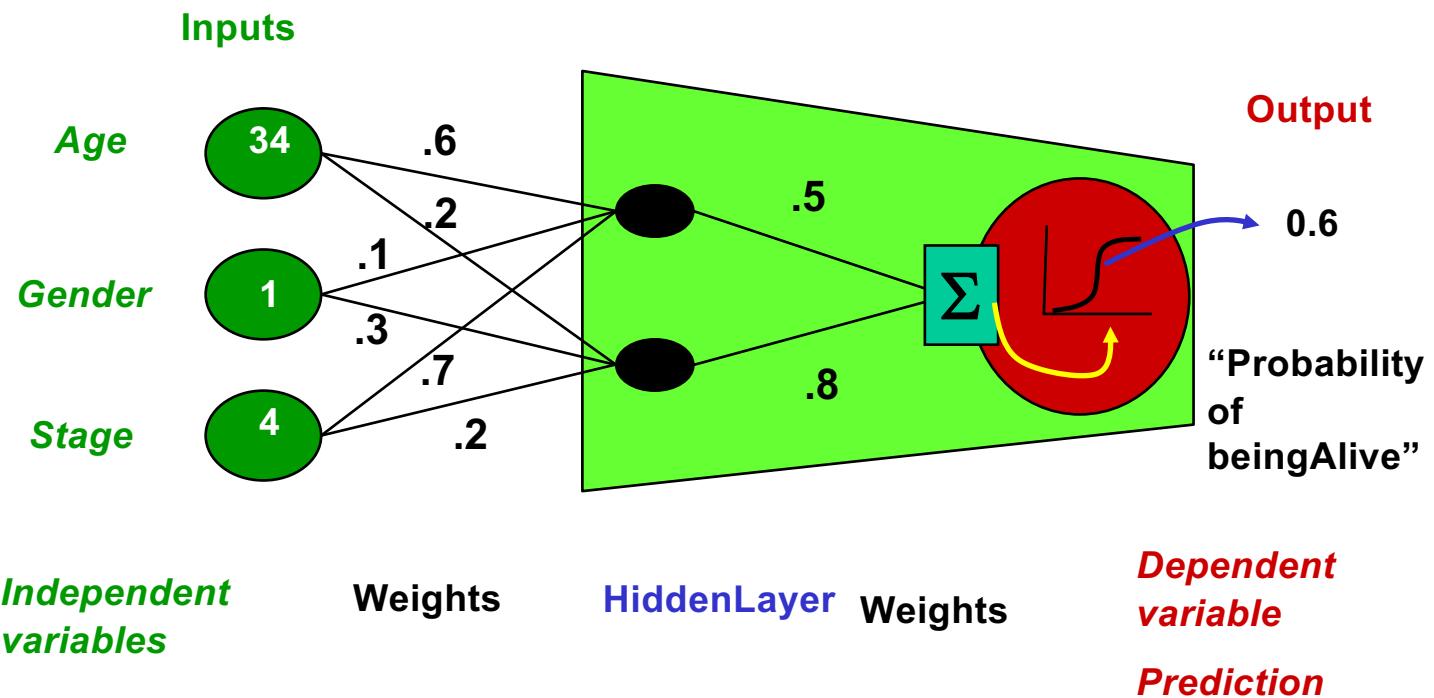
# Neural Network Model



# “Combined logistic models”

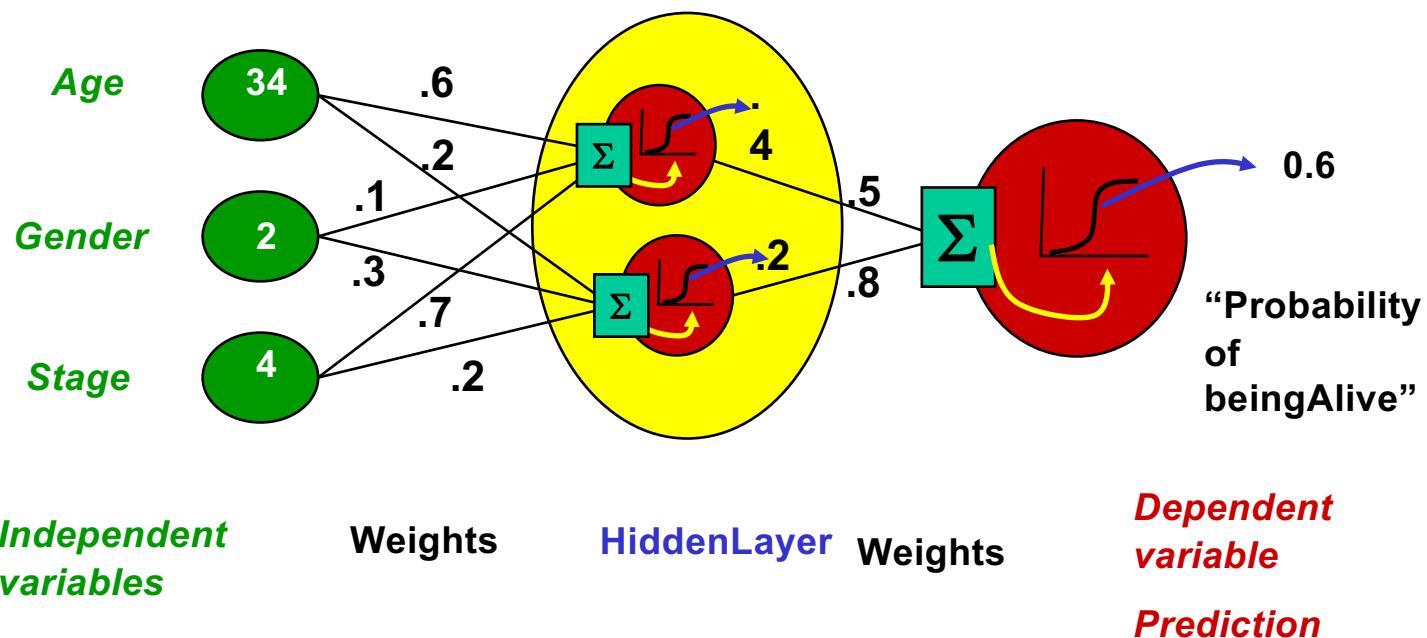






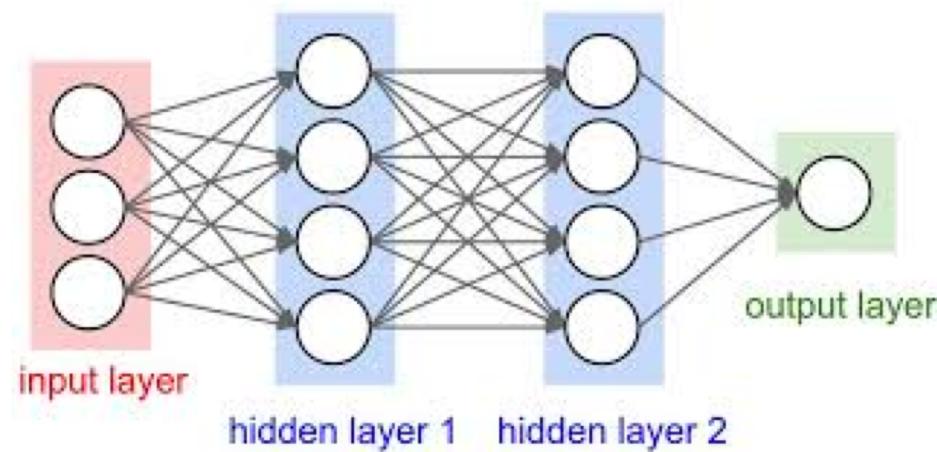
# Not really, no target for hidden units...

Hidden neurons help us model feature interactions  
See manifold section below for more details



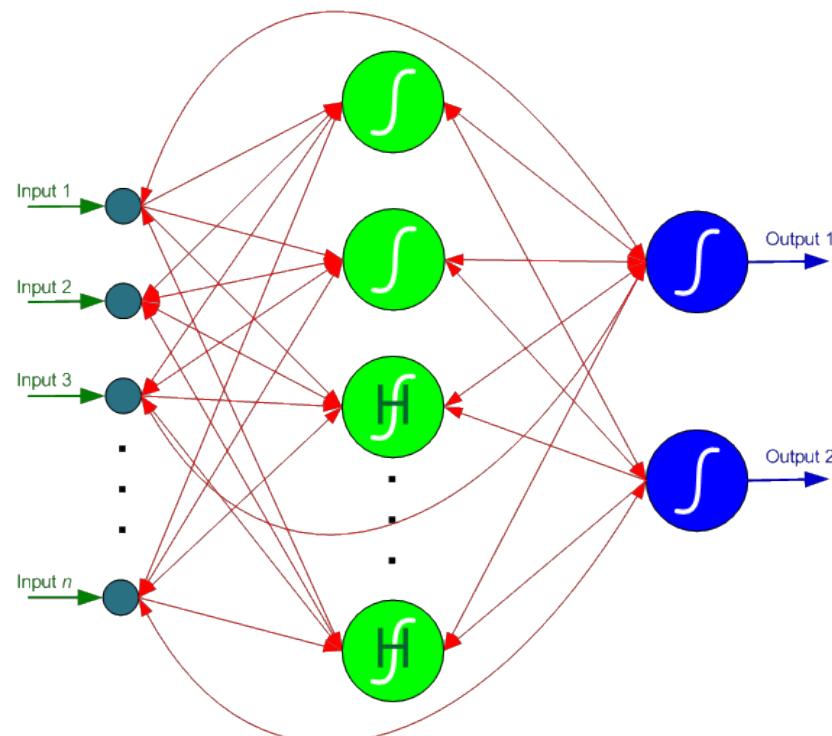
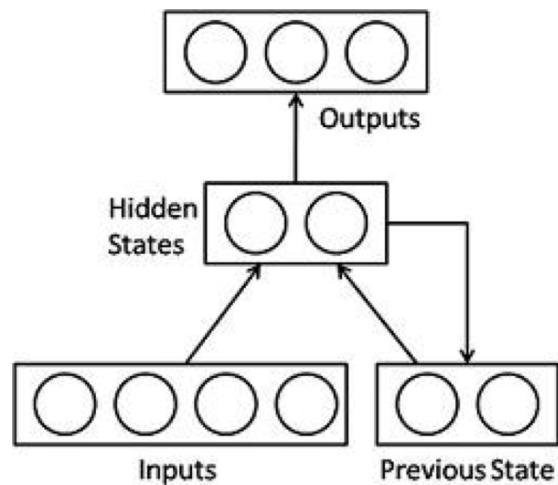
# MLP, feed forward neural network

---



# Neural Network Architectures

## Recurrent Neural Networks



# Neural Network Architectures

---

- **Architecture of a neural network is driven by the task it is intended to address**
  - Classification, regression, clustering, general optimization, association, ....
- **Most popular architecture: Feedforward, multi-layered perceptron with backpropagation learning algorithm**
  - Used for both classification and regression type problems

# Learning Algos

---

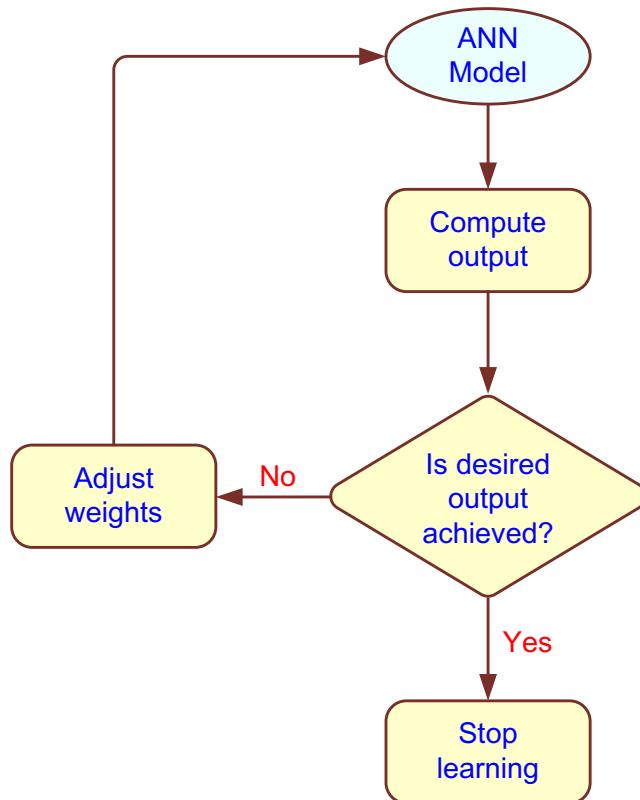
- **Supervised**

- **Unsupervised**

- **Semi-Supervised**

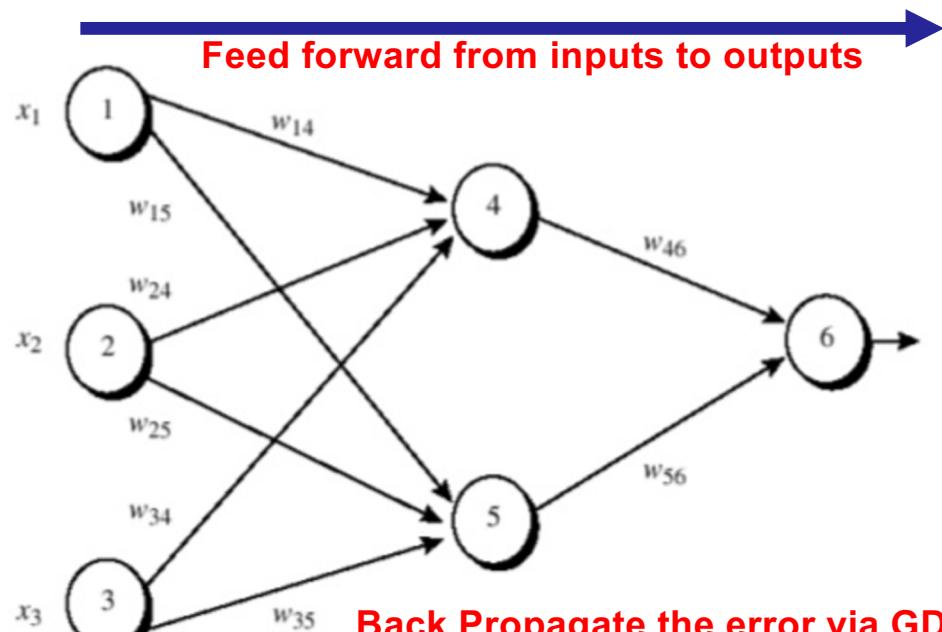
- **BackProp**

# A Supervised Learning Process



## Three-step process:

1. Compute temporary outputs
2. Compare outputs with desired targets
3. Adjust the weights and repeat the process



Initial input, weight, and bias values.

$\Delta$  rule

change weights to  
decrease the error

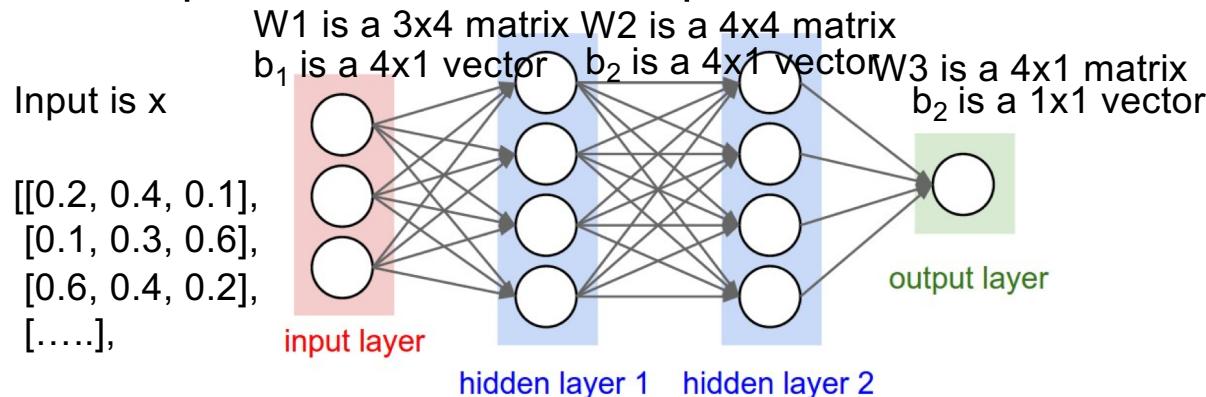
Bias terms for  
nodes 4,5,6

$x_1$	$x_2$	$x_3$	$w_{14}$	$w_{15}$	$w_{24}$	$w_{25}$	$w_{34}$	$w_{35}$	$w_{46}$	$w_{56}$	$\theta_4$	$\theta_5$	$\theta_6$
1	0	1	0.2	-0.3	0.4	0.1	-0.5	0.2	-0.3	-0.2	-0.4	0.2	0.1
*													

Unit $j$	Net input, $I_j$	Output, $O_j$
4	$0.2 + 0 - 0.5 - 0.4 = -0.7$	$1/(1 + e^{-0.7}) = 0.332$
5	$-0.3 + 0 + 0.2 + 0.2 = 0.1$	$1/(1 + e^{-0.1}) = 0.525$
6	$(-0.3)(0.332) - (0.2)(0.525) + 0.1 = -0.105$	$1/(1 + e^{0.105}) = 0.474$

# Layers enable vectorized computations

Example Feed-forward computation of a Neural Network



]

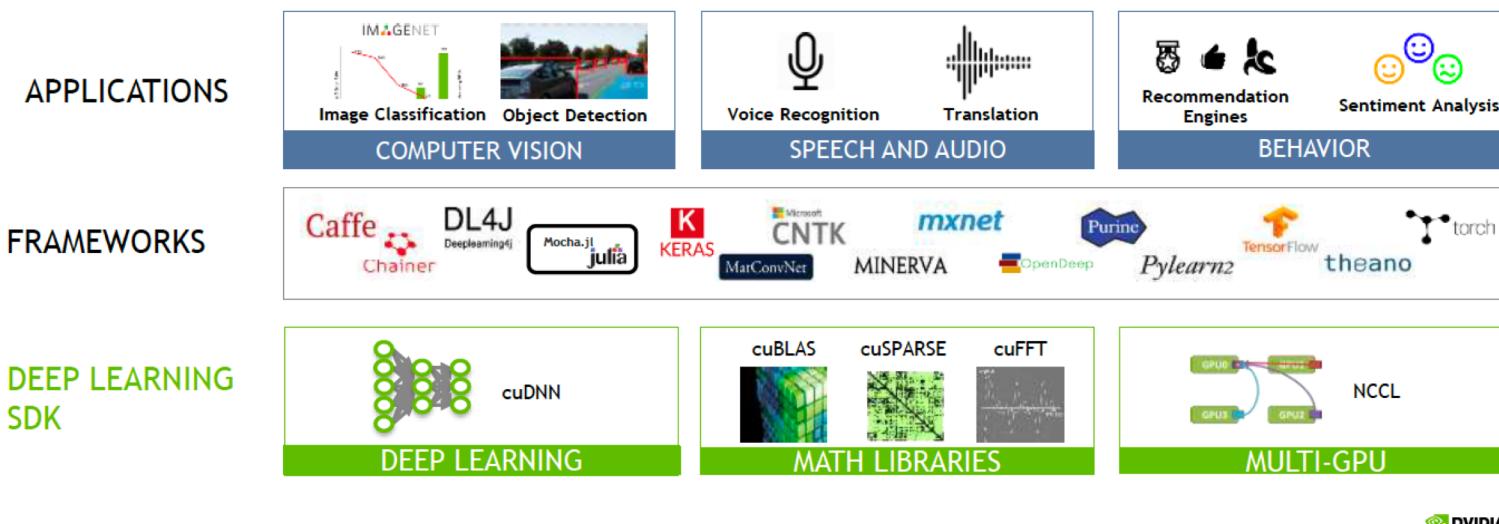
```
# forward-pass of a 3-layer neural network:  
f = lambda x: 1.0/(1.0 + np.exp(-x)) # activation function (use sigmoid)  
x = np.random.randn(3, 1) # random input vector of three numbers (3x1)  
h1 = f(np.dot(W1, x) + b1) # calculate first hidden layer activations (4x1)  
h2 = f(np.dot(W2, h1) + b2) # calculate second hidden layer activations (4x1)  
out = np.dot(W3, h2) + b3 # output neuron (1x1)
```

# Outline

- 1. Introduction**
- 2. ML and deep learning review**
  - 1. Perceptron → MLP
  - 2. Keras
  - 3. BackPropagation
- 3. What is computer vision?**
  - 1. Computer vision
  - 2. Convolutional Neural Nets (CNNs)
- 4. Deep NN Architectures for CV Tasks**
  - 1. Backbone networks
- 5. Solving edge-based IoT**
- 6. Conclusions and Next steps**

# Deep Learning EcoSystem

- High Performance GPU-Acceleration for Deep Learning



- 
- **What is Keras?**
    - Neural Network library written in Python
    - Designed to be minimalistic & straight forward yet extensive Built on top of either Theano or newly TensorFlow
  - **Why use Keras?**
    - Simple to get started, simple to keep going
    - Written in python and highly modular; easy to expand Deep enough to build serious models



# Keras is minimal, modular and extensible

---

- Keras is an [open source neural network](#) library written in [Python](#). It is capable of running on top of either [Tensorflow](#) or [Theano](#).<sup>[1]</sup> Designed to enable fast experimentation with [deep neural networks](#), it focuses on being minimal, modular and extensible.
- The library contains numerous implementations of commonly used neural network building blocks such as layers, [objectives](#), [activation functions](#), [optimizers](#), and a host of tools to make working with image and text data easier.
- As of Sept 16, 2016, Keras is the second fastest growing deep learning framework after Google's Tensorflow.
- As of September 2017, Keras a core part of Tensorflow

# Keras: a model-level library

---

- Keras is a model-level library, providing high-level building blocks for developing deep learning models.
  - It does not handle its own low-level operations such as tensor products, convolutions and so on.
  - Instead, it relies on a specialized, well-optimized tensor manipulation library to do so, serving as the "backend engine" of Keras. Rather than picking one single tensor library and making the implementation of Keras tied to that library, Keras handles the problem in a modular way, and several different backend engines can be plugged seamlessly into Keras.
- At this time, Keras has two backend implementations available:
  - the TensorFlow backend and the Theano backend.
    - [TensorFlow](#) is an open-source symbolic tensor manipulation framework developed by Google, Inc.
    - [Theano](#) is an open-source symbolic tensor manipulation framework developed by LISA/MILA Lab at Université de Montréal.

# Switching from one backend to another

---

If you have run Keras at least once, you will find the Keras configuration file at:

```
~/keras/keras.json
```

If it isn't there, you can create it.

The default configuration file looks like this:

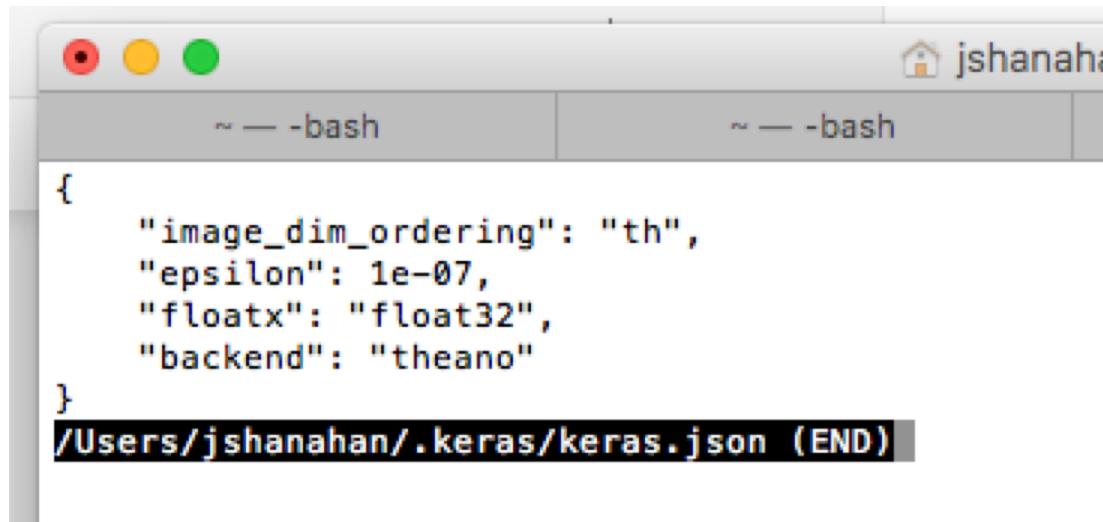
```
{  
    "image_dim_ordering": "tf",  
    "epsilon": 1e-07,  
    "floatx": "float32",  
    "backend": "tensorflow"  
}
```

Simply change the field `backend` to either `"theano"` or `"tensorflow"`, and Keras will use the new configuration next time you run any Keras code.

# Switching from one backend to another

---

- Simply change the field backend to either "theano" or "tensorflow", and Keras will use the new configuration next time you run any Keras code.
- 



The image shows a screenshot of a Mac OS X desktop with two terminal windows side-by-side. Both windows have the title bar "jshanahan" and the path "~ — -bash". The left terminal window displays the JSON configuration file:

```
{  
    "image_dim_ordering": "th",  
    "epsilon": 1e-07,  
    "floatx": "float32",  
    "backend": "theano"  
}
```

The right terminal window shows the command that was run to view the file:

```
/Users/jshanahan/.keras/keras.json (END)
```

# Outline

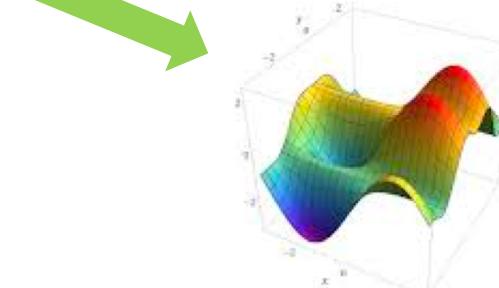
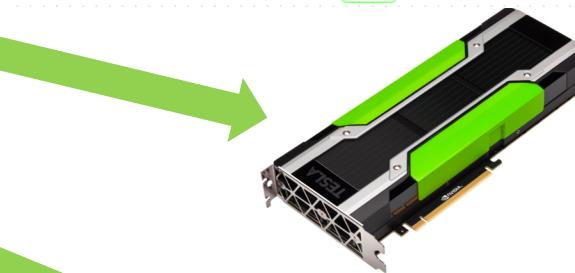
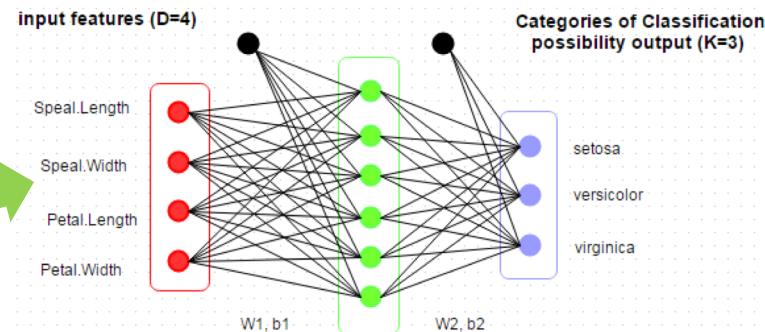
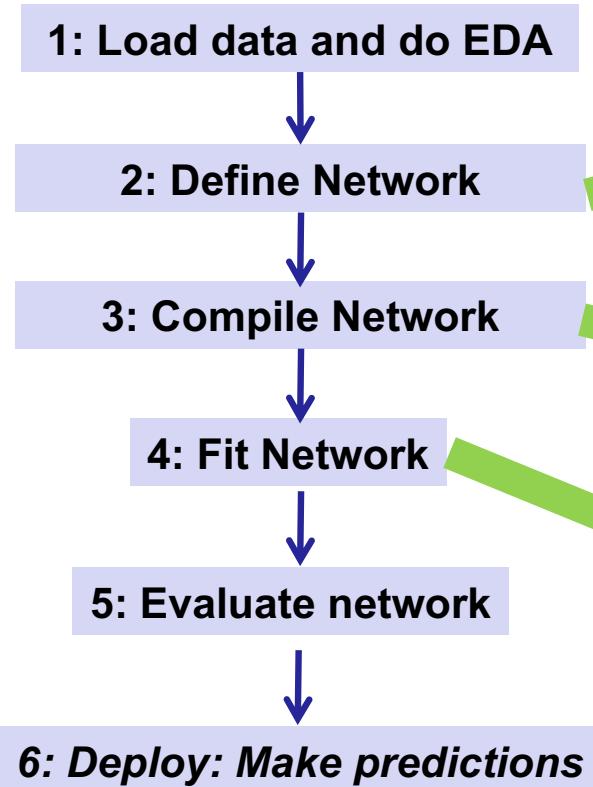
---

- **Neural Networks Introduction**
  - Perceptron
  - MLP
  - Learning via Gradient Descent
- **Keras**
- **5 steps to create a NN/deep learning model**
- **Labs**
- **BackProp Algorithm**

---

# 5 steps to create a NN/deep learning model in Keras

# 5 steps to create a NN/deep learning model



# 5 steps to create a NN/deep learning model

---

- Create a neural network model using the powerful Keras Python library for deep learning

## 5 Steps

1. How to load data.
2. How to define neural network in Keras.
3. How to compile a Keras model using the efficient numerical backend (such as Tensorflow or Theano)
4. How to train a model on data.
5. How to evaluate a model on data.
6. *Make predictions (deploy)*

# Binary cross entropy example

---

- **Keras library**
  - Keras Tutorial
  - Example Notebook
    - Binary cross entropy example
      - <http://machinelearningmastery.com/binary-classification-tutorial-with-the-keras-deep-learning-library/>
    - MLP

# Classification Notebooks

---

- **Pima diabetes: binary classification problem**
  - <http://localhost:8890/notebooks/shared/Dropbox/Projects/UniversityOFGhent-2017/Src/Untitled%20Folder/Pima-Diabetes-First-NN.ipynb>
- **Iris: Three class classification problem**
  - <http://localhost:8890/notebooks/shared/Dropbox/Projects/UniversityOFGhent-2017/Src/U>

Notebook hidden (quiz)

# Main Steps in learning a NN

## 1. Load data and do EDA

```
1 # load pima indians dataset
2 dataset = numpy.loadtxt("pima-indians-diabetes.csv", delimiter=",")
3 # split into input (X) and output (Y) variables
4 X = dataset[:,0:8]
5 Y = dataset[:,8]
```

## 2. Define the model architecture

```
1 # create model
2 model = Sequential()
3 model.add(Dense(12, input_dim=8, init='uniform', activation='relu'))
4 model.add(Dense(8, init='uniform', activation='relu'))
5 model.add(Dense(1, init='uniform', activation='sigmoid'))
```

SoftMax

## 3. Compile the model using a numerical backend (and set objective)

```
1 # Compile model
2 model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
```

## 4. Fit the model

```
1 # Fit the model
2 model.fit(X, Y, nb_epoch=150, batch_size=10)
```

## 5. Evaluate

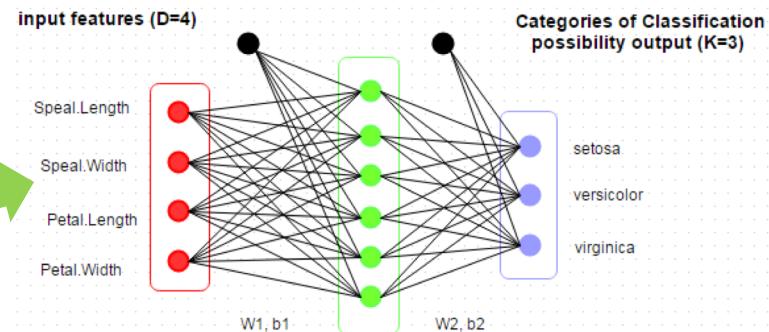
```
1 # evaluate the model
2 scores = model.evaluate(X, Y)
3 print("%s: %.2f%%" % (model.metrics_names[1], scores[1]*100))
```

# 5 steps to create a NN/deep learning model

1: Load data and do EDA



2: Define Network



3: Compile Network



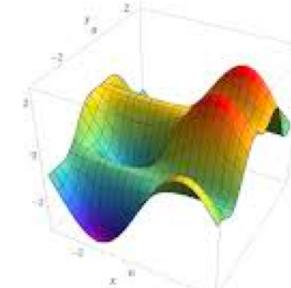
4: Fit Network



5: Evaluate network



6: Make predictions



---

# Datasets

File size: 23.31 KB

Description: Predict the Onset of Diabetes

Attributes:

**Pregnancies:** Number of times pregnant

**Glucose:** Plasma glucose concentration a 2 hours in an oral glucose tolerance test

**BloodPressure:** Diastolic blood pressure (mm Hg)

**SkinThickness:** Triceps skin fold thickness (mm)

**Insulin:** 2-Hour serum insulin (mu U/ml)

**BMI:** Body mass index (weight in kg/(height in m)^2)

**DiabetesPedigreeFunction:** Diabetes pedigree function

**Age:** Age (years)

**Outcome:** Class variable (0 or 1)

# PIMA diabetes

Table [768 females, 9]

<https://www.kaggle.com/uciml/pima-indians-diabetes-database>



File Preview:

① Edit column descriptions ▾

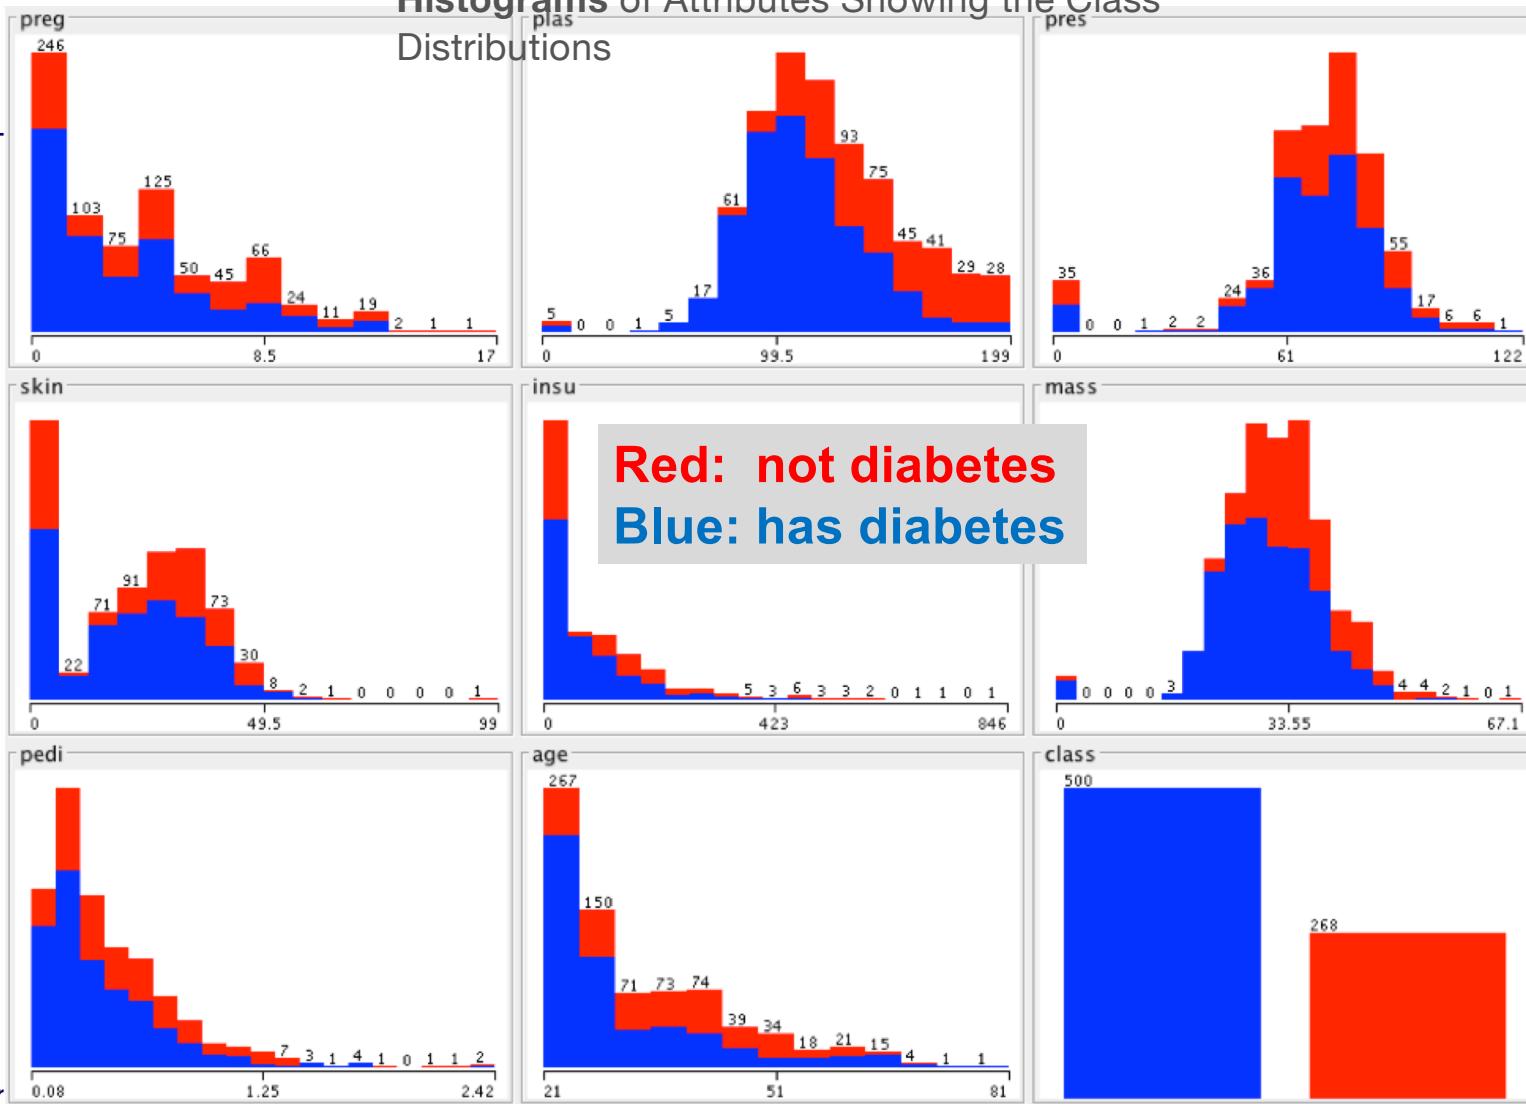
Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
6	148	72	35	0	33.6	0.627	50	1
1	85	66	29	0	26.6	0.351	31	0
8	183	64	0	0	23.3	0.672	32	1
1	89	66	23	94	28.1	0.167	21	0
0	137	40	35	168	43.1	2.288		

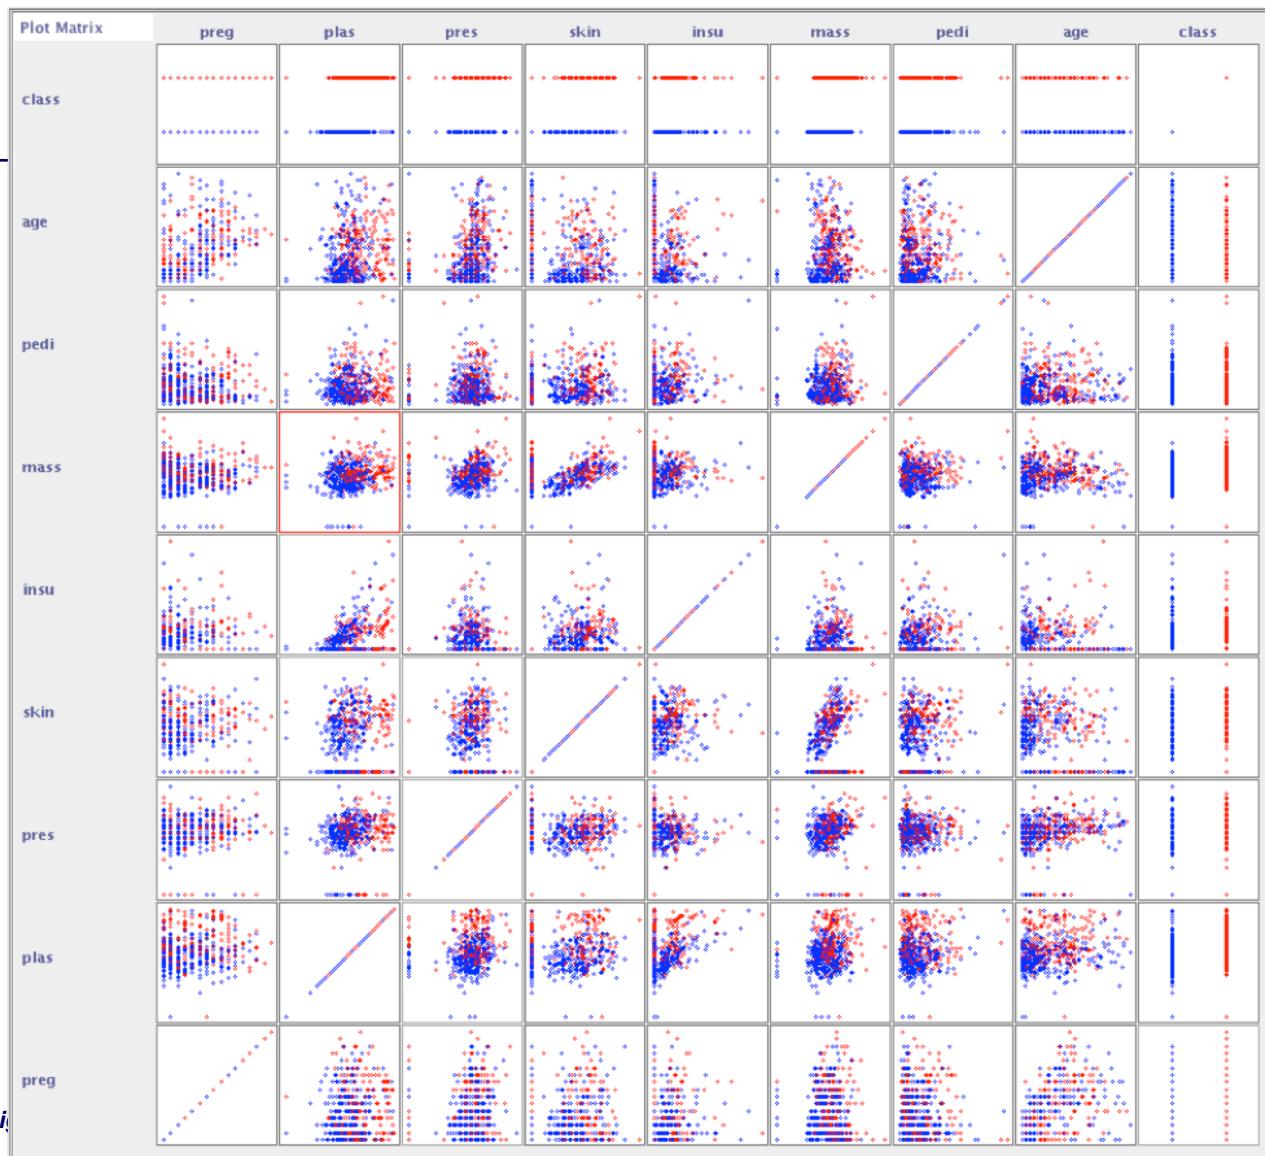
# Evaluation of various ML Algorithms on PIMA Diabetes

---

- See this article
  - <http://machinelearningmastery.com/case-study-predicting-the-onset-of-diabetes-within-five-years-part-2-of-3/>

Histograms of Attributes Showing the Class Distributions

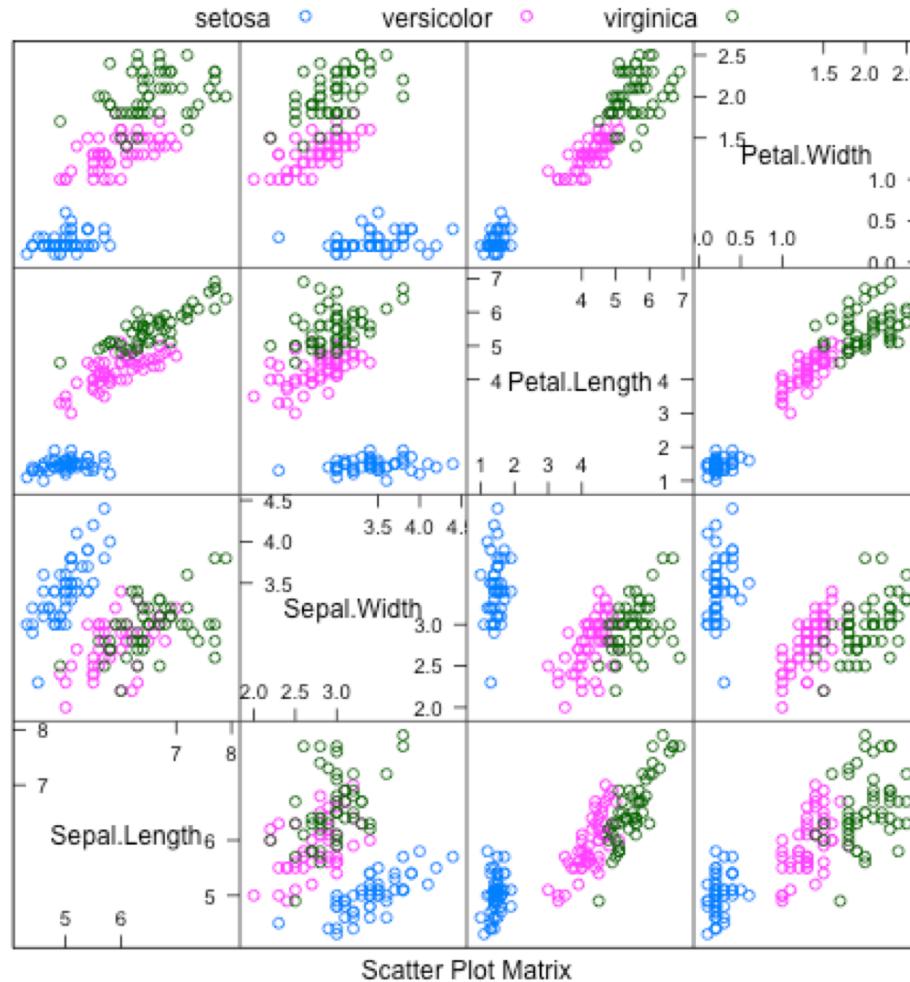




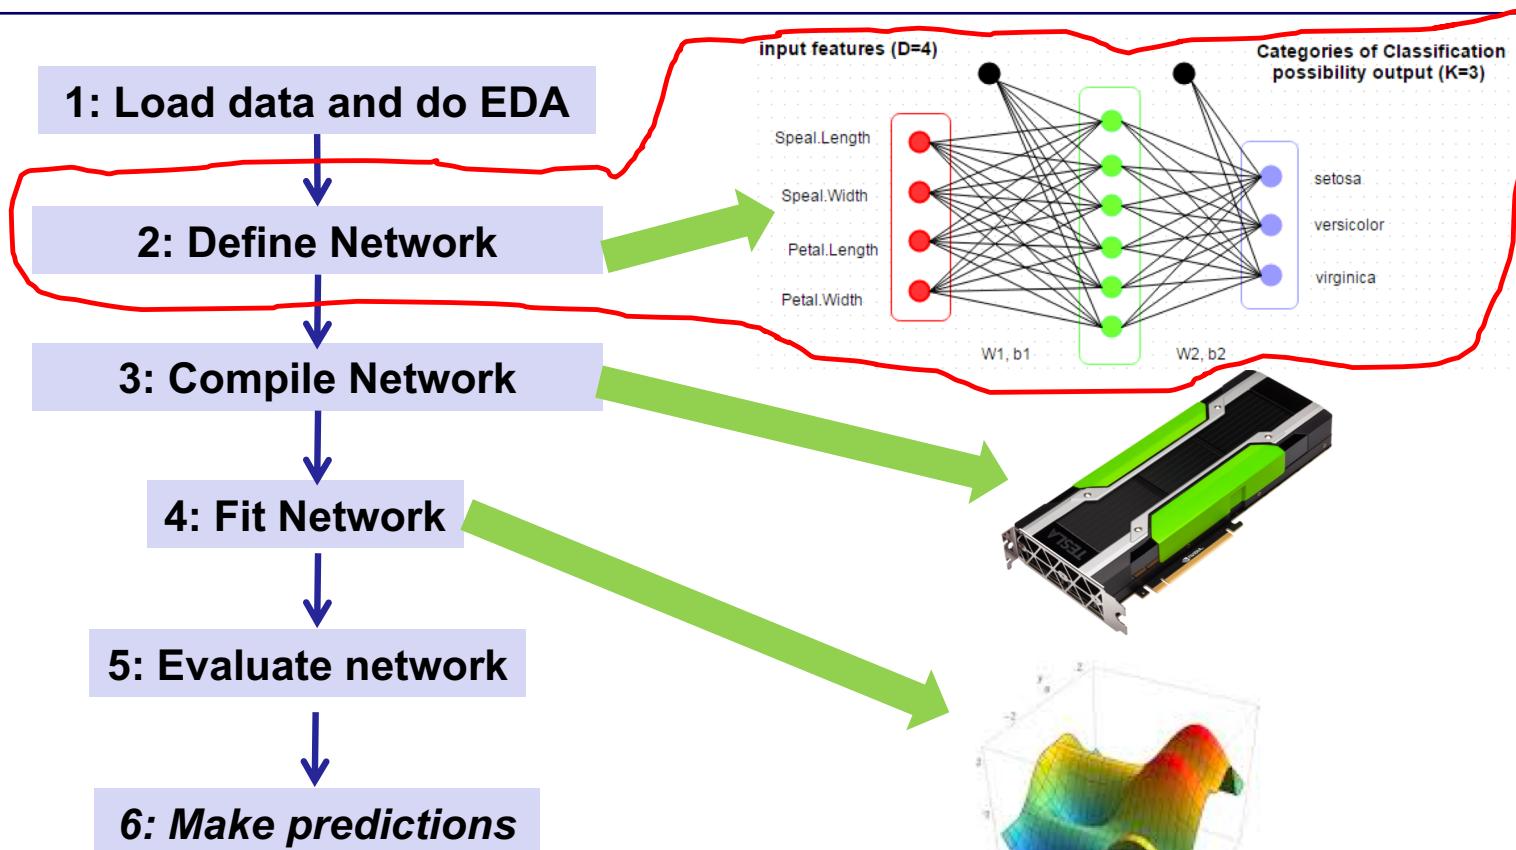
## Iris Dataset: 4 real-val<sup>d</sup>. inputs; 150 examples; 3 categories

---





# 5 steps to create a NN/deep learning model



# Define your neural network in Keras: Sequence container PLUS Layers

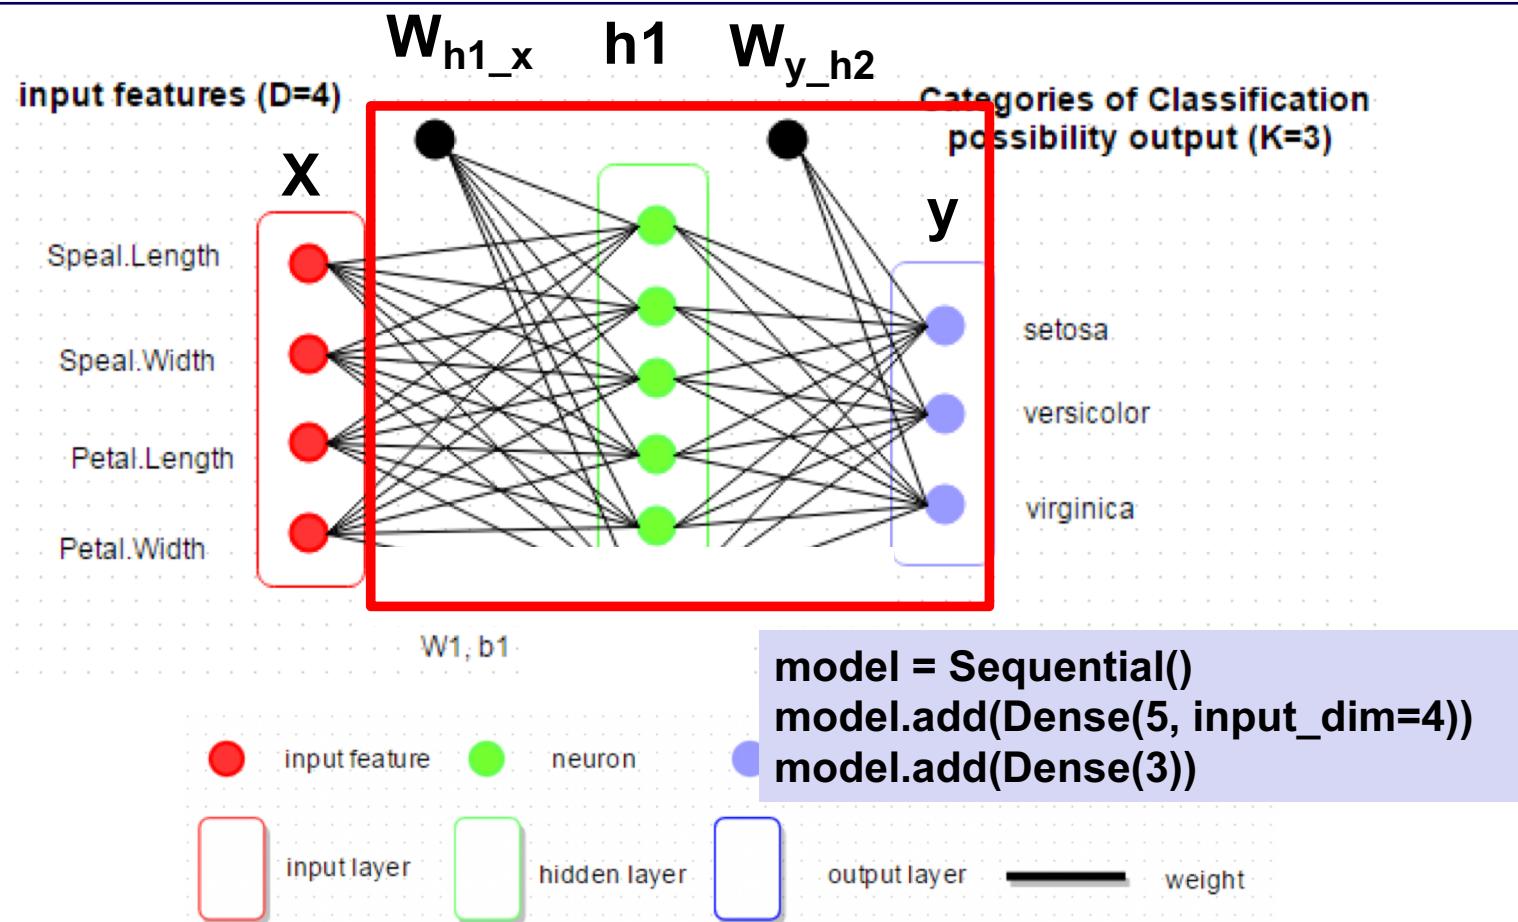
---

- **The first step is to define your neural network as a Sequence container PLUS Layers.**
  - Sequence of layers
    - Neural networks are defined in Keras as a sequence of layers.
  - Sequential container
    - The container for these layers is the Sequential class.
- 1. **The first step is to create an instance of the Sequential class.**
- 2. **Then you can create your layers and add them in the order that they should be connected.**
- **For example,**
  - a small Multilayer Perceptron model with 2 inputs in the visible layer, 5 neurons in the hidden layer and one neuron in the output layer can be defined as:

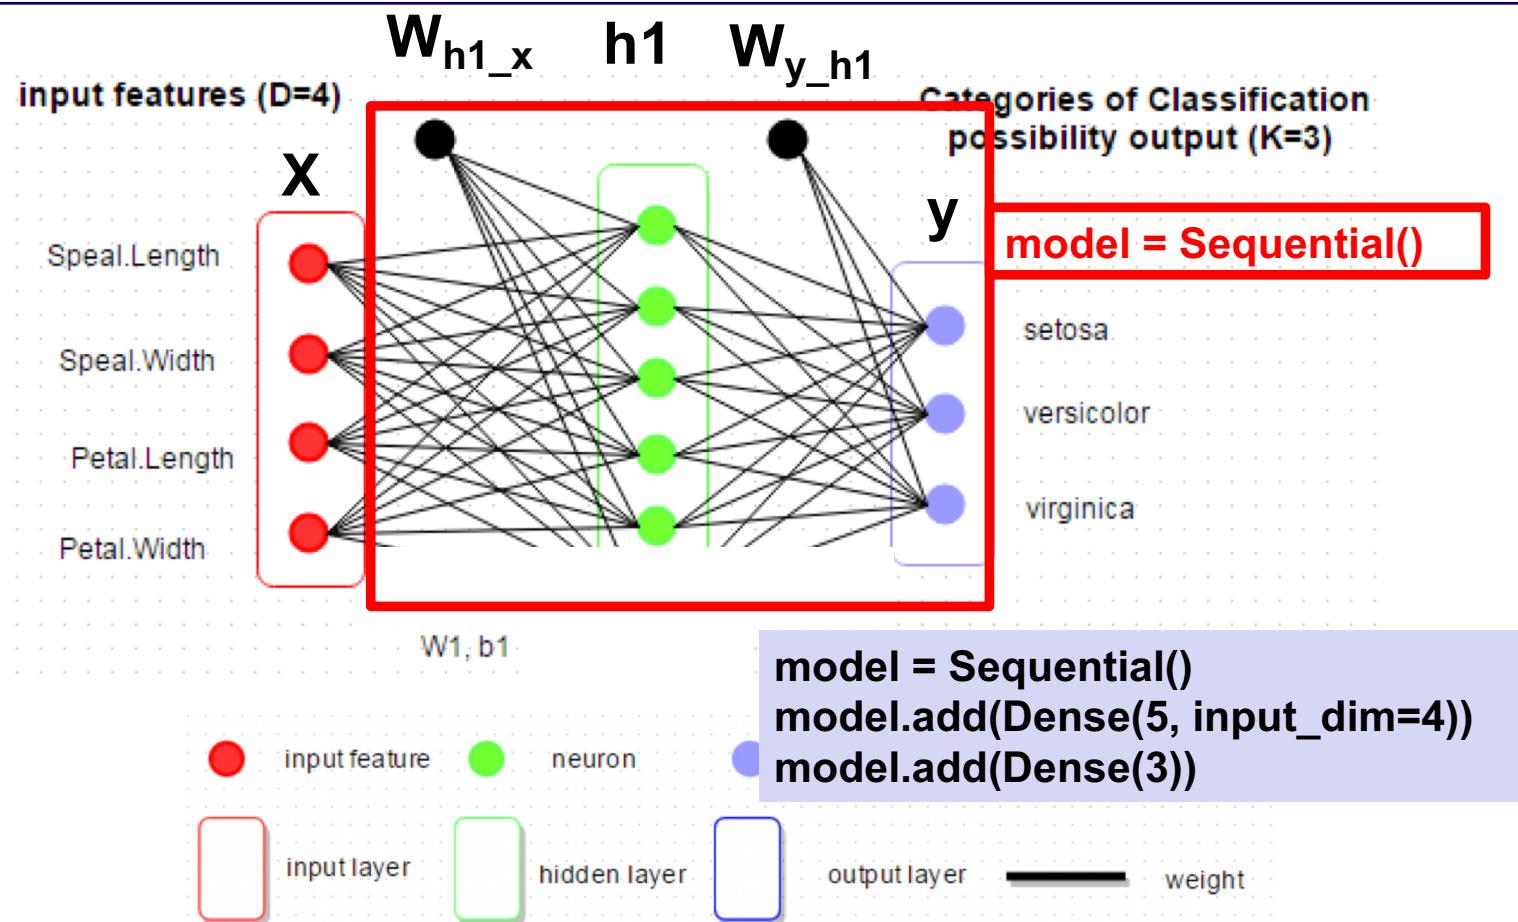
2-5-1

```
1 model = Sequential()  
2 model.add(Dense(5, input_dim=2))  
3 model.add(Dense(1))
```

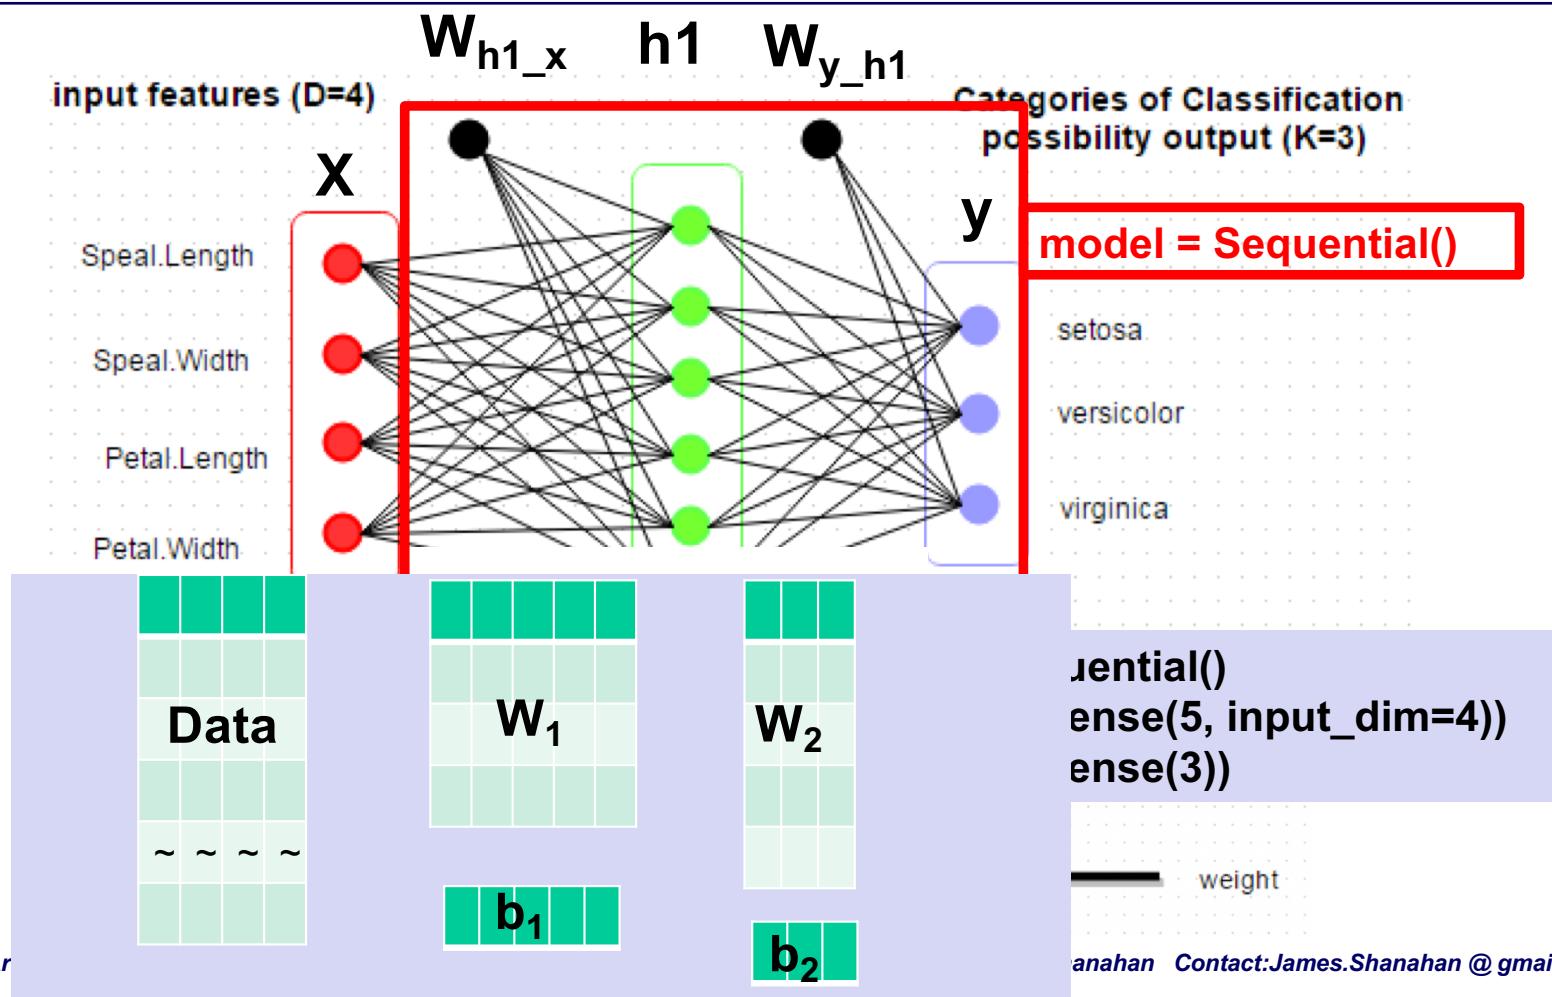
# Let's work with a 4 – 5 – 3 MLP



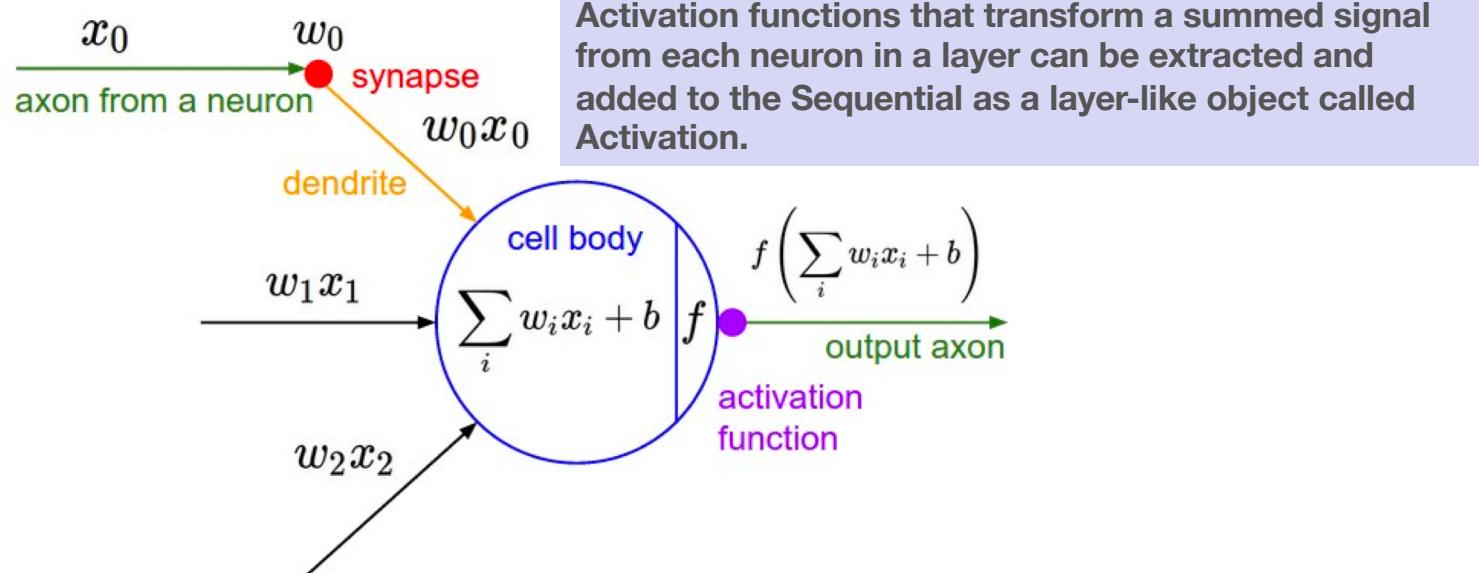
# Let's work with a 4 – 5 – 3 MLP



# Let's work with a 4 – 5 – 3 MLP

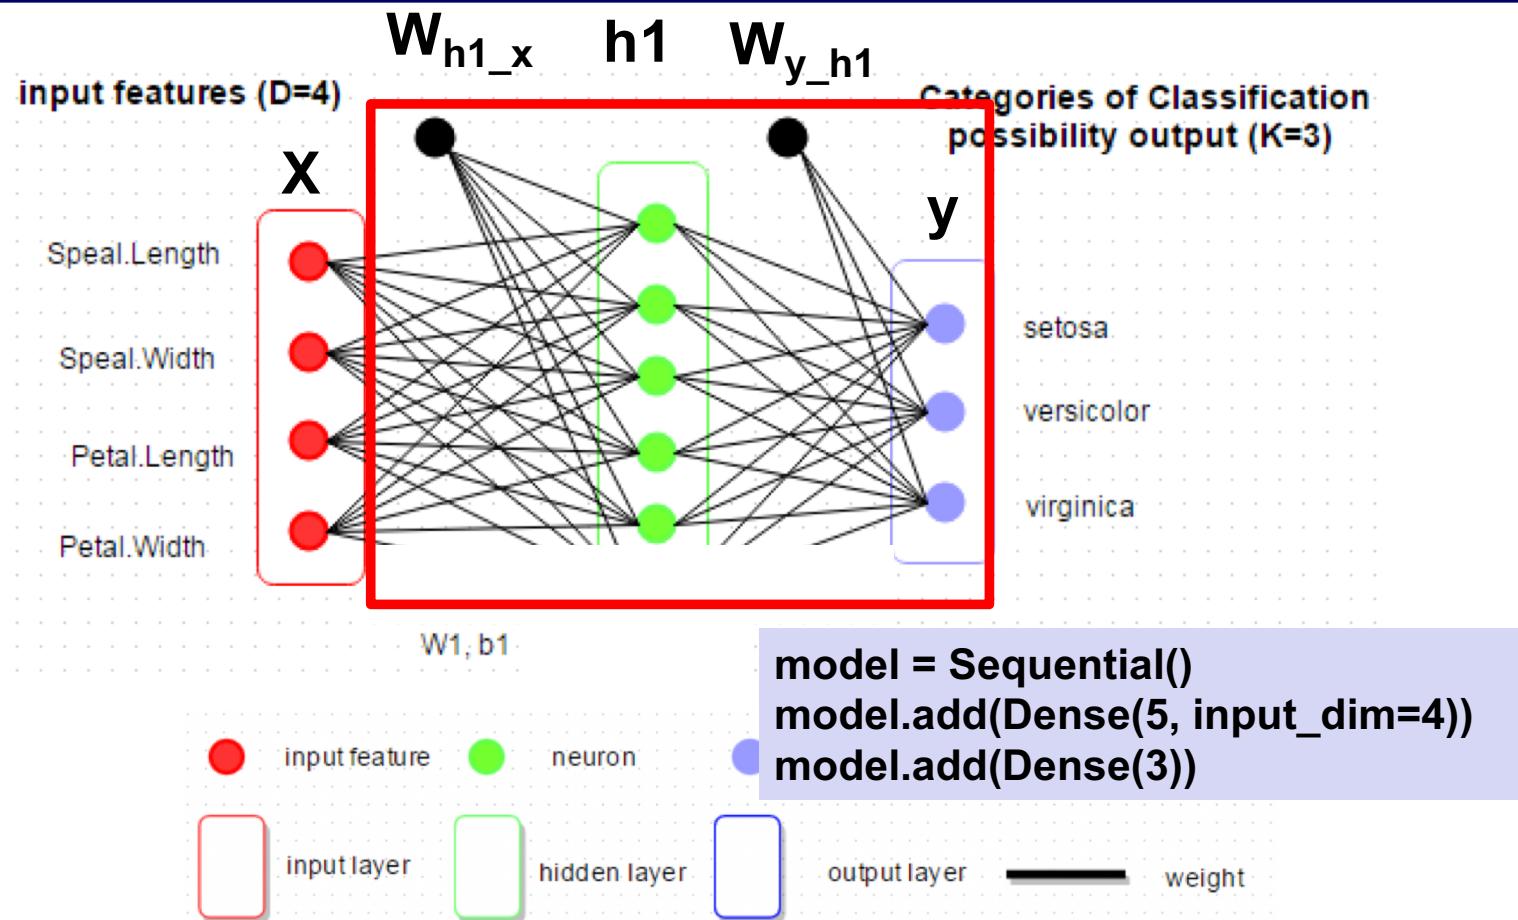


# Layer in the activation functions



```
1 model = Sequential()  
2 model.add(Dense(5, input_dim=2))  
3 model.add(Activation('relu'))  
4 model.add(Dense(1))  
5 model.add(Activation('sigmoid'))
```

# Activation functions for hidden outer nodes



# Activation functions for hidden nodes versus outer nodes

---

- **Inner node activation functions**
  - Sigmoid, Tanh, ReLU, etc.
- **Outer node activation functions (and loss function objective)**
  - **Regression**: Linear activation function or ‘linear’ and the number of neurons matching the number of outputs.
  - **Binary Classification** (2 class): Logistic activation function or ‘sigmoid’ and one neuron the output layer.
  - **Multiclass Classification** (>2 class): Softmax activation function ‘Cross entropy’ and one output neuron per class value, assuming a one-hot encoded output pattern.

Cross entropy

# **Output layer activation function: Regression versus Classification**

---

- **The choice of activation function is most important for the output layer as it will define the format that predictions will take.**
- **For example, some common predictive modeling problem types and the structure and standard activation function that you can use in the output layer:**
  - **Regression:** Linear activation function or ‘linear’ and the number of neurons matching the number of outputs.
  - **Binary Classification** (2 class): Logistic activation function or ‘sigmoid’ and one neuron the output layer.
  - **Multiclass Classification** (>2 class): Softmax activation function or ‘softmax’ and one output neuron per class value, assuming a one-hot encoded output pattern.

# Main Steps in learning a NN

## 1. Load data and do EDA

```
1 # load pima indians dataset
2 dataset = numpy.loadtxt("pima-indians-diabetes.csv", delimiter=",")
3 # split into input (X) and output (Y) variables
4 X = dataset[:,0:8]
5 Y = dataset[:,8]
```

## 2. Define the model architecture

```
1 # create model
2 model = Sequential()
3 model.add(Dense(12, input_dim=8, init='uniform', activation='relu'))
4 model.add(Dense(8, init='uniform', activation='relu'))
5 model.add(Dense(1, init='uniform', activation='sigmoid'))
```

SoftMax

## 3. Compile the model using a numerical backend (and set objective)

```
1 # Compile model
2 model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
```

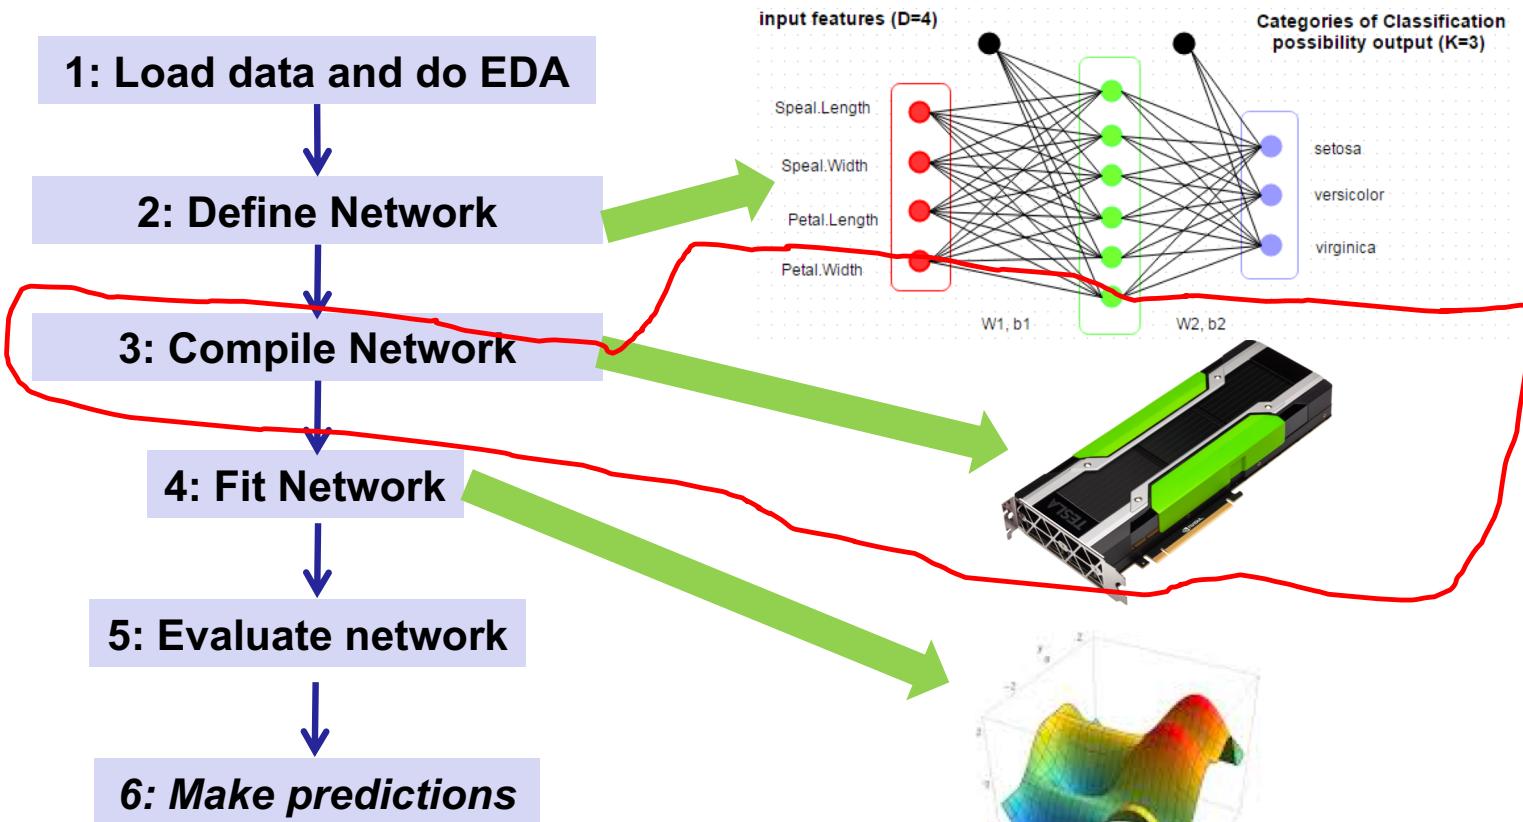
## 4. Fit the model

```
1 # Fit the model
2 model.fit(X, Y, nb_epoch=150, batch_size=10)
```

## 5. Evaluate

```
1 # evaluate the model
2 scores = model.evaluate(X, Y)
3 print("%s: %.2f%%" % (model.metrics_names[1], scores[1]*100))
```

# 5 steps to create a NN/deep learning model



## Step 3: Compile the model

---

- **Compile the model using a numerical backend (and set objective)**
- **Compilation is an efficiency step.**
  - It transforms the simple sequence of layers that we defined into a highly efficient series of matrix transforms in a format intended to be executed on your GPU or CPU, depending on how Keras is configured.
- **Think of compilation as a precompute step for your network.**

# Compilation is done before training

---

- **Compilation is always required after defining a model**
  1. Compile to fit hardware
  2. Optimization and Loss
- **Compilation is done before training and before transfer learning (CPU/GPU efficiencies)**
  - This includes both before training it using an optimization scheme as well as loading a set of pre-trained weights from a save file.
  - The reason is that the compilation step prepares an efficient representation of the network that is also required to make predictions on your hardware.
- **Compiling requires the optimization algorithm and loss**
  - Specify the optimization algorithm to use to train the network and
  - Specify the loss function used to evaluate the network that is minimized by the optimization algorithm.

# Compilation is done before training

---

- **Compilation is always required after defining a model.**
- **Compilation is done before training and before transfer learning (CPU/GPU efficiencies)**
  - This includes both before training it using an optimization scheme as well as loading a set of pre-trained weights from a save file.
  - The reason is that the compilation step prepares an efficient representation of the network that is also required to make predictions on your hardware.
- **Compiling requires the optimization algorithm and loss**
  - Specify the optimization algorithm to use to train the network and
  - Specify the loss function used to evaluate the network that is minimized by the optimization algorithm

```
model.compile(optimizer='sgd', loss='mse')
```

# Loss functions for DL

---

- The type of predictive modeling problem imposes constraints on the type of loss function that can be used.
- For example, below are some standard loss functions for different predictive model types:
  - Regression: Mean Squared Error or '**mse**'.
  - Binary Classification (2 class): Logarithmic Loss, also called cross entropy or '**binary\_crossentropy**'.
  - Multiclass Classification (>2 class): Multiclass Logarithmic Loss or '**categorical\_crossentropy**' (CX).

```
model.compile(optimizer='sgd', loss='mse')
```

# Objective Functions (for output layer)

---

- Objectives
  - <https://keras.io/objectives>

## Available objectives

- `mean_squared_error` / `mse`
- `mean_absolute_error` / `mae`
- `mean_absolute_percentage_error` / `mape`
- `mean_squared_logarithmic_error` / `msle`
- `squared_hinge`
- `hinge`
- `binary_crossentropy`: Also known as logloss.
- `categorical_crossentropy`: Also known as multiclass logloss. **Note:** using this objective requires that your labels are binary arrays of shape `(nb_samples, nb_classes)`.
- `sparse_categorical_crossentropy`: As above but accepts sparse labels. **Note:** this objective still requires that your labels have the same number of dimensions as your outputs; you may need to add a length-1 dimension to the shape of your labels, e.g with `np.expand_dims(y, -1)`.
- `kullback_leibler_divergence` / `kld`: Information gain from a predicted probability distribution Q to a true probability distribution P. Gives a measure of difference between both distributions.
- `poisson`: Mean of `(predictions - targets * log(predictions))`
- `cosine_proximity`: The opposite (negative) of the mean cosine proximity between predictions and targets.

# One hot encode (OHE) multi-class output variable

---

Note: when using the `categorical_crossentropy` objective, your targets should be in categorical format (e.g. if you have 10 classes, the target for each sample should be a 10-dimensional vector that is all-zeros except for a 1 at the index corresponding to the class of the sample). In order to convert *integer targets* into *categorical targets*, you can use the Keras utility `to_categorical`:

```
from keras.utils.np_utils import to_categorical  
  
categorical_labels = to_categorical(int_labels, nb_classes=None)
```

# Optimization Algorithms and hyperparameters

---

- Perhaps the most commonly used optimization algorithms because of their generally better performance are:
  - **Stochastic Gradient Descent** or ‘*sgd*’ that requires the tuning of a learning rate
  - **ADAM** or ‘*adam*’ that requires the tuning of learning rate and momentum.
  - **RMSprop** or ‘*rmsprop*’ that requires the tuning of learning rate and momentum.

# Track learning by collecting metrics

---

- Finally, you can also specify metrics to collect while fitting your model in addition to the loss function. Generally, the most useful additional metric to collect is accuracy for classification problems.
- The metrics to collect are specified by name in an array.

```
1     model.compile(optimizer='sgd', loss='mse', metrics=['accuracy'])
```

# Main Steps in learning a NN

## 1. Load data and do EDA

```
1 # load pima indians dataset
2 dataset = numpy.loadtxt("pima-indians-diabetes.csv", delimiter=",")
3 # split into input (X) and output (Y) variables
4 X = dataset[:,0:8]
5 Y = dataset[:,8]
```

## 2. Define the model architecture

```
1 # create model
2 model = Sequential()
3 model.add(Dense(12, input_dim=8, init='uniform', activation='relu'))
4 model.add(Dense(8, init='uniform', activation='relu'))
5 model.add(Dense(1, init='uniform', activation='sigmoid'))
```

SoftMax

## 3. Compile the model using a numerical backend (and set objective)

```
1 # Compile model
2 model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
```

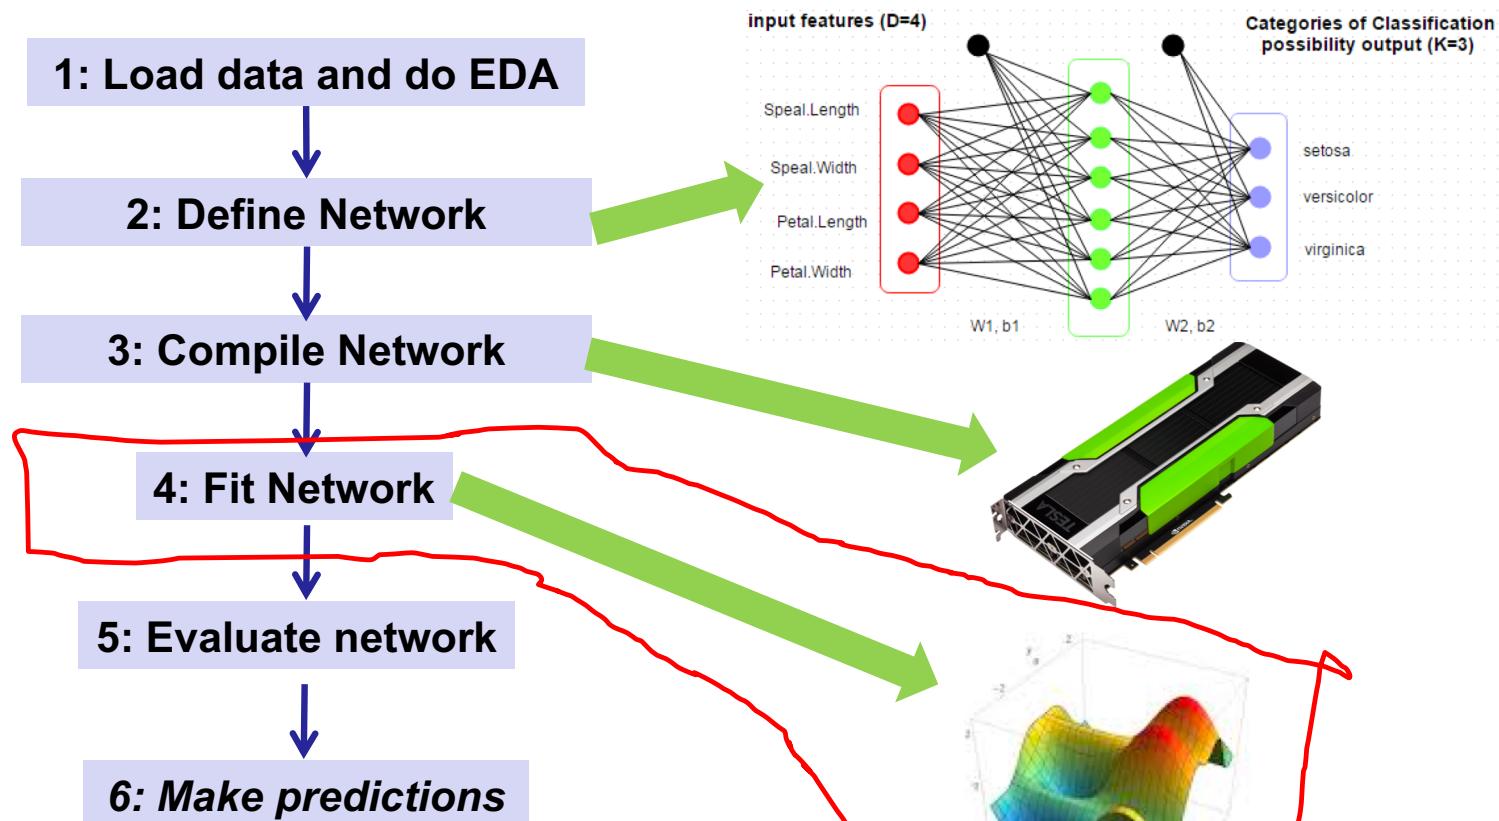
## 4. Fit/Train the model

```
1 # Fit the model
2 model.fit(X, Y, nb_epoch=150, batch_size=10)
```

## 5. Evaluate

```
1 # evaluate the model
2 scores = model.evaluate(X, Y)
3 print("%s: %.2f%%" % (model.metrics_names[1], scores[1]*100))
```

# 5 steps to create a NN/deep learning model

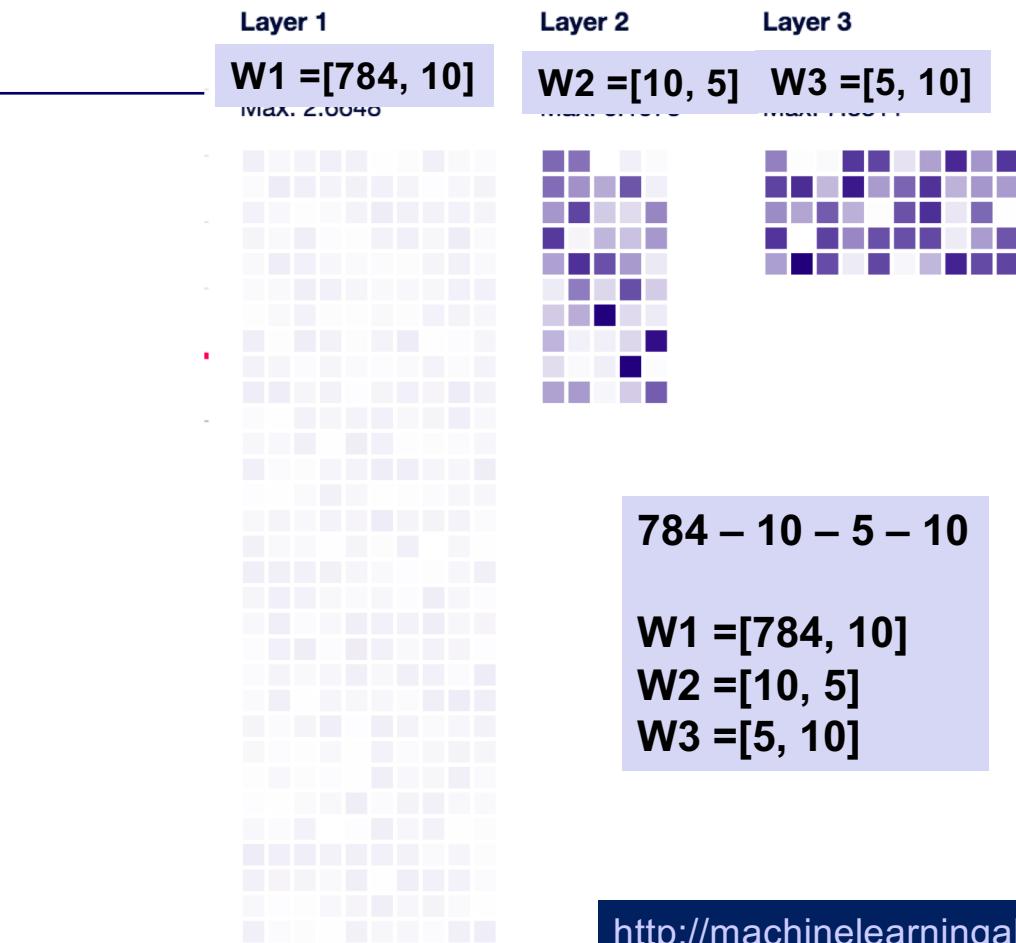


## Step 4: Fit the model(AKA Train)

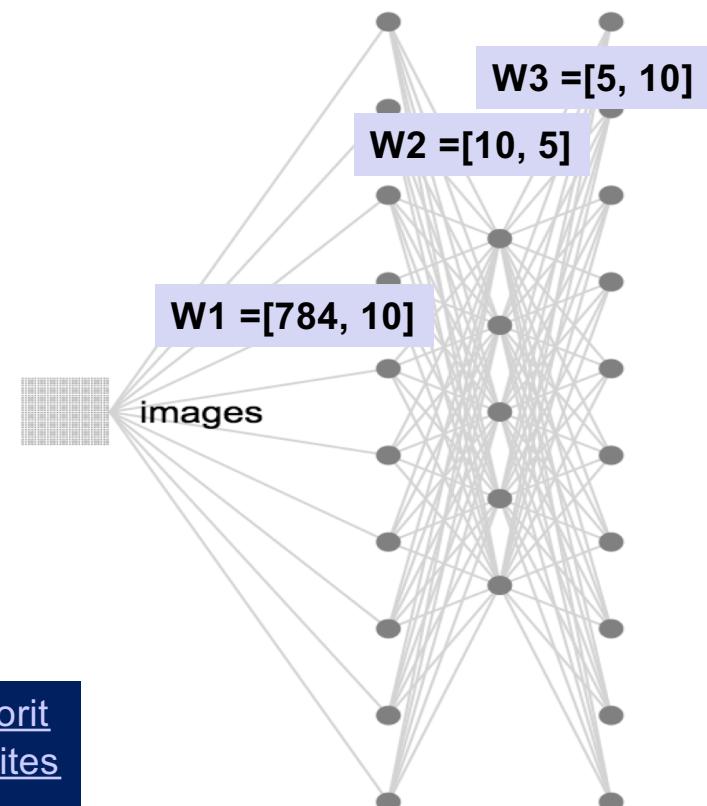
---

- Once the network is compiled, it can be fit, which means adapt the weights on a training dataset.
- The network is trained using the backpropagation algorithm and optimized according to the optimization algorithm and loss function specified when compiling the model.

```
1     history = model.fit(X, y, batch_size=10, nb_epoch=100)
```



## Weight Matrices for image-based input



# Main Steps in learning a NN

## 1. Load data and do EDA

```
1 # load pima indians dataset
2 dataset = numpy.loadtxt("pima-indians-diabetes.csv", delimiter=",")
3 # split into input (X) and output (Y) variables
4 X = dataset[:,0:8]
5 Y = dataset[:,8]
```

## 2. Define the model architecture

```
1 # create model
2 model = Sequential()
3 model.add(Dense(12, input_dim=8, init='uniform', activation='relu'))
4 model.add(Dense(8, init='uniform', activation='relu'))
5 model.add(Dense(1, init='uniform', activation='sigmoid'))
```

SoftMax

## 3. Compile the model using a numerical backend (and set objective)

```
1 # Compile model
2 model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
```

## 4. Fit/Train the model

```
1 # Fit the model
2 model.fit(X, Y, nb_epoch=150, batch_size=10)
```

## 5. Evaluate

```
1 # evaluate the model
2 scores = model.evaluate(X, Y)
3 print("%s: %.2f%%" % (model.metrics_names[1], scores[1]*100))
```

# Binary Classification Notebook

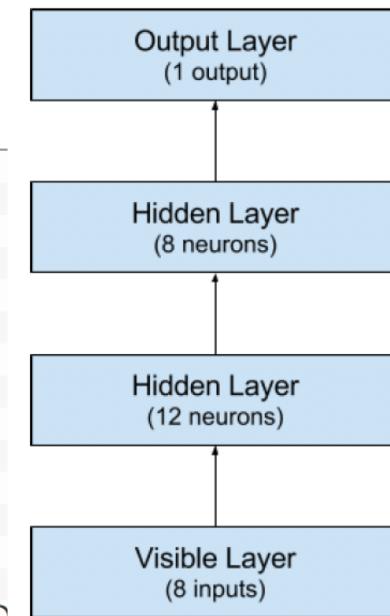
---

- **Pima diabetes: binary classification problem**
  - <http://localhost:8890/notebooks/shared/Dropbox/Projects/UniversityOFGhent-2017/Src/Untitled%20Folder/Pima-Diabetes-First-NN.ipynb>
- <http://localhost:8892/notebooks/shared/Dropbox/Projects/UniversityOFGhent-2017/Src/01-Pima-Diabetes-First-NN/Pima-Diabetes-First-NN.ipynb>

# Pima Diabetes Detector: Binary Classifier

---

```
1 # Create first network with Keras
2 from keras.models import Sequential
3 from keras.layers import Dense
4 import numpy
5 # fix random seed for reproducibility
6 seed = 7
7 numpy.random.seed(seed)
8 # load pima indians dataset
9 dataset = numpy.loadtxt("pima-indians-diabetes.csv", delimiter=",")
10 # split into input (X) and output (Y) variables
11 X = dataset[:,0:8]
12 Y = dataset[:,8]
13 # create model
14 model = Sequential()
15 model.add(Dense(12, input_dim=8, init='uniform', activation='relu'))
16 model.add(Dense(8, init='uniform', activation='relu'))
17 model.add(Dense(1, init='uniform', activation='sigmoid'))
18 # Compile model
19 model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
20 # Fit the model
21 model.fit(X, Y, nb_epoch=150, batch_size=10)
22 # evaluate the model
23 scores = model.evaluate(X, Y)
24 print("%s: %.2f%%" % (model.metrics_names[1], scores[1]*100))
```



---

```
1 ...
2 Epoch 143/150
3 768/768 [=====] - 0s - loss: 0.4614 - acc: 0.7878
4 Epoch 144/150
5 768/768 [=====] - 0s - loss: 0.4508 - acc: 0.7969
6 Epoch 145/150
7 768/768 [=====] - 0s - loss: 0.4580 - acc: 0.7747
8 Epoch 146/150
9 768/768 [=====] - 0s - loss: 0.4627 - acc: 0.7812
10 Epoch 147/150
11 768/768 [=====] - 0s - loss: 0.4531 - acc: 0.7943
12 Epoch 148/150
13 768/768 [=====] - 0s - loss: 0.4656 - acc: 0.7734
14 Epoch 149/150
15 768/768 [=====] - 0s - loss: 0.4566 - acc: 0.7839
16 Epoch 150/150
17 768/768 [=====] - 0s - loss: 0.4593 - acc: 0.7839
18 768/768 [=====] - 0s
19 acc: 79.56%
```

# Main Steps in learning a NN

## 1. Load data and do EDA

```
1 # load pima indians dataset
2 dataset = numpy.loadtxt("pima-indians-diabetes.csv", delimiter=",")
3 # split into input (X) and output (Y) variables
4 X = dataset[:,0:8]
5 Y = dataset[:,8]
```

## 2. Define the model architecture

```
1 # create model
2 model = Sequential()
3 model.add(Dense(12, input_dim=8, init='uniform', activation='relu'))
4 model.add(Dense(8, init='uniform', activation='relu'))
5 model.add(Dense(1, init='uniform', activation='sigmoid'))
```

SoftMax

## 3. Compile the model using a numerical backend (and set objective)

```
1 # Compile model
2 model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
```

## 4. Fit/Train the model

```
1 # Fit the model
2 model.fit(X, Y, nb_epoch=150, batch_size=10)
```

## 5. Evaluate

```
1 # evaluate the model
2 scores = model.evaluate(X, Y)
3 print("%s: %.2f%%" % (model.metrics_names[1], scores[1]*100))
```

# Vanilla logistic regression

---

## 4-3 for iris data

have 3 output units.

```
: inp = Input(shape=(4,))
out = Dense(3, activation="softmax")(inp)
```

Multinomial Logistic Regression uses softmax activation in multiclass case instead of binary logit function. Also we use categorical\_crossentropy as a loss function.

```
: model = Model(inputs=[inp], outputs=[out])
model.compile(loss="categorical_crossentropy", optimizer="sgd", metrics=["accuracy"])
```

Split into train and test set

# Python For Data Science Cheat Sheet

## Keras

Learn Python for data science interactively at [www.DataCamp.com](http://www.DataCamp.com)



### Keras

Keras is a powerful and easy-to-use deep learning library for Theano and TensorFlow that provides a high-level neural networks API to develop and evaluate deep learning models.

#### A Basic Example

```
>>> import numpy as np
>>> from keras.models import Sequential
>>> from keras.layers import Dense
>>> data = np.random.random((1000,100))
>>> labels = np.random.randint(2,size=(1000,1))
>>> model = Sequential()
>>> model.add(Dense(32,
    activation='relu',
    input_dim=100))
>>> model.add(Dense(1, activation='sigmoid'))
>>> model.compile(optimizer='rmsprop',
    loss='binary_crossentropy',
    metrics=['accuracy'])
>>> model.fit(data,labels,epochs=10,batch_size=32)
>>> predictions = model.predict(data)
```

#### Data

#### Also see NumPy, Pandas & Scikit-Learn

Your data needs to be stored as NumPy arrays or as a list of NumPy arrays. Ideally, you split the data in training and test sets, for which you can also resort to the `train_test_split` module of `sklearn.cross_validation`.

#### Keras Data Sets

```
>>> from keras.datasets import boston_housing,
    mnist,
    cifar10,
    imdb
>>> (x_train,y_train),(x_test,y_test) = mnist.load_data()
>>> (x_train2,y_train2),(x_test2,y_test2) = boston_housing.load_data()
>>> (x_train3,y_train3),(x_test3,y_test3) = cifar10.load_data()
>>> (x_train4,y_train4),(x_test4,y_test4) = imdb.load_data(num_words=20000)
>>> num_classes = 10
```

#### Other

```
>>> from urllib.request import urlopen
>>> data = np.loadtxt(urlopen("http://archive.ics.uci.edu/ml/machine-learning-databases/pima-indians-diabetes/pima-indians-diabetes.data"),delimiter=",")
>>> X = data[:,0:8]
>>> y = data[:,8]
```

#### Preprocessing

##### Sequence Padding

```
>>> from keras.preprocessing import sequence
>>> x_train4 = sequence.pad_sequences(x_train4,maxlen=80)
>>> x_test4 = sequence.pad_sequences(x_test4,maxlen=80)
```

##### One-Hot Encoding

```
>>> from keras.utils import to_categorical
>>> y_train = to_categorical(y_train, num_classes)
>>> y_test = to_categorical(y_test, num_classes)
>>> y_train3 = to_categorical(y_train3, num_classes)
>>> y_test3 = to_categorical(y_test3, num_classes)
```

## Model Architecture

### Sequential Model

```
>>> from keras.models import Sequential
>>> model = Sequential()
>>> model2 = Sequential()
>>> model3 = Sequential()
```

### Multilayer Perceptron (MLP)

#### Binary Classification

```
>>> from keras.layers import Dense
>>> model.add(Dense(12,
    input_dim=8,
    kernel_initializer='uniform',
    activation='relu'))
>>> model.add(Dense(8,kernel_initializer='uniform',activation='relu'))
>>> model.add(Dense(1,kernel_initializer='uniform',activation='sigmoid'))
```

#### Multi-Class Classification

```
>>> from keras.layers import Dropout
>>> model.add(Dropout(0.2))
>>> model.add(Dense(512,activation='relu',input_shape=(784,)))
>>> model.add(Dense(512,activation='relu'))
>>> model.add(Dense(10,activation='softmax'))
```

#### Regression

```
>>> model.add(Dense(64,activation='relu',input_dim=train_data.shape[1]))
>>> model.add(Dense(1))
```

### Convolutional Neural Network (CNN)

```
>>> from keras.layers import Activation,Conv2D,MaxPooling2D,Flatten
>>> model2.add(Conv2D(32,(3,3),padding='same',input_shape=x_train.shape[1:]))
>>> model2.add(Activation('relu'))
>>> model2.add(Conv2D(32,(3,3)))
>>> model2.add(Activation('relu'))
>>> model2.add(MaxPooling2D(pool_size=(2,2)))
>>> model2.add(Dropout(0.25))
>>> model2.add(Conv2D(64,(3,3),padding='same'))
>>> model2.add(Activation('relu'))
>>> model2.add(Conv2D(64,(3,3)))
>>> model2.add(Activation('relu'))
>>> model2.add(MaxPooling2D(pool_size=(2,2)))
>>> model2.add(Dropout(0.25))
>>> model2.add(Flatten())
>>> model2.add(Dense(512))
>>> model2.add(Activation('relu'))
>>> model2.add(Dropout(0.5))
>>> model2.add(Dense(num_classes))
>>> model2.add(Activation('softmax'))
```

### Recurrent Neural Network (RNN)

```
>>> from keras.layers import Embedding,LSTM
>>> model3.add(Embedding(20000,128))
>>> model3.add(LSTM(128,dropout=0.2,recurrent_dropout=0.2))
>>> model3.add(Dense(1,activation='sigmoid'))
```

#### Also see NumPy & Scikit-Learn

### Train and Test Sets

```
>>> from sklearn.model_selection import train_test_split
>>> X_train5,X_test5,y_train5,y_test5 = train_test_split(X,
    y,
    test_size=0.33,
    random_state=42)
```

### Standardization/Normalization

```
>>> from sklearn.preprocessing import StandardScaler
>>> scaler = StandardScaler().fit(x_train)
>>> standardized_X = scaler.transform(x_train2)
>>> standardized_X_test = scaler.transform(x_test2)
```

## Inspect Model

```
>>> model.output_shape
>>> model.summary()
>>> model.get_config()
>>> model.get_weights()
```

Model output shape  
Model summary representation  
Model configuration  
List all weight tensors in the model

## Compile Model

MLP: Binary Classification  
>>> model.compile(optimizer='adam',  
loss='binary\_crossentropy',  
metrics=['accuracy'])  
MLP: Multi-Class Classification  
>>> model.compile(optimizer='rmsprop',  
loss='categorical\_crossentropy',  
metrics=['accuracy'])  
MLP: Regression  
>>> model.compile(optimizer='rmsprop',  
loss='mse',  
metrics=['mae'])  
Recurrent Neural Network  
>>> model3.compile(loss='binary\_crossentropy',  
optimizer='adam',  
metrics=['accuracy'])

## Model Training

```
>>> model3.fit(x_train4,
    y_train4,
    batch_size=32,
    epochs=15,
    verbose=1,
    validation_data=(x_test4,y_test4))
```

## Evaluate Your Model's Performance

```
>>> score = model3.evaluate(x_test,
    y_test,
    batch_size=32)
```

## Prediction

```
>>> model3.predict(x_test4, batch_size=32)
>>> model3.predict_classes(x_test4, batch_size=32)
```

## Save / Reload Models

```
>>> from keras.models import load_model
>>> model3.save('model_file.h5')
>>> my_model = load_model('my_model.h5')
```

## Model Fine-tuning

### Optimization Parameters

```
>>> from keras.optimizers import RMSprop
>>> opt = RMSprop(lr=0.0001, decay=1e-6)
>>> model2.compile(loss='categorical_crossentropy',
    optimizer=opt,
    metrics=['accuracy'])
```

### Early Stopping

```
>>> from keras.callbacks import EarlyStopping
>>> early_stopping_monitor = EarlyStopping(patience=2)
>>> model3.fit(x_train4,
    y_train4,
    batch_size=32,
    epochs=15,
    validation_data=(x_test4,y_test4),
    callbacks=[early_stopping_monitor])
```

# Convert Networks from TF → MXNet/etc.

---

- **MMdnn**
- <https://github.com/Microsoft/MMdnn>

- 
- Lab: Sequential API versus Functional/Model API

Regression: Predict the price of houses in Boston

# Programming Deep Neural Networks in Keras or TensorFlow

The screenshot shows the Keras Documentation website. The header features a large red 'K' logo and the text 'Keras Documentation'. Below the header is a search bar labeled 'Search docs'. The main navigation menu includes 'Home', 'Getting started', 'Guide to the Sequential model' (which is highlighted in dark grey), and 'Guide to the Functional API'. Under 'Guide to the Functional API', there is a list of topics: 'Getting started with the Keras functional API', 'First example: a densely-connected network', 'All models are callable, just like layers', 'Multi-input and multi-output models', 'Shared layers', 'The concept of layer "node"', and 'More examples'.

- **Sequential programming**
  - The Sequential model is a linear stack of layers.
- **Functional programming**
  - The Keras functional API is the way to go for defining complex models, such as multi-output models, directed acyclic graphs, or models with shared layers.

<https://keras.io/getting-started/sequential-model-guide/>

# Sequential definition of a model

---

## 4 Train a Neural Network Model (8-12-8-1)

- Train a 8-12-8-1 network
- Extract probabilities

```
: 1 ## Train a Neural Network Model (8-12-8-1)
: 2 # Sample Multilayer Perceptron Neural Network in Keras
: 3 from keras.models import Sequential
: 4 from keras.layers import Dense
: 5 import numpy
: 6
: 7 X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.30, random_state=42)
: 8
: 9 # 1. define the network
:10 model = Sequential()
:11 model.add(Dense(12, input_dim=8, init='uniform', activation='relu'))
:12 model.add(Dense(8, init='uniform', activation='relu'))
:13 model.add(Dense(1, init='uniform', activation='sigmoid'))
:14 # 2. compile the network
:15 #time it
:16 %time model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
```

# Keras functional API

## First example: a densely-connected network

The `Sequential` model is probably a better choice to implement such a network, but it helps to start with something really simple.

**784 – 64 – 64 – 10**

- A layer instance is callable (on a tensor), and it returns a tensor

**x = Dense(64, activation =“relu”)(inputs)**  
1. Creates layer `Dense(64, activation =“relu”)` creates a layer  
2. Calls the callable method on this Dense layer `Dense(64, activation =“relu”)(inputs)`  
Connect the layer inputs to to the Dense 64 layer and returns Dense(64)

```
from keras.models import Model

# This returns a tensor
inputs = Input(shape=(784,))

# a layer instance is callable on a tensor, and returns a tensor
x = Dense(64, activation='relu')(inputs)
x = Dense(64, activation='relu')(x)
predictions = Dense(10, activation='softmax')(x)

# This creates a model that includes
# the Input layer and three Dense layers
model = Model(inputs=inputs, outputs=predictions)
model.compile(optimizer='rmsprop',
              loss='categorical_crossentropy',
              metrics=['accuracy'])
model.fit(data, labels) # starts training
```

# Reuse model architecture and weights

---

- All models are callable, just like layers
- With the functional API, it is easy to re-use trained models: you can treat any model as if it were a layer, by calling it on a tensor.
- Note that by calling a model you aren't just re-using the **architecture** of the model, you are also re-using its **weights**.

```
x = Input(shape=(784,))
# This works, and returns the 10-way softmax we defined above.
y = model(x)
```

# Callable object

---

- A **callable object**, in [computer programming](#), is any object that can be called like a [function](#).

In Python any object with a `__call__()` method can be called using function-call syntax.

```
# Python 3 Syntax
class Foo(object):
    def __call__(self):
        print("Called.")

foo_instance = Foo()
foo_instance() # This will output "Called." to the screen.
```

[2]

Another example:

```
class Accumulator(object):
    def __init__(self, n):
        self.n = n

    def __call__(self, x):
        self.n += x
        return self.n
```

# Callable object

---

- A **callable object**, in [computer programming](#), is any object that can be called like a [function](#).

`__call__` makes any object be callable as a function.

This example will output 8:

```
class Adder(object):
    def __init__(self, val):
        self.val = val

    def __call__(self, val):
        return self.val + val

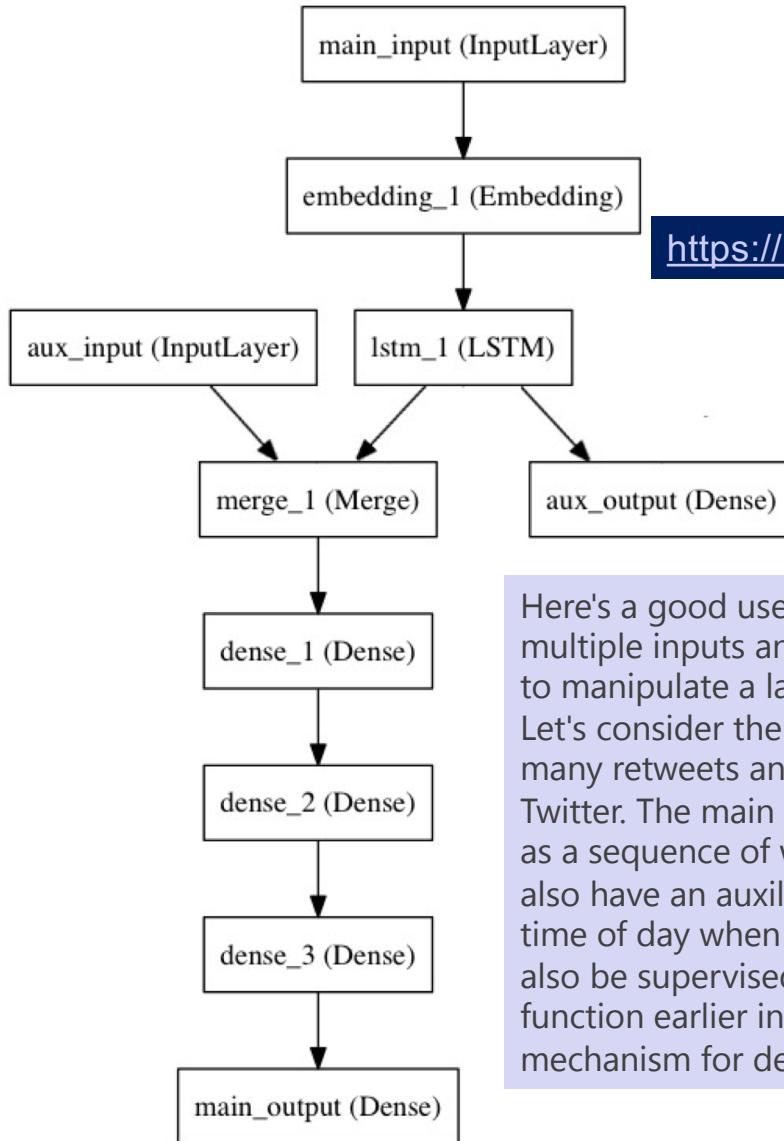
func = Adder(5)
print func(3)
```

# Dense Layer

---

- <https://github.com/keras-team/keras/blob/master/keras/layers/core.py>

```
878     def call(self, inputs):
879         output = K.dot(inputs, self.kernel)
880         if self.use_bias:
881             output = K.bias_add(output, self.bias, data_format='channels_last')
882         if self.activation is not None:
883             output = self.activation(output)
884         return output
885
```



## Multi-input and multi-output models

<https://keras.io/getting-started/functional-api-guide/>

Here's a good use case for the functional API: models with multiple inputs and outputs. The functional API makes it easy to manipulate a large number of intertwined datastreams. Let's consider the following model. We seek to predict how many retweets and likes a news headline will receive on Twitter. The main input to the model will be the headline itself, as a sequence of words, but to spice things up, our model will also have an auxiliary input, receiving extra data such as the time of day when the headline was posted, etc. The model will also be supervised via two loss functions. Using the main loss function earlier in a model is a good regularization mechanism for deep models.

# Model Class API

The screenshot shows a documentation page for the Keras Model Class API. The left sidebar has a red header "Documentation" and a dark grey section labeled "model API". The main content area has a breadcrumb navigation "Docs » Models » Model (functional API)" and a "Edit on GitHub" button. The title "Model class API" is followed by a paragraph explaining how to instantiate a `Model` via functional API. A code snippet shows creating an input tensor `a`, a dense layer `b`, and a `Model` with `inputs=a` and `outputs=b`. Below it, a note says the model will include all layers required in the computation of `b` given `a`. Another section for multi-input or multi-output models is mentioned. A "Keras Linear Regression" section follows, with code for a linear model and notes about its implementation.

Model class API

In the functional API, given some input tensor(s) and output tensor(s), you can instantiate a `Model` via:

```
from keras.models import Model
from keras.layers import Input, Dense

a = Input(shape=(32,))
b = Dense(32)(a)
model = Model(inputs=a, outputs=b)
```

This model will include all layers required in the computation of `b` given `a`.

In the case of multi-input or multi-output models, you can use

```
model = Model(inputs=[a1, a2], outputs=[b1, b3, b3])
```

For a detailed introduction of what `Model` can do, read [this guide](#).

## Useful attributes of Model

- `model.layers` is a flattened list of the layers comprising the model.
- `model.inputs` is the list of input tensors.
- `model.outputs` is the list of output tensors.

## Keras Linear Regression

Linear Regression model can be seen as a neural network without hidden layer and any activation function

```
from keras.layers import Dense, Input
from keras.models import Model
```

Using TensorFlow backend.

Number of features in Boston dataset is 13. Thus the input layer has 13 units. We want to predict only one value - price of the house. Therefore output layer has only one unit.

```
: inp = Input(shape=(13,))
out = Dense(1)(inp)
```

Usual Linear Regression is done by means of solving Ordinary Least Square Error problem which means that one should use squared (mse) loss.

Here we also choose the simplest optimizer - usual stochastic gradient descent.

```
: model_keras = Model(inputs=[inp], outputs=[out])
model_keras.compile(loss="mse", optimizer="sgd", metrics=["mae"])
```

Split into train and test set (with the same random\_state which means we can compare results)

```
: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

Scaling

```
: scaler = MinMaxScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

# **\_\_call\_\_ method for an object: default execution is to execution**

---

- Used in SciKit Learn
- In Keras to connect layers into a network
- myInstance(var1, var2)
  - Will result in the method call being executed

```
...     __call__  
839     def call(self, inputs):  
840         output = K.dot(inputs, self.kernel)  
841         if self.use_bias:  
842             output = K.bias_add(output, self.bias)  
843         if self.activation is not None:  
844             output = self.activation(output)  
845         return output  
846
```

## Keras Linear Regression

Linear Regression model can be seen as a neural network without hidden layer and any activation function

```
In [31]: from keras.layers import Dense, Input  
from keras.models import Model  
  
Using TensorFlow backend.
```

Number of features in Boston dataset is 13. Thus the input layer has 13 units. We want to predict only one value - price of the house. Therefore output layer has only one unit.

```
In [32]: inp = Input(shape=(13,))  
out = Dense(1)(inp)
```

inp is callable

Usual Linear Regression is done by means of solving Ordinary Least Square Error problem which means that one should use squared (mse) loss.

Here we also choose the simplest optimizer - usual stochastic gradient descent.

```
In [33]: model_keras = Model(inputs=[inp], outputs=[out])  
model_keras.compile(loss="mse", optimizer="sgd", metrics=["mae"])
```

Split into train and test set (with the same random\_state which means we can compare results)

```
In [34]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

Scaling

```
In [35]: scaler = MinMaxScaler()  
X_train = scaler.fit_transform(X_train)  
X_test = scaler.transform(X_test)
```

Scale data

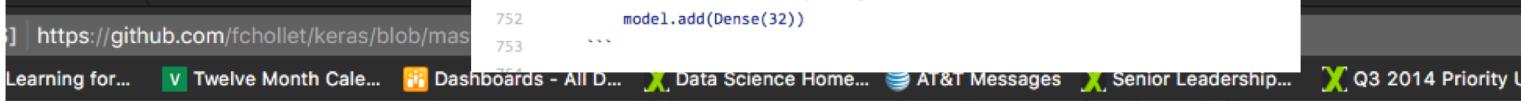
Fitting

```
In [36]: hist = model_keras.fit(X_train, y_train,  
                           validation_data=(X_test, y_test),  
                           epochs=300, verbose=0)
```

```
/opt/conda/lib/python3.6/site-packages/keras/backend/tensorflow_backend.py:2252: UserWarning: Expected no kwargs, you passed 1  
kwargs passed to function are ignored with Tensorflow backend  
warnings.warn('\n'.join(msg))
```

Sequential assumes a single input and a single output and single path thru graph. However, using Keras Model/functional API, we can build sophisticated graphs with multiple inputs and outputs

```
    ...
735     created by the layer, and `bias` is a bias vector created by the layer
736     (only applicable if `use_bias` is `True`).
737
738     Note: if the input to the layer has a rank greater than 2, then
739     it is flattened prior to the initial dot product with `kernel`.
740
741     # Example
742
743     ```python
744         # as first layer in a sequential model:
745         model = Sequential()
746         model.add(Dense(32, input_shape=(16,)))
747         # now the model will take as input arrays of shape (*, 16)
748         # and output arrays of shape (*, 32)
749
750         # after the first layer, you don't need to specify
751         # the size of the input anymore:
752         model.add(Dense(32))
753     ...
```



```
836     self.input_spec = InputSpec(min_ndim=2, axes=(-1: input_dim))
837     self.built = True
838
839     def call(self, inputs):
840         output = K.dot(inputs, self.kernel)
841         if self.use_bias:
842             output = K.bias_add(output, self.bias)
843         if self.activation is not None:
844             output = self.activation(output)
845         return output
846
847     def compute_output_shape(self, input_shape):
848         assert input_shape and len(input_shape) >= 2
849         assert input_shape[-1]
850         output_shape = list(input_shape)
851         output_shape[-1] = self.units
852         return tuple(output_shape)
853
854     def get_config(self):
```

- 
- Model functional API versus sequential
  - <https://www.quora.com/Why-would-one-use-the-special-method-call--instead-of-defining-a-new-method-for-an-object-in-Python>

---

# **Lab: Iris: Three class classification problem**

# Quizzes and HW Assignments

 Quiz 4.1: Pima Number of parameter

1 pt | 1 Question

 Quiz 4.3: Side of hyperplane

1 Question

 Quiz 4.4: Label probability

1 Question

 Quiz 4.5: Perceptron Loss

1 Question

 Quiz 4.6: Output activation function

1 Question

 Quiz 4.7: BackProp for MLP

1 Question

 Quiz 4.8: Keras MLP for Boston

2 pts | 2 Questions

 Quiz 4.9: Keras MLP for Iris

2 pts | 2 Questions

 Assignment 1: Linear Regression

40 pts

 Assignment 2: KNN and Logistic Regression

30 pts

 Assignment 3: Logistic Regression

20 pts

 Assignment 4: BackProp

30 pts

# Quiz 4.9

---

## Quiz Instructions

Run the [MLPIris](#) and do the task in the end of the notebook.



### Question 1

What is the test Accuracy for MLP with the described architecture?

Report an answer with 2 decimal places (using `np.round()`).



### Question 2

How many parameters does the described model have?

## Quiz: Bonus exercise

---

- Implement vanilla logistic regression in Keras (use the code from you 4-20-3 network as a basis)

# Vanilla logistic regression

---

## 4-3 for iris data

have 3 output units.

```
] : inp = Input(shape=(4,))
out = Dense(3, activation="softmax")(inp)
```

Multinomial Logistic Regression uses softmax activation in multiclass case instead of binary logit function. Also we use categorical\_crossentropy as a loss function.

```
] : model = Model(inputs=[inp], outputs=[out])
model.compile(loss="categorical_crossentropy", optimizer="sgd", metrics=["accuracy"])
```

Split into train and test set

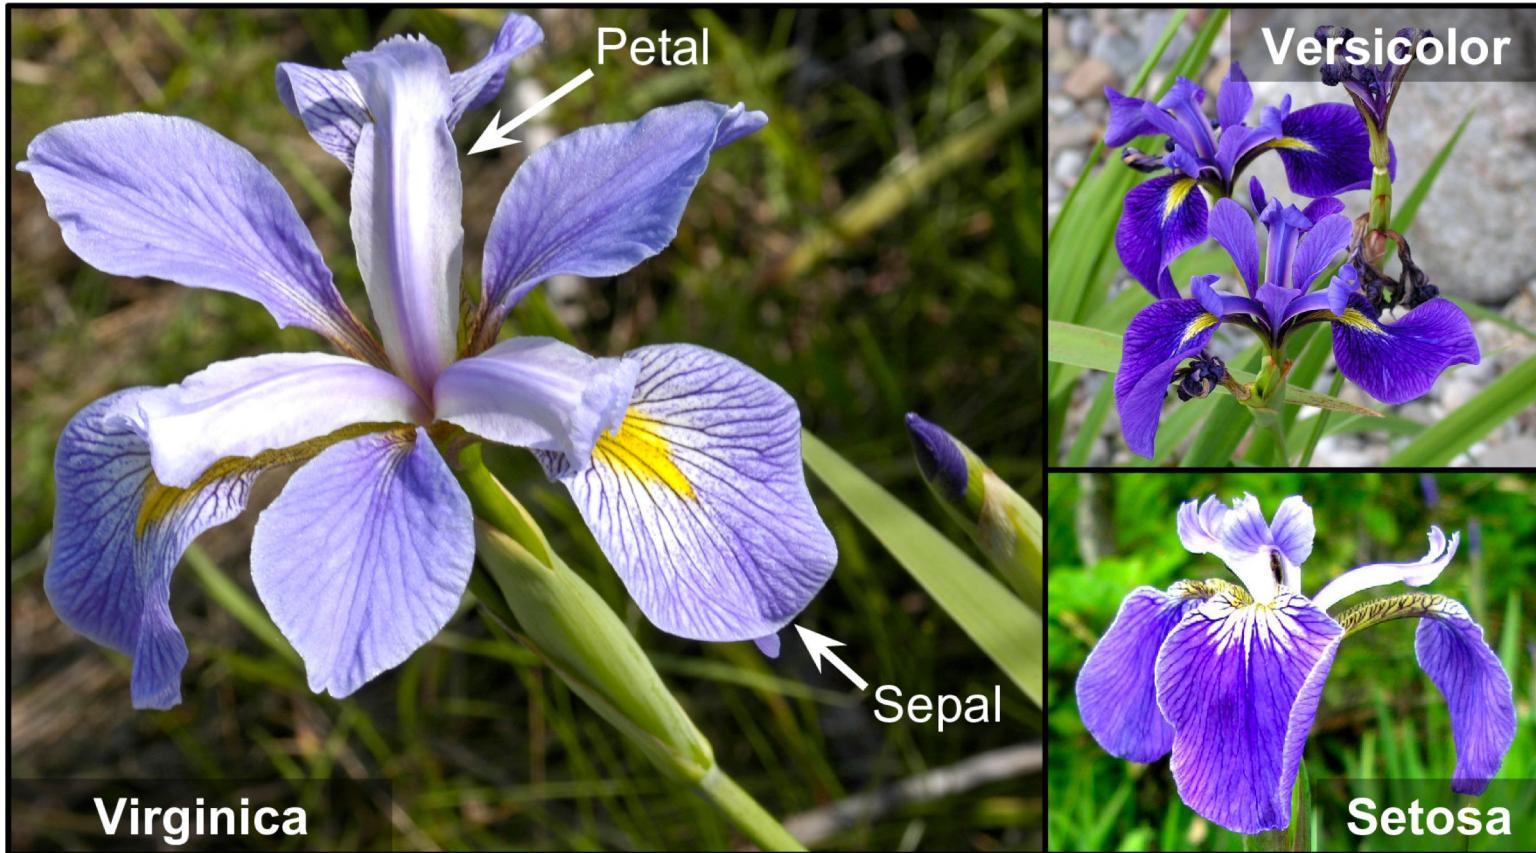
# Labs

---

- **Here are the links to the notebooks for in-class work:**
  - [Pima](#)
  - [Boston](#)
  - [Iris](#)

**A sepal is a part of the flower of angiosperms (flowering plants).  
Usually green, sepals typically function as protection for the flower in bud,  
and often as support for the petals when in bloom**

---



# Quiz: IrisTask

---

## Keras MLP (TASK)

Iris dataset has 4 features per sample. Thus we have 4 units in the input layer. Our problem has 3 classes and therefore we should have 3 output units.

Now let's add one more hidden layer with 20 unit between input and output and tanh activation function.

The syntax is the following (Keras API):

<name> = <Layer>(<args>)(<name of layer to connect to>)

```
In [ ]: inp = Input(shape=(4,))
# add Dense layer called "hidden" with 20 neurons and "relu" activation and connect it to the "in
p" layer
out = Dense(3, activation="softmax")(hidden)
```

MLP

```
In [ ]: model = Model(inputs=[inp], outputs=[out])
model.compile(loss="categorical_crossentropy", optimizer="adam", metrics=["accuracy"])
```

Model description

```
In [ ]: model.summary()
```

# Notebooks

- **Pima diabetes: binary classification problem**

- GitHub: /Labs/Unit-16
  - Src/Dev/Labs/Unit-16

- **QUIZ: Iris Notebook Colab**

- Iris: Three class classification problem
  - <https://drive.google.com/file/d/1AcVGV64AU0-gFM8qW0t0Zy5JspTQi1vk/view?usp=sharing>

- ▼ **Labs-16-Deep-Learning-CrashCourse**
  - [Unit-04-MLP-BostonHousePrices](#)
  - [Unit-04-MLP-IrisFlowers](#)
  - [Unit-04-MLP-Keras-Hyperparameter-Tuning](#)
  - [Unit-04-MLP-PimaDiabets](#)



---

# **Lab: Regression: Predict the price of houses in Boston**

## Regression: Predict the price of houses in Boston

---

- The problem that we will look at in this tutorial is the [Boston house price dataset](#). *Shape [506, 14]*
- The dataset describes 13 numerical properties of houses in Boston suburbs and is concerned with modeling the price of houses in those suburbs in thousands of dollars.
- Input attributes include things like crime rate, proportion of nonretail business acres, chemical concentrations and more.



# Quiz 4.8 Boston MLP regression

---

## Quiz Instructions

Run the `MLPBoston` and do the task in the end of the notebook.



### Question 1

What is the test Accuracy for MLP with the described architecture?

Report an answer with 2 decimal places (using `np.round()`).



### Question 2

How many parameters does the described model have?

# MLP Regression Model 13-13Relu-1

---

- Boston dataset

```
1 # define base model
2 def baseline_model():
3     # create model
4     model = Sequential()
5     model.add(Dense(13, input_dim=13, kernel_initializer='normal', activation='relu'))
6     model.add(Dense(1, kernel_initializer='normal'))
7     # Compile model
8     model.compile(loss='mean_squared_error', optimizer='adam')
9     return model
```

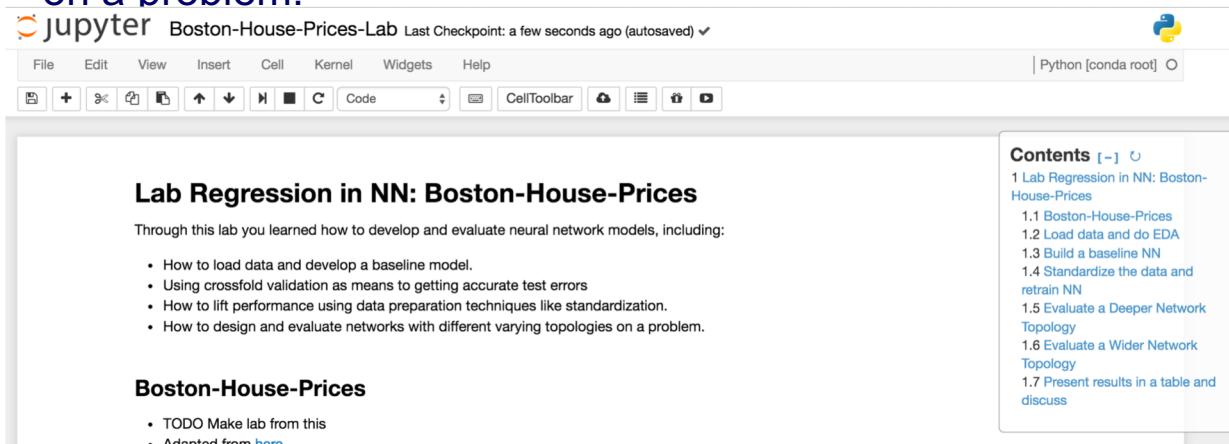
# MSE for Boston Housing dataset

---

- Reasonable performance for models evaluated using Mean Squared Error (MSE) are around 20 in squared thousands of dollars
- Or a mean absolute error of \$4,500 if you take the square root of MSE (not totally accurate).
- This is a nice target to aim for with our neural network model.

# In class Lab: NN for regression

- Through this lab you learned how to develop and evaluate neural network models, including:
  - How to load data and develop a baseline model.
  - Using crossfold validation as means to getting accurate test errors
  - How to lift performance using data preparation techniques like standardization.
  - How to design and evaluate networks with different varying topologies on a problem.



# Quiz: Boston Task

---

## Keras MLP (TASK)

Number of features in Boston dataset is 13. Thus the input layer has 13 units. We want to predict only one value - price of the house. Therefore output layer has only one unit.

Now let's add one more hidden layer with 20 unit between input and output and tanh activation function.

The syntax is the following (Keras API):

```
$\text{
```

```
: inp = Input(shape=(13,))
# add Dense layer called "hidden" with 20 neurons and "relu" activation and connect it to the "inp" layer
out = Dense(1, activation="linear")(hidden)
```

MLP

```
: model_mlp = Model(inputs=[inp], outputs=[out])
model_mlp.compile(loss="mse", optimizer="sgd", metrics=[ "mae" ])
```

Model description

---

# Quiz: Boston housing in linear regression

---

- Add home linear regression model in here
- [https://github.com/cadgip/DL\\_course/blob/master/Basics/Regression/LinRegrBoston.ipynb](https://github.com/cadgip/DL_course/blob/master/Basics/Regression/LinRegrBoston.ipynb)

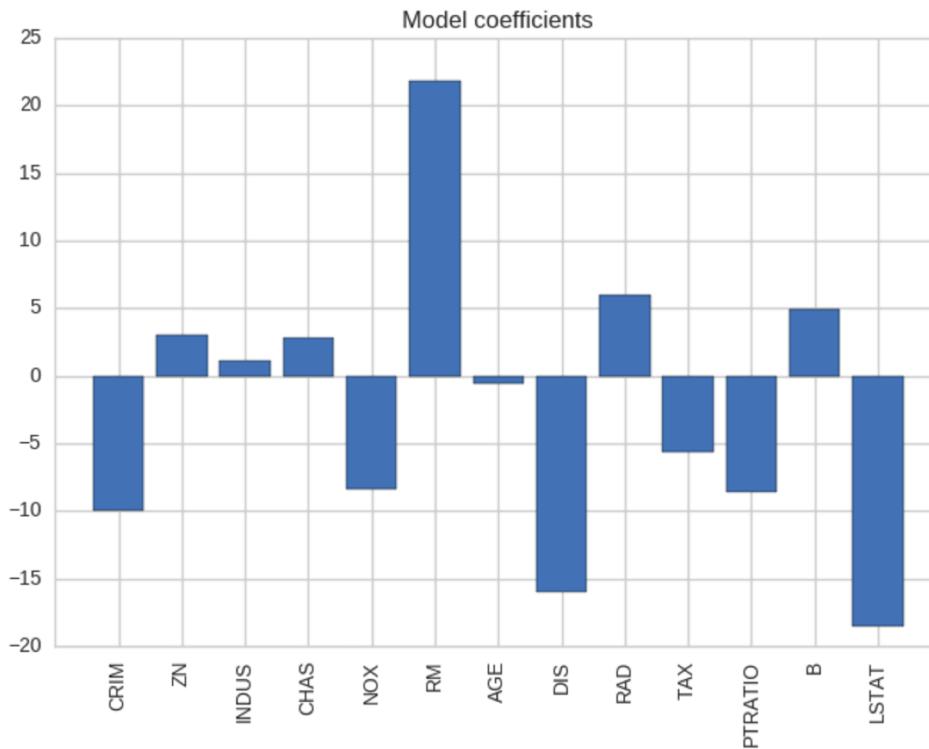
```
In [15]: model.fit(X_train, y_train)  
Out[15]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1, normalize=False)
```

---

## Evaluation

Let's see what features are significant for the model

```
In [16]: plt.bar(np.arange(model.coef_.shape[0]) - 0.4, model.coef_)  
plt.xticks(np.arange(model.coef_.shape[0]), X.columns, rotation='vertical')  
plt.xlim([-1, model.coef_.shape[0]])  
plt.title("Model coefficients")  
plt.show()
```



# KerasClassifier class

---

- Keras in pipeline fashion
- Uses the KerasClassifier class
  - Iris data example
    - <https://machinelearningmastery.com/iris-dataset-classification-tutorial-python/>

## Wrappers for the Scikit-Learn API

You can use `Sequential` Keras models (single-input only) as part of your Scikit-Learn workflow via the wrappers found at `keras.wrappers.scikit_learn.py`.

There are two wrappers available:

`keras.wrappers.scikit_learn.KerasClassifier(build_fn=None, **sk_params)`, which implements the Scikit-Learn classifier interface,

`keras.wrappers.scikit_learn.KerasRegressor(build_fn=None, **sk_params)`, which implements the Scikit-Learn regressor interface.

### Arguments

- `build_fn`: callable function or class instance
- `sk_params`: model parameters & fitting parameters

`build_fn` should construct, compile and return a Keras model, which will then be used to fit/predict. One of the following three values could be passed to `build_fn`:

---

The network topology of this simple one-layer neural network can be summarized as:

```
1 4 inputs -> [8 hidden nodes] -> 3 outputs
```

Note that we use a “softmax” activation function in the output layer. This is to ensure the output values are in the range of 0 and 1 and may be used as predicted probabilities.

Finally, the network uses the efficient Adam gradient descent optimization algorithm with a logarithmic loss function, which is called “*categorical\_crossentropy*” in Keras.

```
1 # define baseline model
2 def baseline_model():
3     # create model
4     model = Sequential()
5     model.add(Dense(8, input_dim=4, activation='relu'))
6     model.add(Dense(3, activation='softmax'))
7     # Compile model
8     model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accu
9     return model
```

We can now create our KerasClassifier for use in scikit-learn.

We can also pass arguments in the construction of the KerasClassifier class that will be passed on to the fit() function internally used to train the neural network. Here, we pass the number of epochs as 200 and batch size as 5 to use when training the model. Debugging is also turned off when training by setting verbose to 0.

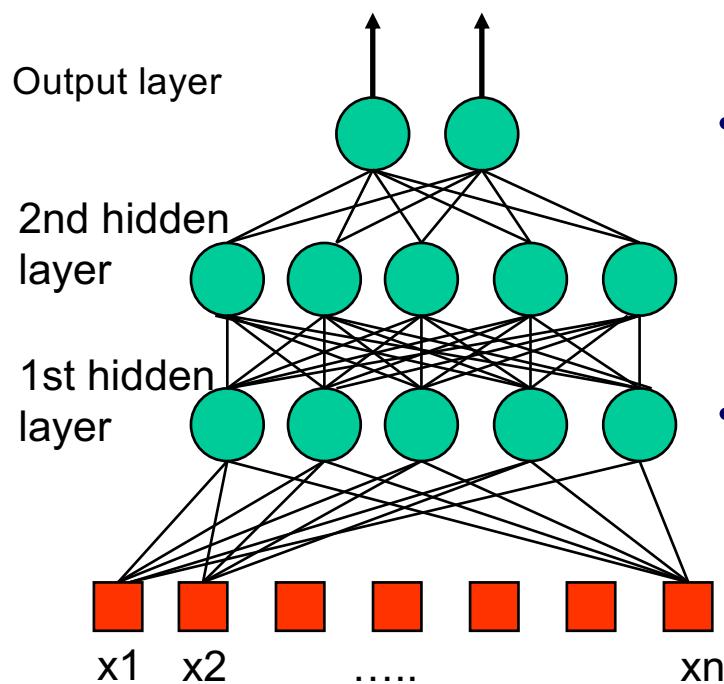
```
1 estimator = KerasClassifier(build_fn=baseline_model, epochs=200, batch_size=5, verbose=0)
```



# Outline

- 1. Introduction**
- 2. ML and deep learning review**
  1. Perceptron → MLP
  2. Keras
  - 3. BackPropagation**
- 3. What is computer vision?**
  1. Computer vision
  2. Convolutional Neural Nets (CNNs)
- 4. Deep NN Architectures for CV Tasks**
  1. Backbone networks
- 5. Solving edge-based IoT**
- 6. Conclusions and Next steps**

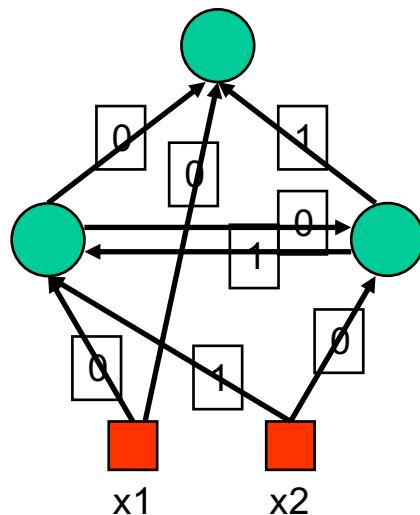
# Feed Forward Neural Networks



- The information is propagated from the inputs to the outputs
- Computations of **No non-linear functions from  $n$  input variables by compositions of  $N_c$  algebraic functions**
- Time has no role (NO cycle between outputs and inputs)

# Recurrent Neural Networks

- Can have arbitrary topologies
- Can model systems with internal states (dynamic ones)
- Delays are associated to a specific weight
- Training is more difficult
- Performance may be problematic
  - Stable Outputs may be more difficult to evaluate
  - Unexpected behavior (oscillation, chaos, ...)



# MLP Regression Model: 13-13Relu-1

---

- Boston dataset

```
1 # define base model
2 def baseline_model():
3     # create model
4     model = Sequential()
5     model.add(Dense(13, input_dim=13, kernel_initializer='normal', activation='relu'))
6     model.add(Dense(1, kernel_initializer='normal'))
7     # Compile model
8     model.compile(loss='mean_squared_error', optimizer='adam')
9     return model
```

$$\text{MSE}(X; \text{MLP}_{\text{Regression}}) = \frac{1}{2n} (I(W_2 (\sigma(W_1 X))) - t_k)^2 \text{ where } I \text{ is identity}$$

Input. W1 Hidden W2 output

# MLP Softmax Model

$$J(\Theta) = -\frac{1}{m} \sum_{i=1}^m \sum_{k=1}^K y_k^{(i)} \log (\hat{p}_k^{(i)})$$

$$\nabla_{\theta_k} J(\Theta) = \frac{1}{m} \sum_{i=1}^m (\hat{p}_k^{(i)} - y_k^{(i)}) \mathbf{x}^{(i)}$$

MSE(X; MLP<sub>Regression</sub>) =  $\frac{1}{2n}(I(W_2(\sigma(W_1 X))) - t_k)^2$  where I is identity  
CrossEntropy(X; MLP<sub>Regression</sub>) =  $\frac{1}{2n}(\sigma(W_2(\sigma(W_1 X)))_k - t_k)^2$ .  
where  $\sigma(\cdot)_k$  is defined as follows:

$$\hat{p}_k = \sigma(\mathbf{s}(\mathbf{x}))_k = \frac{\exp(s_k(\mathbf{x}))}{\sum_{j=1}^K \exp(s_j(\mathbf{x}))}$$

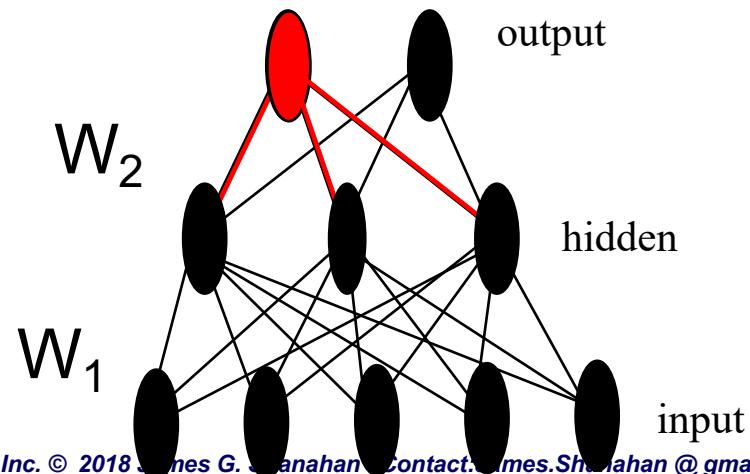
- $K$  is the number of classes.
- $\mathbf{s}(\mathbf{x})$  is a vector containing the scores of each class for the instance  $\mathbf{x}$ .
- $\sigma(\mathbf{s}(\mathbf{x}))_k$  is the estimated probability that the instance  $\mathbf{x}$  belongs to class  $k$  given the scores of each class for that instance.

# General (two-layer) feedforward network

---

$$\text{MLP}(X) = g( f(XW_1) W_2 )$$

- The hidden units with their activation functions can express non-linear functions
- The activation functions can be different at neurons (but the same one is used in practice)



# Backpropagation Algorithm

- Initialize weights (typically random!)
- Keep doing epochs
  - For each example  $e$  in training set do
    - forward pass to compute
      - $O = \text{neural-net-output}(\text{network}, e)$
      - miss =  $(T - O)$  at each output unit
    - backward pass to calculate deltas to weights
    - update all weights
  - end
- until tuning set error stops improving

Forward pass explained earlier

Backward pass explained in next slide

# Backpropagation

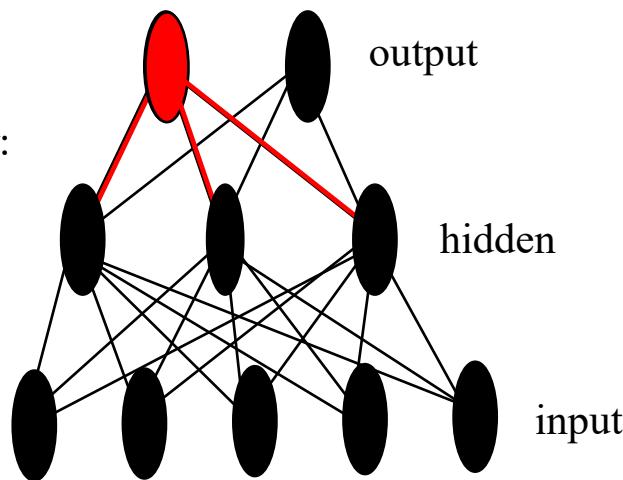
---

**Calculate the error signal for the output neurons and update the weights between the output and hidden layers**

$$\delta_k = (t_k - z_k) f'(net_k)$$

update the weights to  $k$ :

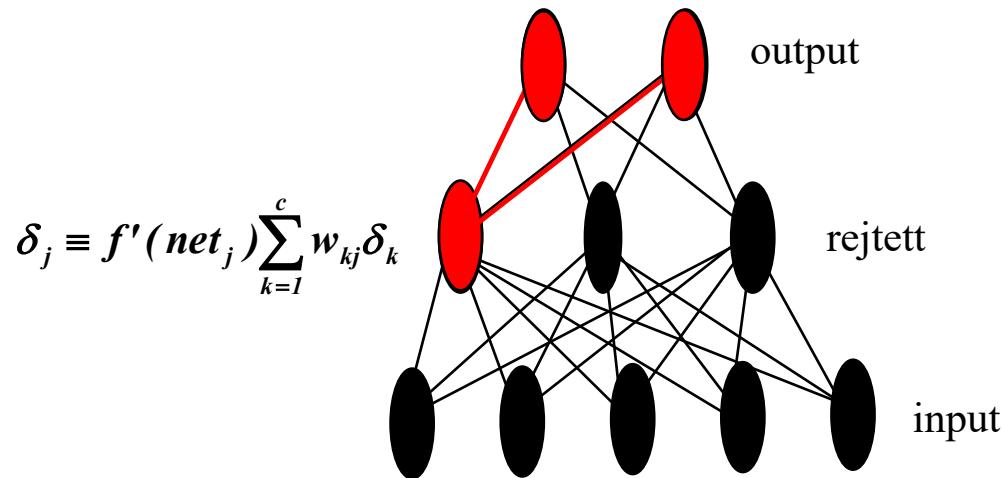
$$\Delta w_{ki} = \eta \delta_k z_i$$



# Backpropagation

---

Calculate the error signal for hidden neurons



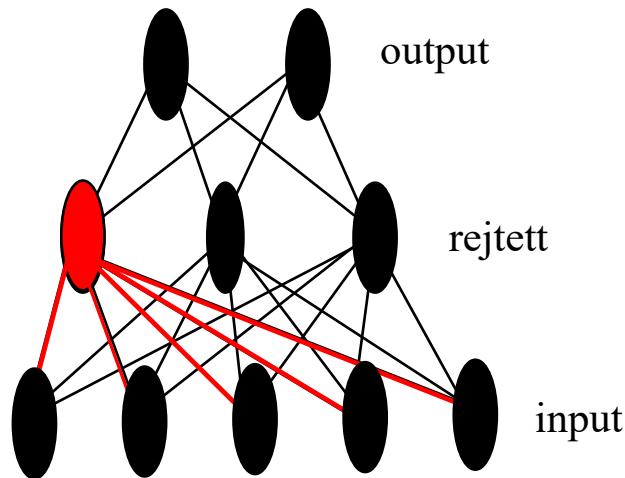
# Backpropagation

---

Update the weights between the input and hidden neurons

updating the ones to  $j$

$$\Delta w_{ji} = \eta \delta_j x_i$$



---

# BackProp Algo.

# Backpropagation algorithm: A History

---

- The backpropagation algorithm was originally introduced in the 1970s, but its importance wasn't fully appreciated until a [famous 1986 paper](#) by [David Rumelhart](#), [Geoffrey Hinton](#), and [Ronald Williams](#).
  - That paper describes several neural networks where backpropagation works far faster than earlier approaches to learning, making it possible to use neural nets to solve problems which had previously been insoluble.
- Today, the backpropagation algorithm is the workhorse of learning in neural networks.

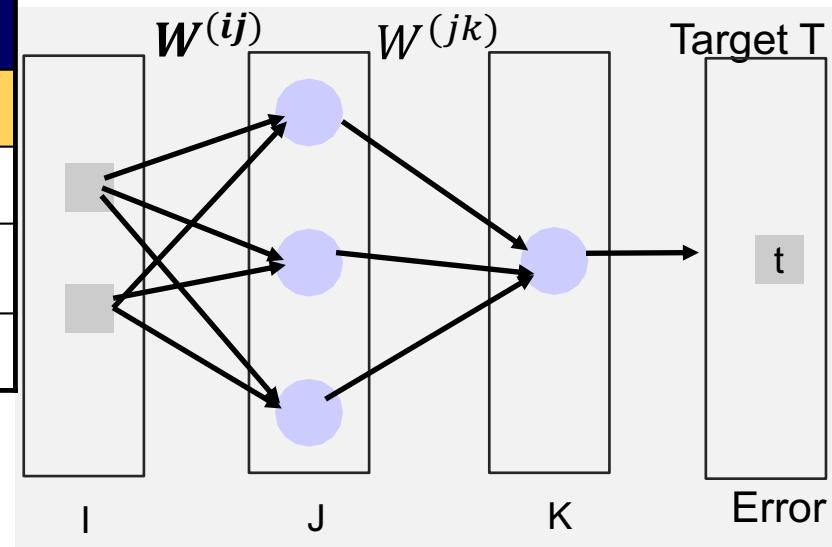
# Why study backprop?

---

- At the heart of backpropagation is an expression for the partial derivative  $\partial C/\partial w$  of the cost function  $C$  with respect to any weight  $w$  (or bias  $b$ ) in the network.
- The expression tells us how quickly the cost changes when we change the weights and biases.
  - And while the expression is somewhat complex, it also has a beauty to it, with each element having a natural, intuitive interpretation.
- And so backpropagation isn't just a fast algorithm for learning. It actually gives us detailed insights into how changing the weights and biases changes the overall behaviour of the network.
- That's well worth studying in detail.

# Regression via MLP

<i>Instance\Attr</i>	$x_1$	$x_2$	...	$x_n$	$t$
1	3	0	..	7	-1
2					-4
...	...	...	...	...	...
$m$	0	4	...	8	23



$$MLP(x) = g(f(XW^{(ij)})W^{(jk)})$$

$$MLP\_regression(X) = I(\sigma(XW^{(ij)})W^{(jk)})$$

$$MSE = \frac{1}{2m} \sum_{i=1:m} (O_k - t_k)^2$$

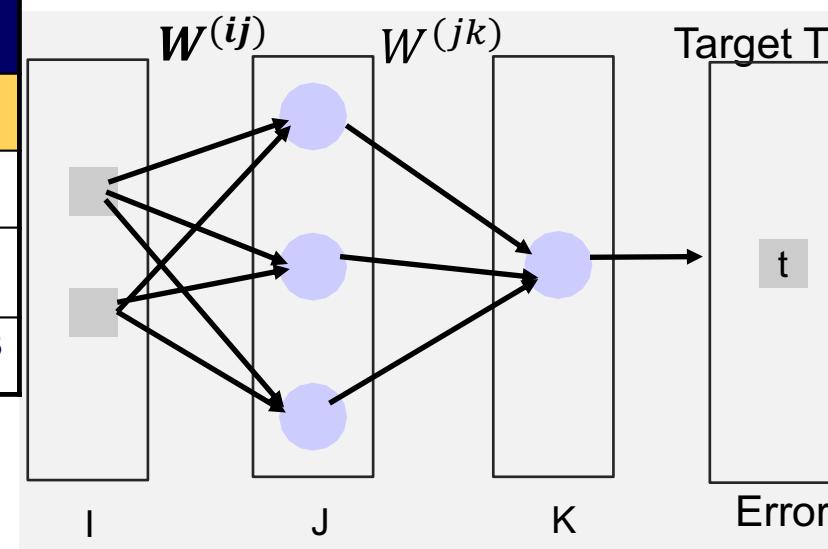
$$MSE = \frac{1}{2m} \sum_{i=1:m} (MLP\_regression(X_i) - t_i)^2$$

#MLP Regression

#old fashioned sigmoid

# Learn an MLP for Regression via BackProp (Gradient descent and chain rule of calculus)

Instance\Attr	$x_1$	$x_2$	...	$x_n$	$t$
1	3	0	..	7	-1
2					-4
...	...	...	...	...	...
$m$	0	4	...	8	23



$$MLP(x) = g(f(XW^{(ij)})W^{(jk)})$$

#MLP Regression

$$MLP\_regression(X) = \sigma(\sigma(XW^{(ij)})W^{(jk)})$$

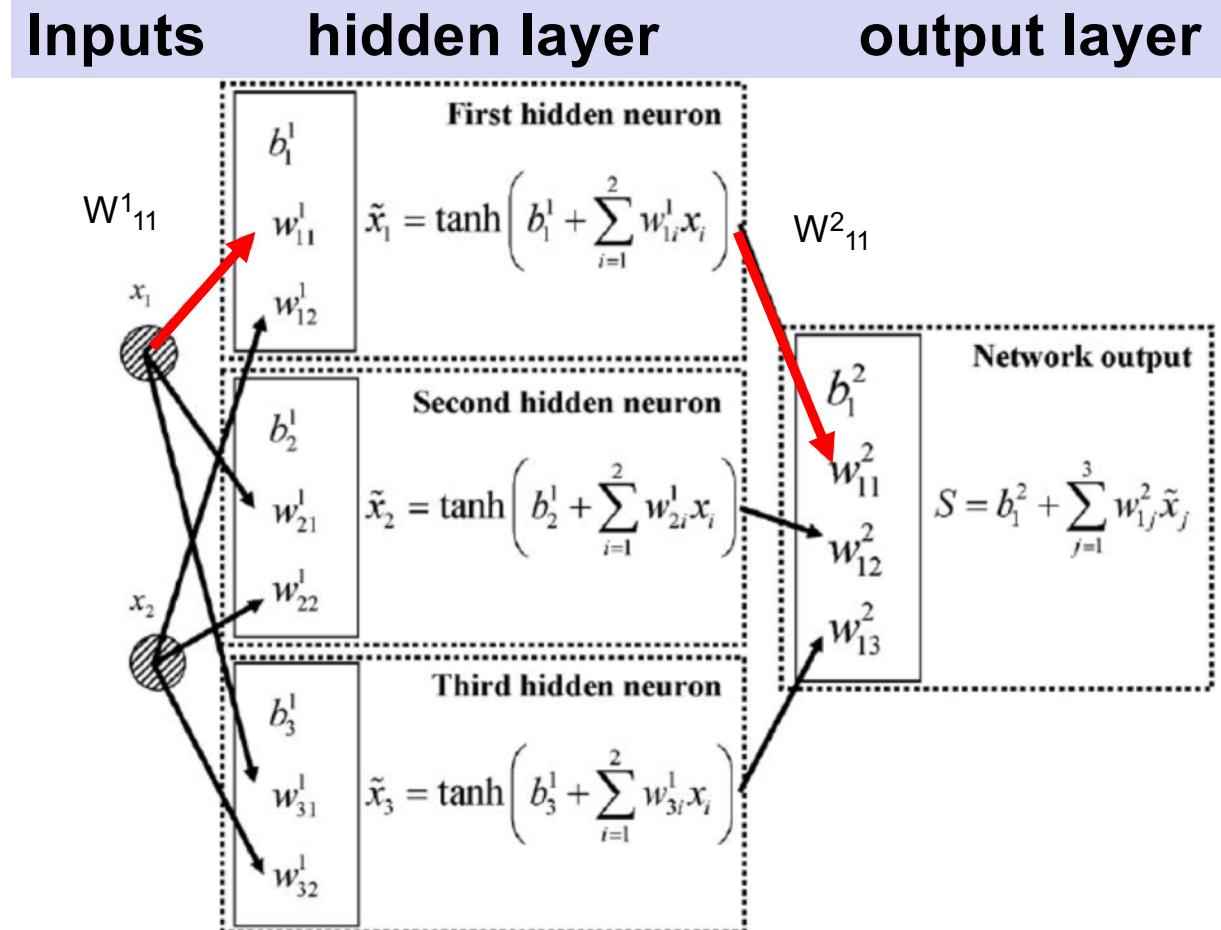
#old fashioned sigmoid

$$MSE = \frac{1}{2n} \sum_{i=1:m} (O_k - t_k)^2$$

Learn via backpropagation  
algorithm (Gradient descent and  
chain rule of calculus)

$$MSE = \frac{1}{2n} \sum_{i=1:m} (MLP\_regression(X_i) - t_i)^2$$

# Activation functions in NN: 2-3-1



# Gradient Descent for NN

---

- Given a set of training data points with target  $t_k$  and output layer output  $O_k$  we can write the error as

$$E = \frac{1}{2} \sum_{k \in K} (O_k - t_k)^2$$

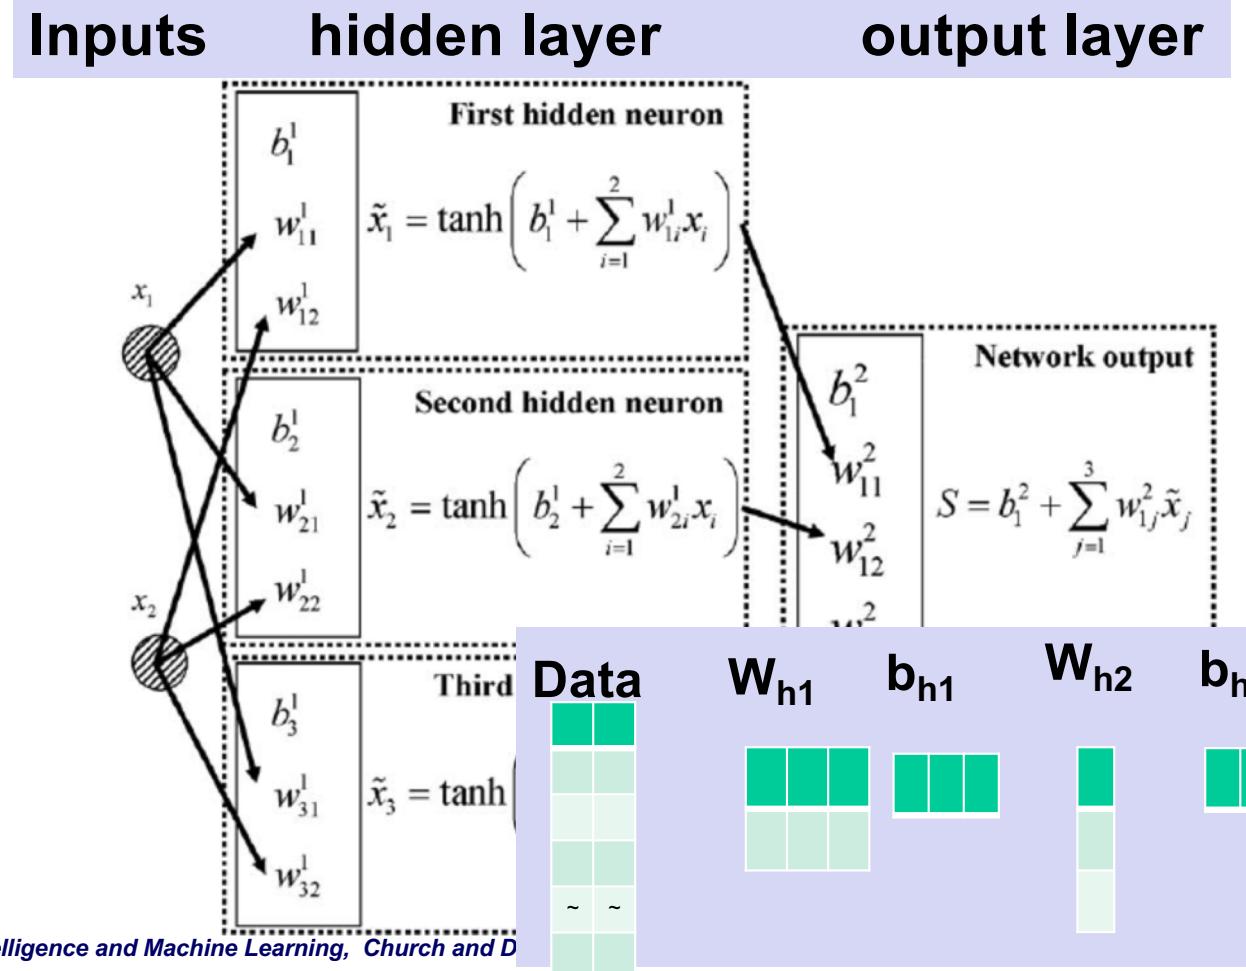
- How can we train the network via Gradient descent? We need to calculate the Gradient  $\nabla$  (the vector of partial derivatives)

$$\frac{\partial E}{\partial W_{ij}}$$

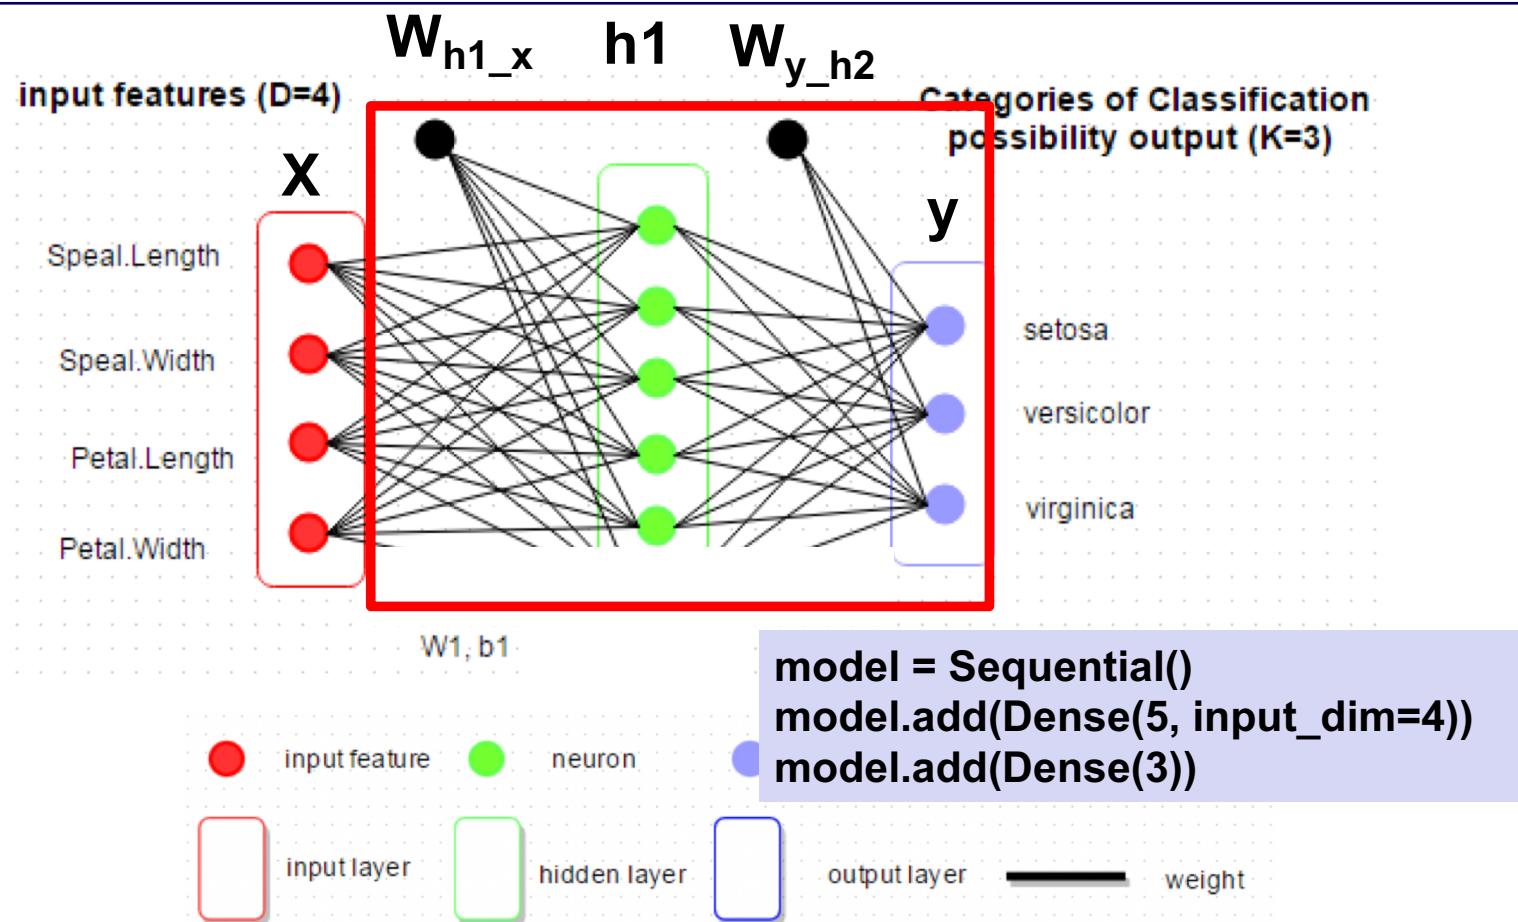
- Now we consider two cases:

- Weights associated with an output layer node
- Weights associated with a hidden layer node...

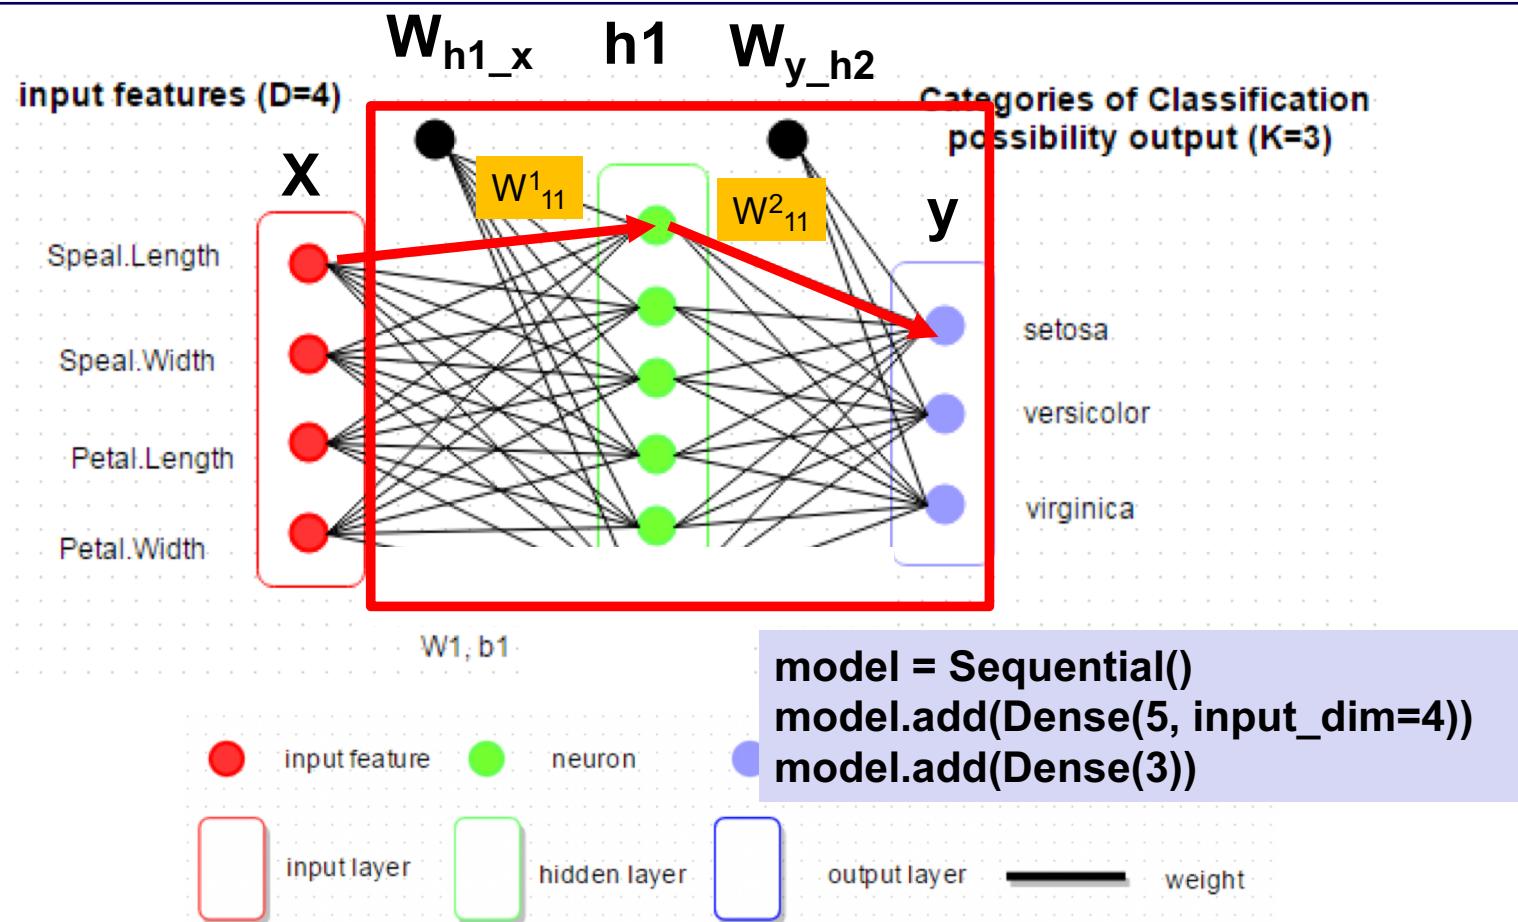
# Activation functions in NN: 2-3-1



# Let's work with a 4 – 5 – 3 MLP



# Let's work with a 4 – 5 – 3 MLP



# Extended multilayer network with Error

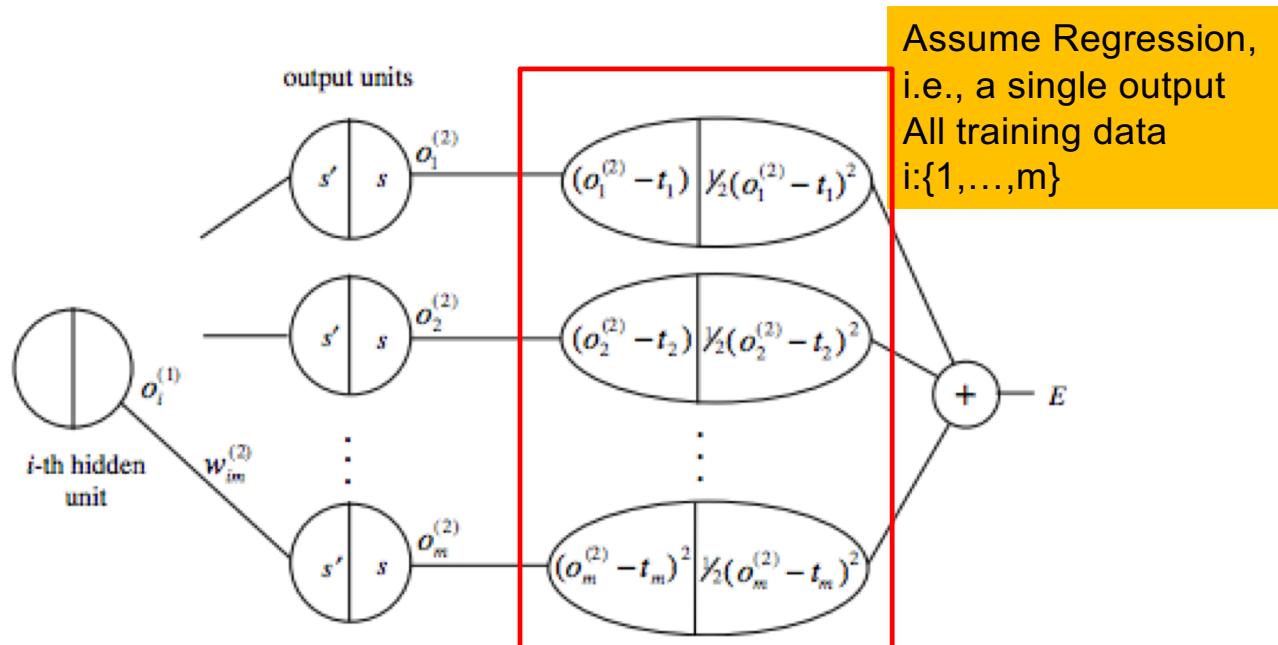
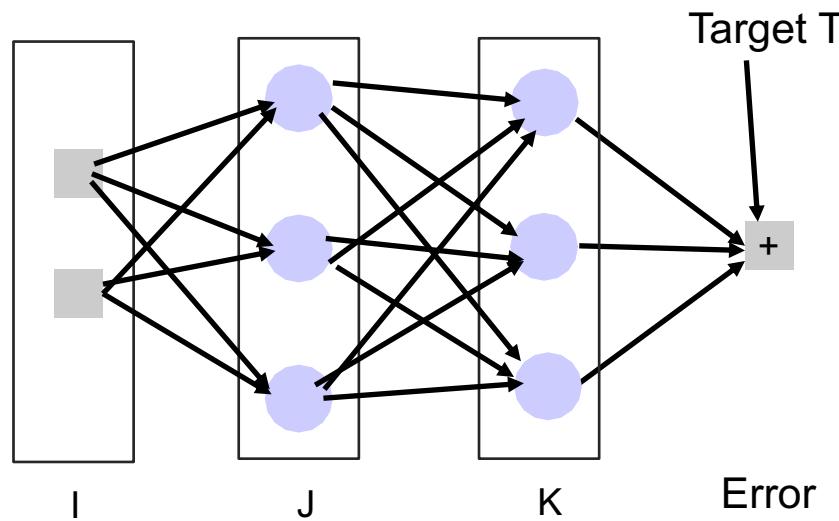


Fig. 7.18. Extended multilayer network for the computation of  $E$

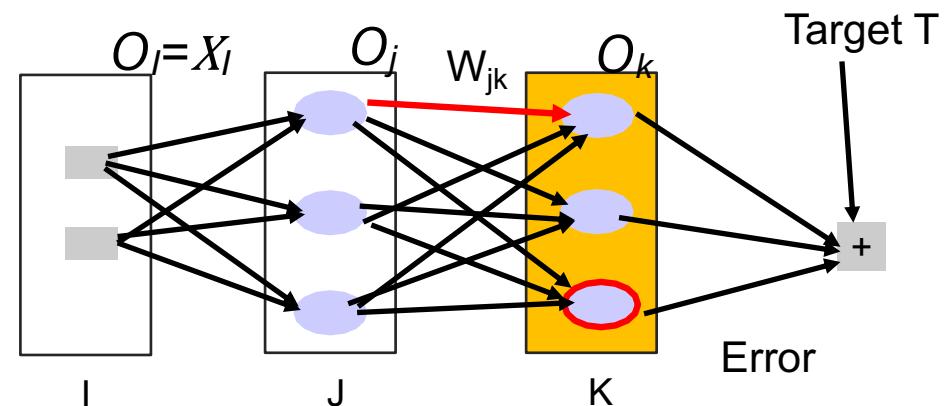
# Extended multilayer network with Error

---

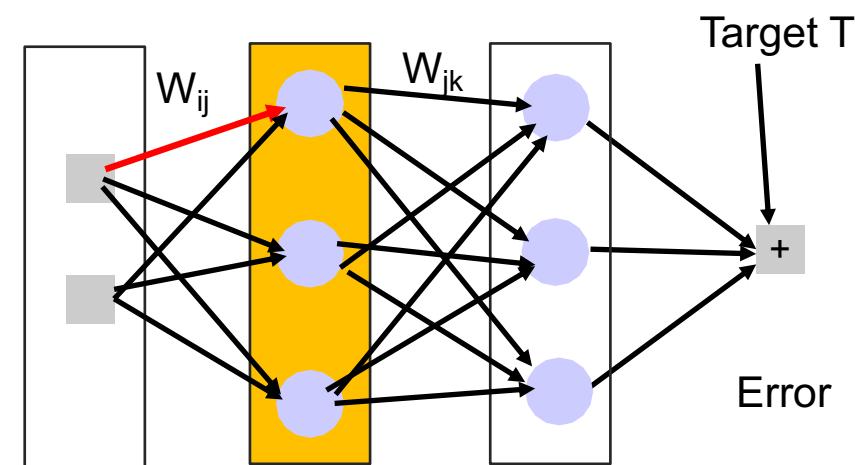


# BackPropagation: Gradient Descent/Chain Rule

- Case 1:Outer node weight



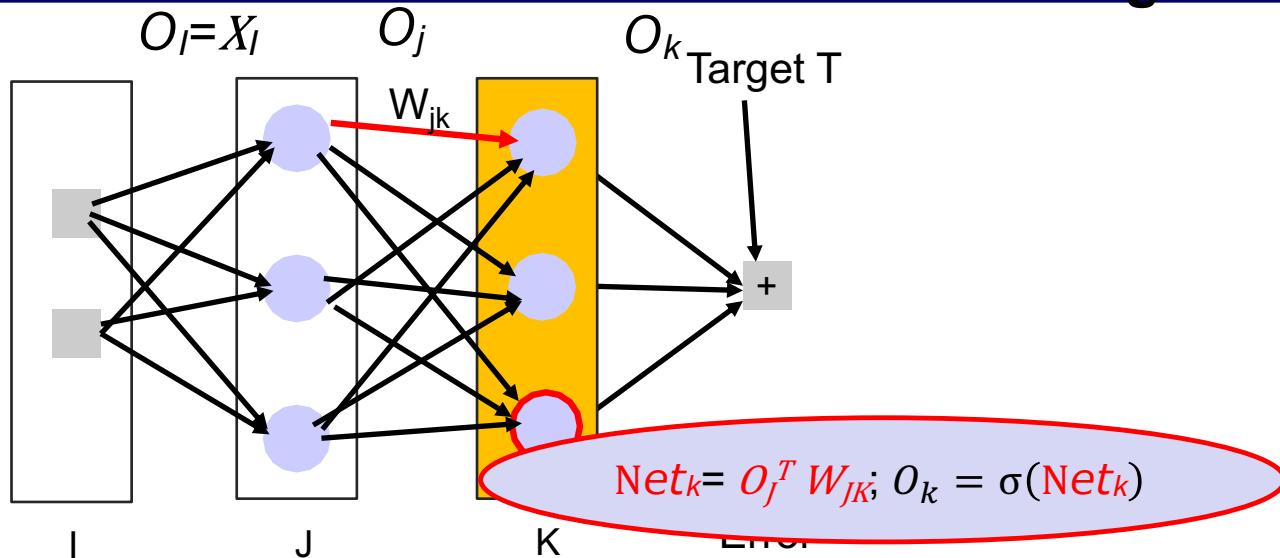
- Case 2: hidden node weight



## Gradient Descent

### Case 1:Outer node weight

---



$$(\sigma(WX) - t_k)^2 = (\sigma(net_k) - t_k)^2$$

Notation:

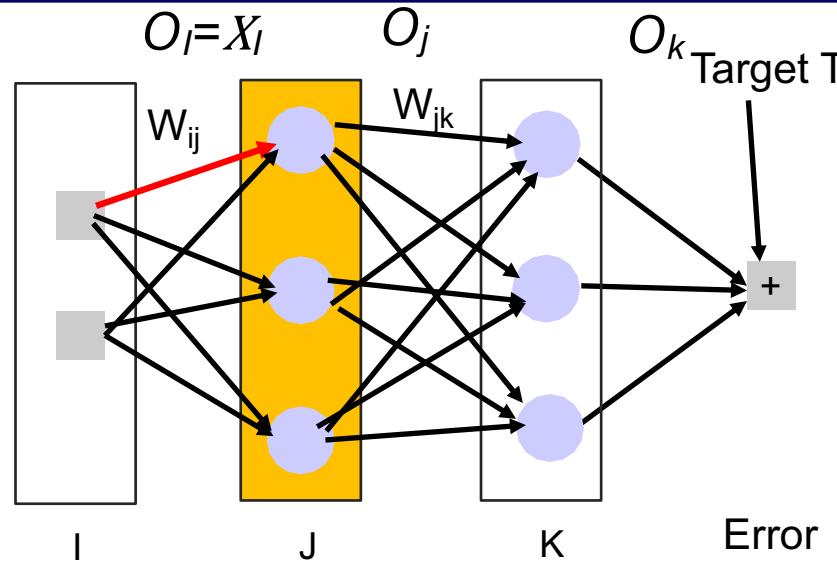
*Each layer, I, J, K, consists of node has:*

- Weighted sum is denoted as  $Net_J$  (i.e.,  $X^T W_J$ )*  
 $\sigma(X_k)$

*Error term,  $\delta_k$*

$\Delta, \delta$  delta

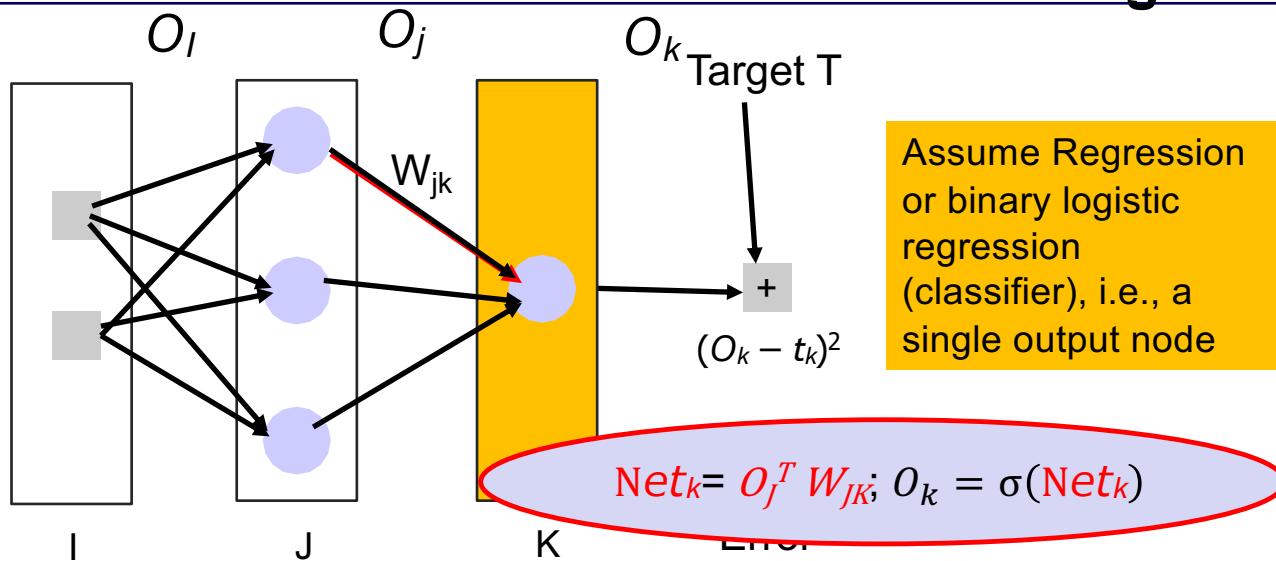
## Gradient Descent Case 2:hidden node weight



## Gradient Descent

### Case 1:Outer node weight

---



Notation:

Each layer,  $I, J$ , and  $K$ , consists of nodes that have in forward prop mode (prediction):

- A Weighted sum is denoted as  $\text{Net}_J$  (i.e.,  $X^T W_J$ );
- an activation  $\sigma(X_k)$  where the activation func is sigmoid

In BackProp mode: each node has error term,  $\delta_k$

<i>Instance\Attr</i>	$x_0$	$x_1$	...	$x_n$	$y$
1	1	-3	..	-7	-1
2	1				+1
...	1	...	...	...	...
<i>L</i> (aka $m$ )	1	0	...	8	-1

$W$
$x_0$
$x_1$
...
$x_n$

## Linear Regression via GD

$$MSE = \sum_{n=1:Train} \frac{1}{2} (\text{Prediction} - \text{target})^2$$

$$MSE = \sum_{n=1:Train} \frac{1}{2} (X_i w - t_i)^2$$

$$MSE = (XW - Y)^2 \quad \#vectorized$$

All training data  $\frac{\partial E}{\partial w} = \frac{\partial}{\partial w} \sum_{n=1:Train} \frac{1}{2} (X_i w - t_i)^2$

One example  $\frac{\partial E}{\partial w} = \frac{\partial}{\partial w} \frac{1}{2} (\mathbf{o} - \mathbf{t})^2 \quad \#Let \mathbf{o} = X_i w - t_i$

$$\frac{\partial E}{\partial w} = (\mathbf{o} - \mathbf{t}) \frac{\partial}{\partial w} (\mathbf{o} - \mathbf{t})$$

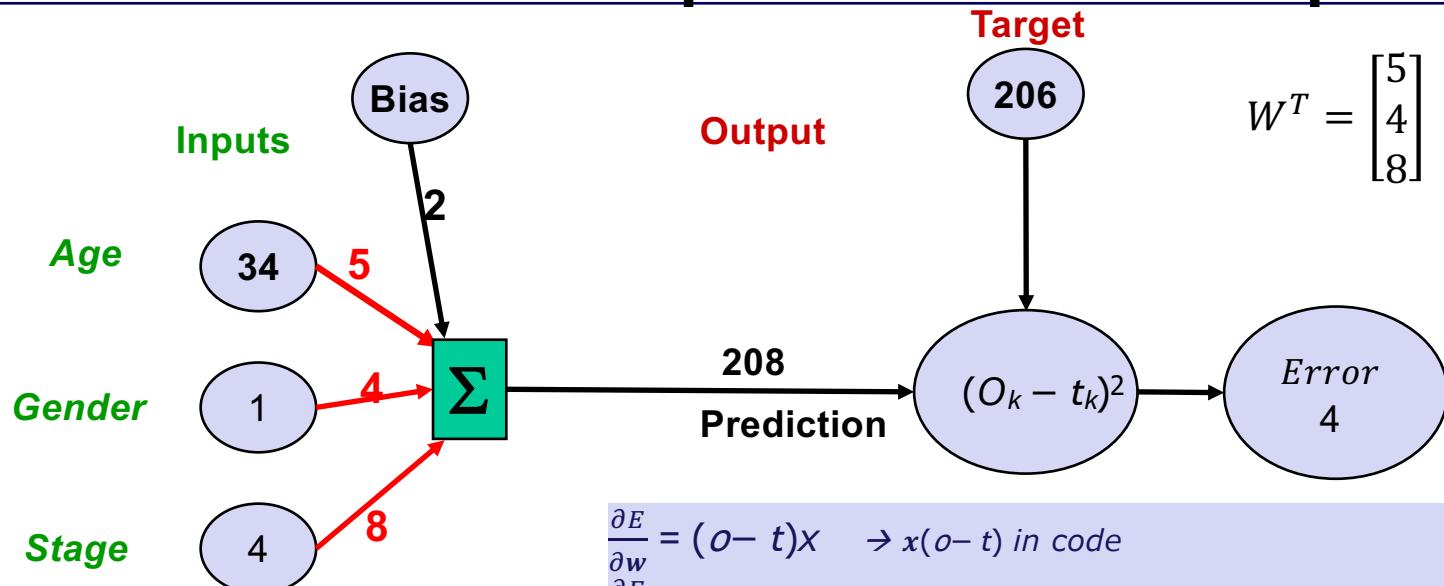
$$\frac{\partial E}{\partial w} = (\mathbf{o} - \mathbf{t}) \frac{\partial}{\partial w} (\mathbf{xw} - \mathbf{t})$$

$$\frac{\partial E}{\partial w} = (\mathbf{o} - \mathbf{t}) \mathbf{x}$$

All training data  $\frac{\partial E}{\partial w} = X^T(\mathbf{o} - \mathbf{t}) \quad \#vectorized; MatrixVector multiplication$

$$\mathbf{w} = \mathbf{w} - \alpha \frac{\partial E}{\partial w} \quad AKA. \quad \mathbf{w} = \mathbf{w} - \alpha \nabla_w$$

# Linear Regression Model: Gradient update for one example

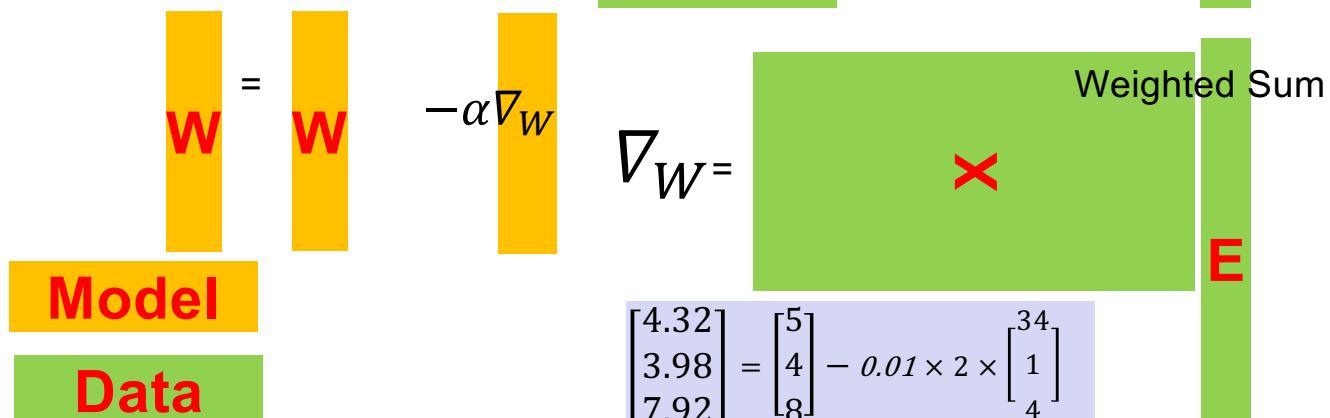
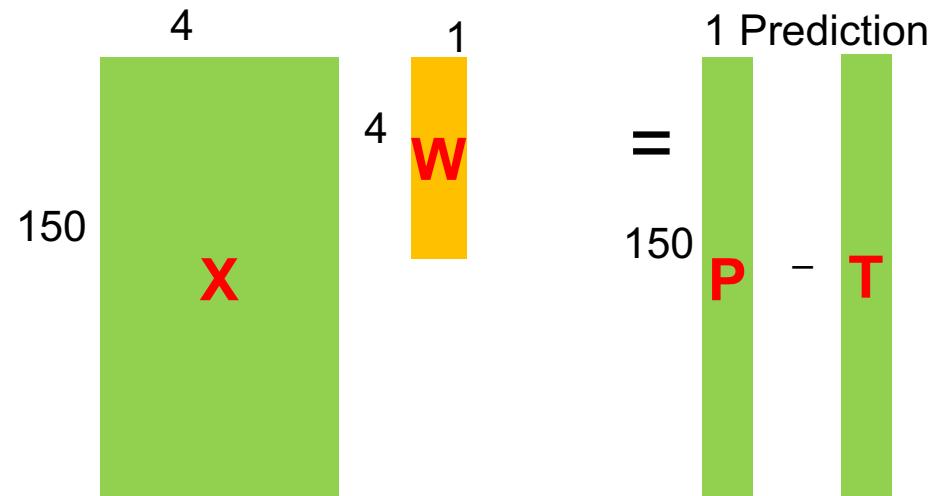
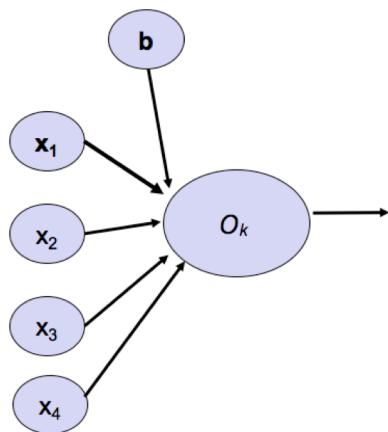


**Weight Gradient for this  
training example [34, 1, 4]; b = 2**

$$\begin{aligned}\frac{\partial E}{\partial w} &= (o - t)x \rightarrow x(o - t) \text{ in code} \\ \frac{\partial E}{\partial w} &= X^T(o - t) \quad \# \text{vectorized; MatrixVector multiplication} \\ w &= w - \alpha \frac{\partial E}{\partial w} \quad \text{AKA. } w = w - \alpha \nabla_w \\ \begin{bmatrix} 4.32 \\ 3.98 \\ 7.92 \end{bmatrix} &= \begin{bmatrix} 5 \\ 4 \\ 8 \end{bmatrix} - 0.01 \times 2 \times \begin{bmatrix} 1 \\ 34 \\ 4 \end{bmatrix} \\ 4.32 &= 5 - 0.01 \times (2 \times 34) \text{ for } i = 1\end{aligned}$$

**Assume a learning rate  $\alpha=0.01$**

# Activation matrices and model matrices for a single target variable



# Chain Rule (Khan Academy)

Chain Rule

$$h(x) = (\sin x)^2$$

↓ Chain Rule

$$h'(x) = \frac{dh}{dx} = 2\sin x \cdot \cos x$$

$\frac{\partial}{\partial x}[x^2] = 2x \quad \frac{\partial}{\partial a}[a^2] = 2a$

$\frac{\partial}{\partial(\sin x)} = (\sin x)^2 = 2\sin x$

$\frac{\partial}{\partial x}[\sin x] = \cos x$

<https://www.khanacademy.org/math/ap-calculus-ab/product-quotient-chain-rules-ab/chain-rule-ab/v/differentiating-composite-functions-2>

Chain Rule

$$h(x) = (\sin x)^2$$

↓ Chain Rule

$$h'(x) = \frac{dh}{dx} = \underbrace{2\sin x \cdot \cos x}_{\frac{d[(\sin x)^2]}{dx} = \frac{d[(\sin x)^2]}{d(\sin x)} \cdot \frac{\partial(\sin x)}{\partial x}}$$

$\frac{\partial}{\partial x}[x^2] = 2x \quad \frac{\partial}{\partial a}[a^2] = 2a$

$\frac{\partial}{\partial(\sin x)} = (\sin x)^2 = 2\sin x$

$\frac{\partial}{\partial x}[\sin x] = \cos x$

## Chain Rule Example 2

---

$$f(x) = \cos^3 x = (\cos x)^3$$

$$f'(x) = \frac{dV}{du} \cdot \frac{du}{dx}$$

$$= \frac{d(\cos x)^3}{d(\cos x)} \cdot \left[ \frac{d(\cos x)}{dx} \right]$$

$$3(\cos x)^2 \cdot -\sin x$$

$x \rightarrow \boxed{\cos} \xrightarrow{u} \boxed{(\quad)^3} \xrightarrow{V(u(x))} \boxed{\cos x} \xrightarrow{V(\cos x)}$

$$f(x) = V(u(x))$$

$$f'(x) = V'(u(x)) \cdot u'(x)$$

$$\frac{d}{dx} [\cos x] = -\sin x$$

$$\frac{d[(\bullet)^3]}{d\bullet} = 3\bullet^2$$

## Chain Rule example 3

$$f(x) = \ln(\sqrt{x})$$

$$f(x) = v(u(x))$$

$$f'(x) = \boxed{v'(u(x))} \boxed{u'(x)}$$

$$= \frac{1}{\sqrt{x}} \cdot \frac{1}{2\sqrt{x}}$$

$$= \boxed{\frac{1}{2x}}$$

$x \rightarrow \boxed{u} \xrightarrow{\sqrt{x}} \boxed{v} \xrightarrow{\ln(x)} \ln(\sqrt{x})$

$u(x) = \sqrt{x} = x^{\frac{1}{2}}$     $v(x) = \ln(x)$

$u'(x) = \frac{1}{2} x^{-\frac{1}{2}}$     $v'(x) = \frac{1}{x}$

$\frac{1}{2} \cdot \frac{1}{x^{\frac{1}{2}}} \quad v'(u(x)) = \frac{1}{u(x)} = \boxed{\frac{1}{\sqrt{x}}}$

$\frac{1}{2} \cdot \frac{1}{\sqrt{x}}$

2 functions composed: step 1 SQRT(X); step 2 ln(SQRT(X))  
 Derivative of the outside function wrt inside function  
 Derivative of the inside function wrt X

# MLP Regression Model: 13-13Relu-1

---

- Boston dataset

```
1 # define base model
2 def baseline_model():
3     # create model
4     model = Sequential()
5     model.add(Dense(13, input_dim=13, kernel_initializer='normal', activation='relu'))
6     model.add(Dense(1, kernel_initializer='normal'))
7     # Compile model
8     model.compile(loss='mean_squared_error', optimizer='adam')
9     return model
```

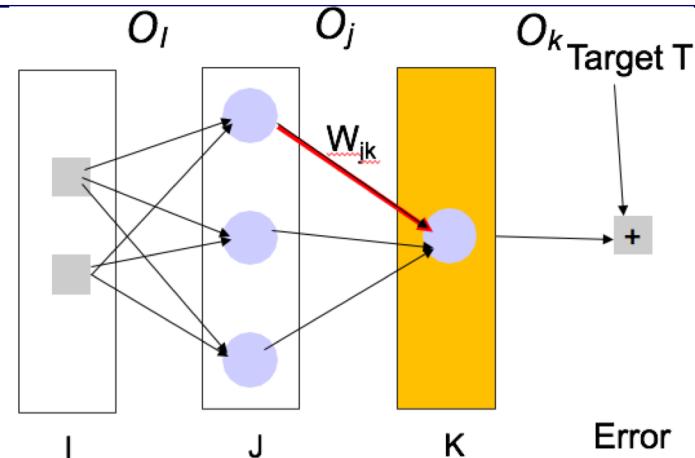
$$\text{MSE}(X; \text{MLP}_{\text{Regression}}) = \frac{1}{2n} (I(W_2 (\sigma(W_1 X))) - t_k)^2 \text{ where } I \text{ is identity}$$

$$\text{MSE}(X; \text{MLP}_{\text{Regression}}) = \frac{1}{2n} (\sigma(W_2 (\sigma(W_1 X))) - t_k)^2 \text{ Where target } \epsilon [0,1]$$

Below we assume a sigmoid activation function in the output layer. HERESY!

# Given a NN where Target $\in [0,1]$

Sigmoid  
Regression!



$$\text{MSE}(X; \text{MLP}_{\text{Regression}}) = \frac{1}{2n} (I(W_2 (\sigma(W_1 X))) - t_k)^2 \quad \# \text{where } I \text{ is identity}$$

$$\text{MSE}(X; \text{MLP}_{\text{Regression}}) = \frac{1}{2n} (\sigma(W_2 (\sigma(W_1 X))) - t_k)^2 \quad \# \text{Here we assume Target } \in [0,1]$$

$$O_k = \frac{1}{2n} (\sigma(W_2 (\sigma(W_1 X))) - t_k)^2$$

$$\frac{\partial E}{\partial W_{jk}} = \frac{\partial}{\partial W_{jk}} \frac{1}{2} (O_k - t_k)^2 \quad \text{MSE}$$

$$\frac{\partial E}{\partial W_{jk}} = (O_k - t_k) \frac{\partial}{\partial W_{jk}} (O_k - t_k) \quad \text{After applying chain rule once}$$

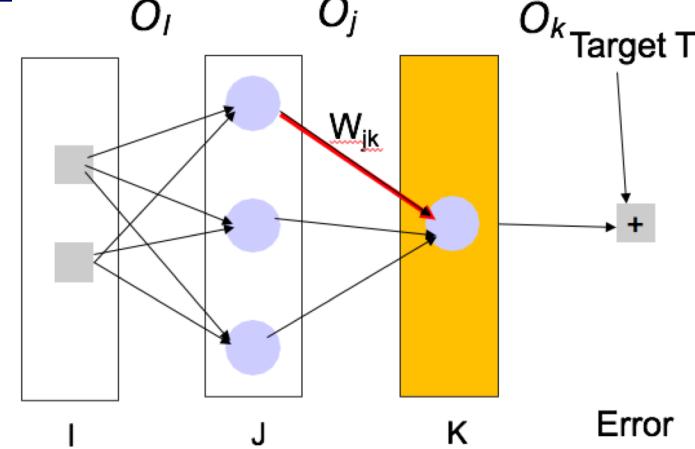
Need to unpack the network using the chain rule

## Given MSE MLP:

*Assume MSE loss where Target  $\epsilon [0,1]$*

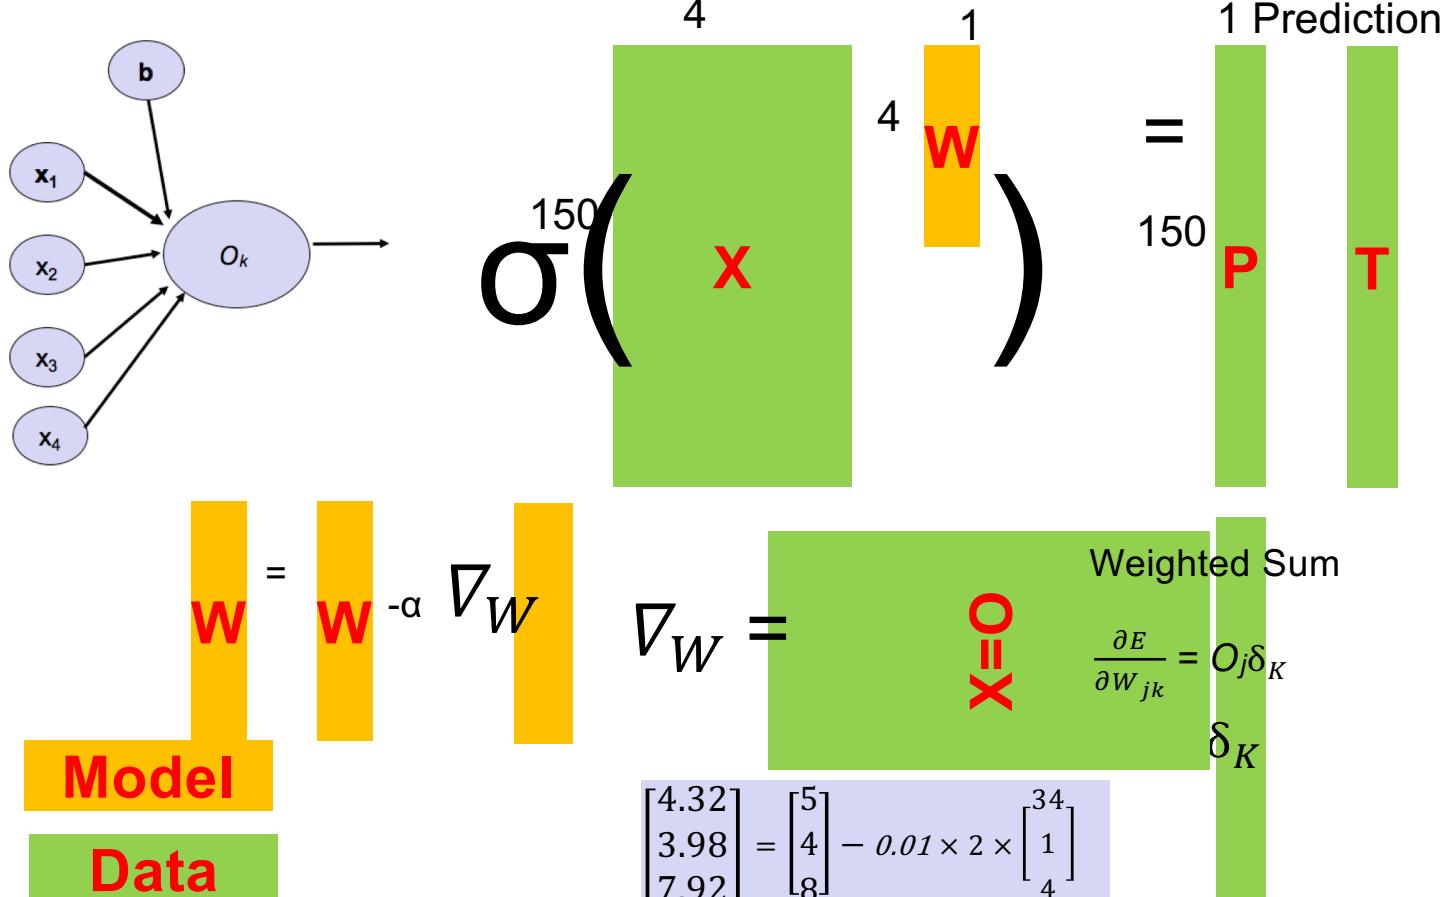
$$\text{Net}_k = O_J W_{JK}; \quad O_k = \sigma(\text{Net}_k)$$

$$\frac{\partial E}{\partial W_{jk}} = (O_k - t_k) \frac{\partial}{\partial W_{jk}} (O_k - t_k)$$



*Assume MSE loss where Target  $\epsilon [0,1]$  ( $\sigma(XW) - t_k)^2 = (\sigma(\text{net}_k) - t_k)^2$*

# Activation matrices and model matrices for a single target variable



Assume MSE loss where Target  $\epsilon[0,1]$   $(\sigma(WX) - t_k)^2 = (\sigma(\text{net}_k) - t_k)^2$

## Output layer node: single output MSE Loss (0-1 range)

$$\frac{\partial E}{\partial W_{jk}} = \frac{\partial}{\partial W_{jk}} \frac{1}{2} (O_k - t_k)^2$$

Chain rule where  $u = (O_k - t_k)^2$

$$\frac{\partial E}{\partial W_{jk}} = (O_k - t_k) \frac{\partial}{\partial W_{jk}} (O_k - t_k)$$

$$\frac{\partial E}{\partial W_{jk}} = (O_k - t_k) \frac{\partial}{\partial W_{jk}} (O_k)$$

$\sigma(\text{net}_k) = \sigma(O_j W_{JK})$

$$\frac{\partial E}{\partial W_{jk}} = (O_k - t_k) \frac{\partial}{\partial W_{jk}} (\sigma(\text{net}_k))$$

$$\frac{\partial E}{\partial W_{jk}} = (O_k - t_k) (\sigma(\text{net}_k)(1 - \sigma(\text{net}_k))) \frac{\partial}{\partial W_{jk}} \text{net}_k$$

Chain rule where  $u = \text{net}_k$

$$\frac{\partial E}{\partial W_{jk}} = (O_k - t_k) (\sigma(\text{net}_k)(1 - \sigma(\text{net}_k))) O_j$$

simplify as  $O_k = \sigma(\text{Net}_k)$

$$\frac{\partial E}{\partial W_{jk}} = (O_k - t_k) O_k (1 - O_k) O_j$$

Simplify

Error Vector (for weighted sum in the gradient calc.)  
for node output node k

$$\text{Let } \delta_K = (O_k - t_k) O_k (1 - O_k)$$

For terms involving k simplify and replace with one term  $\delta_K$

$$\frac{\partial E}{\partial W_{jk}} = O_j \delta_K$$

Net<sub>k</sub> =  $O_j^T W_{JK}$ ;  $O_k = \sigma(\text{Net}_k)$

$$(\sigma(WX) - t_k)^2 = (\sigma(\text{net}_k) - t_k)^2$$

## Output layer node: single output; single example

$$\frac{\partial E}{\partial W_{jk}} = \frac{\partial}{\partial W_{jk}} \frac{1}{2} (O_k - t_k)^2$$

Chain rule where  $u = (O_k - t_k)^2$

$$\frac{\partial E}{\partial W_{jk}} = (O_k - t_k) \frac{\partial}{\partial W_{jk}} (O_k - t_k)$$

$$\frac{\partial E}{\partial W_{jk}} = (O_k - t_k) \frac{\partial}{\partial W_{jk}} (O_k)$$

$$\frac{\partial E}{\partial W_{jk}} = (O_k - t_k) \frac{\partial}{\partial W_{jk}} (\sigma(\text{net}_k))$$

$$\frac{\partial E}{\partial W_{jk}} = (O_k - t_k) (\sigma(\text{net}_k)(1-\sigma(\text{net}_k))) \frac{\partial}{\partial W_{jk}} \text{net}_k$$

Chain rule where  $u = \text{net}_k$

$$\frac{\partial E}{\partial W_{jk}} = (O_k - t_k) (\sigma(\text{net}_k)(1-\sigma(\text{net}_k))) O_j$$

$$\frac{\partial E}{\partial W_{jk}} = (O_k - t_k) O_k (1-O_k) O_j$$

Simplify

$$\text{Let } \delta_K = (O_k - t_k) O_k (1-O_k)$$

For terms involving  $k$  simplify and replace with one term  $\delta_K$

$$\frac{\partial E}{\partial W_{jk}} = O_j \delta_K$$

Weight update

$X \rightarrow Y$  in linear regression

$O_i \rightarrow O_v$  in NN

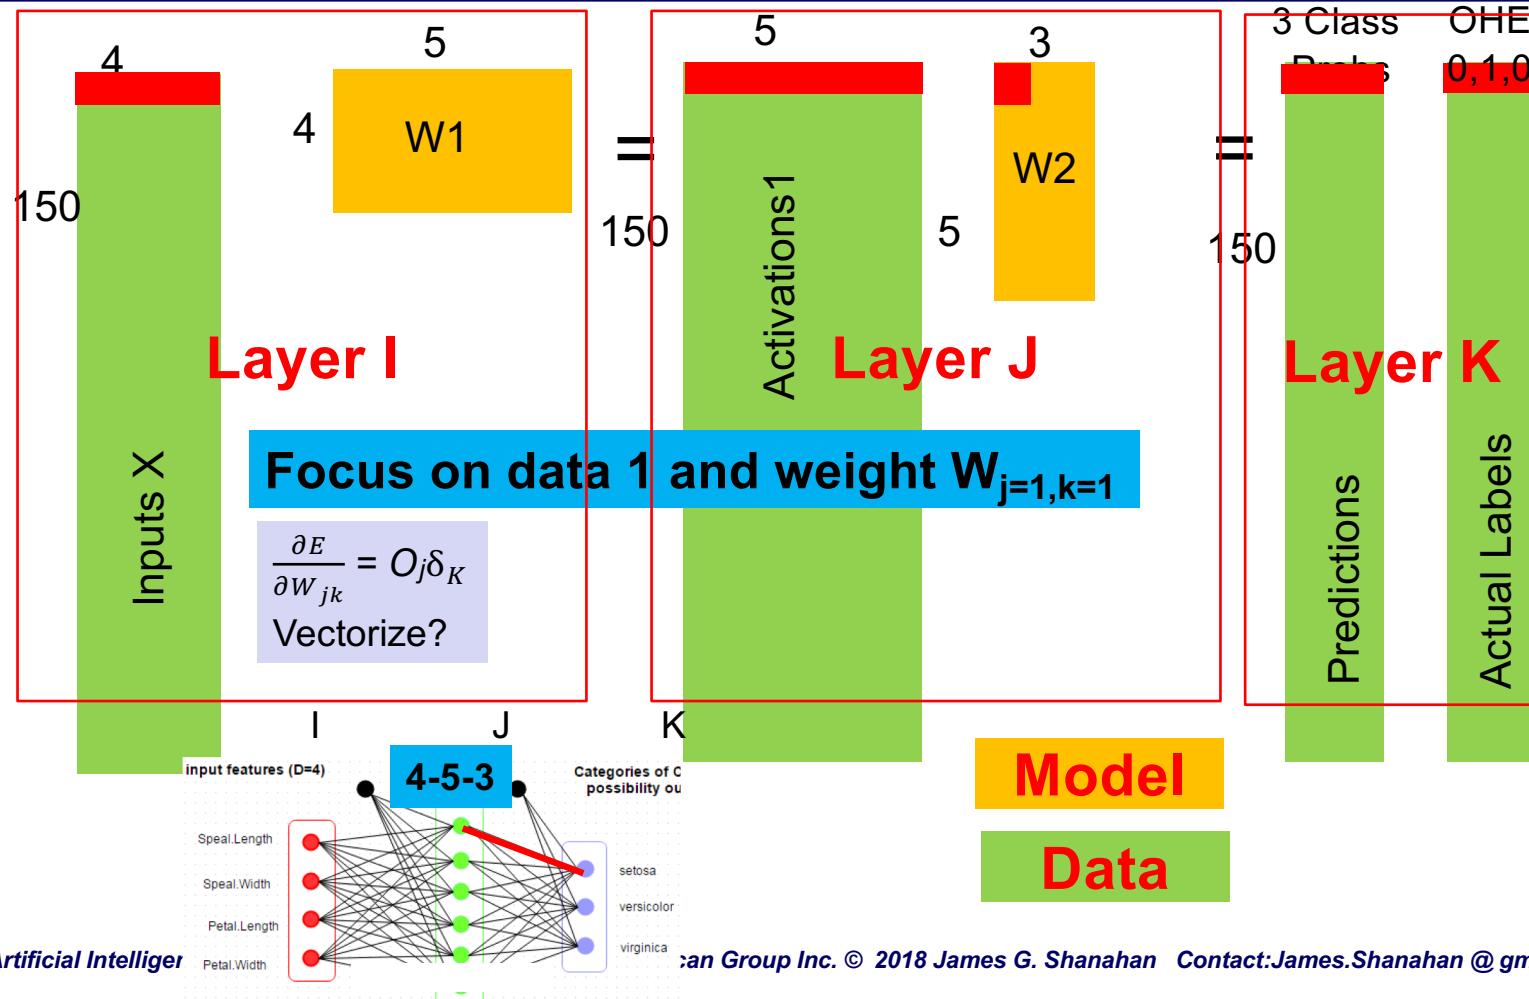
Think of this as error (weight) for output node  $k$

$O$  be the input vector in the case of

regression but now it means the

is from layer  $I$  here

# Activation matrices and weight matrices for Iris NN with a 4-5-3 architecture



# $\nabla$ the gradient chorus

- What is the gradient for linear regression?

- Chorus

- The gradient is the weighted sum of the training data, where the weights are proportional to the error (for each example) !

$$\frac{\partial E}{\partial w} = \text{weight example} (O - t) X$$
$$W = W - \alpha \times \frac{\partial E}{\partial w}$$
$$\begin{bmatrix} 4.32 \\ 3.98 \\ 7.92 \end{bmatrix} = \begin{bmatrix} 5 \\ 4 \\ 8 \end{bmatrix} - 0.01 \times 2 \times \begin{bmatrix} 34 \\ 1 \\ 4 \end{bmatrix}$$
$$4.32 = 5 - 0.01 \times (2 \times 34) \text{ for } i=1$$

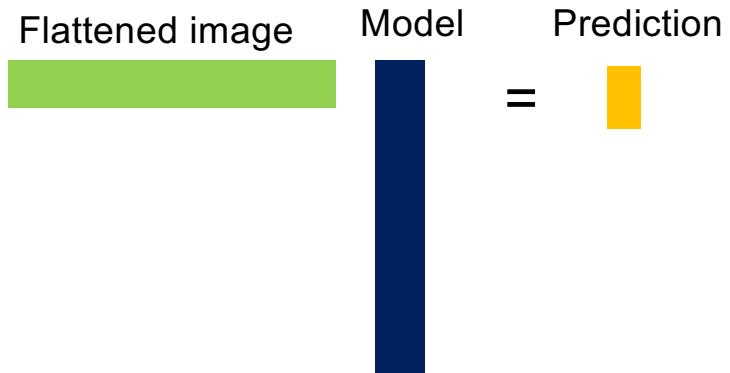
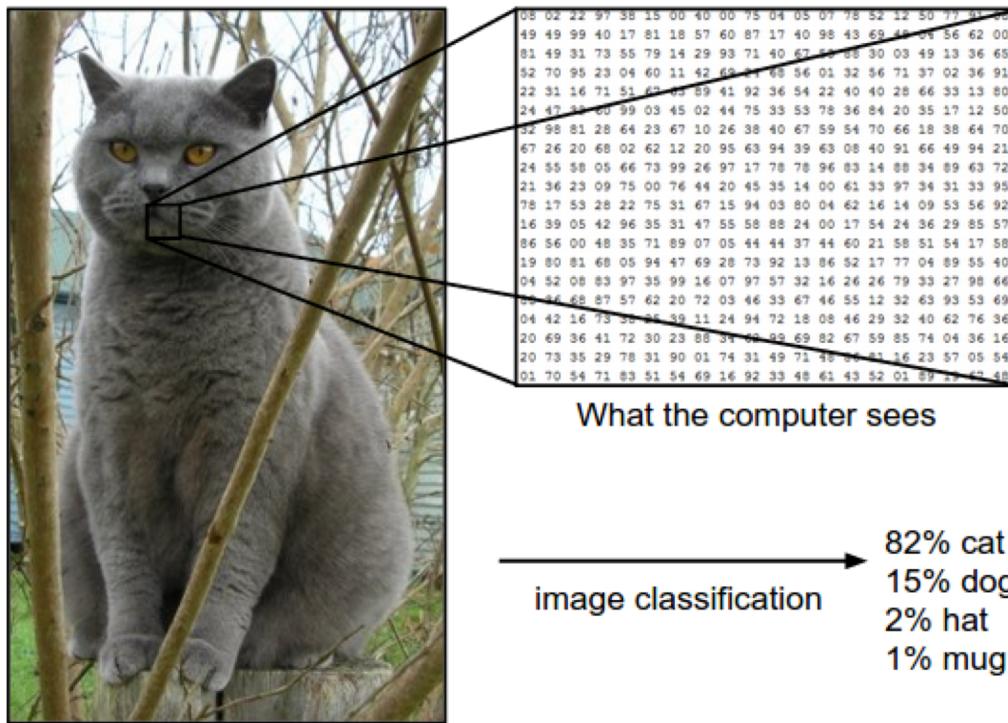


# Outline

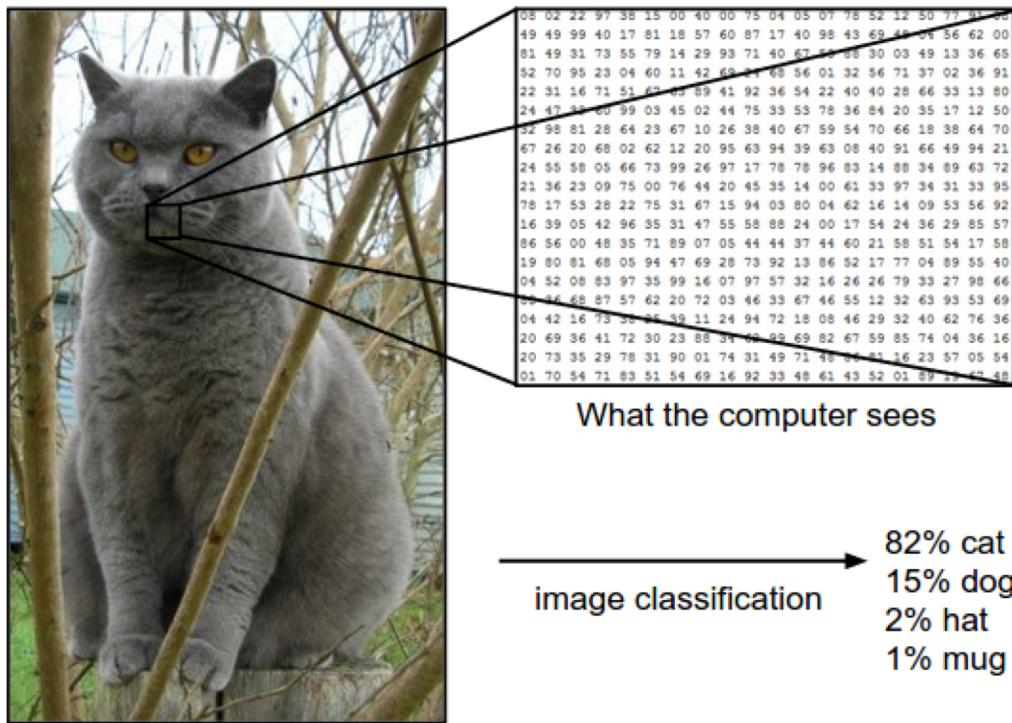
---

- 1. Introduction**
- 2. ML and deep learning review**
- 3. What is computer vision?**
- 1. Computer vision
- 2. Convolutional Neural Nets (CNNs)
- 4. Deep NN Architectures for CV Tasks**
- 1. Backbone networks
- 5. Solving edge-based IoT**
- 6. Conclusions and Next steps**

# Flatten image and learn a logistic regression classifier



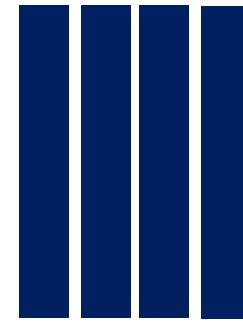
# Flatten image and learn a logistic regression classifier



Flattened image



Model

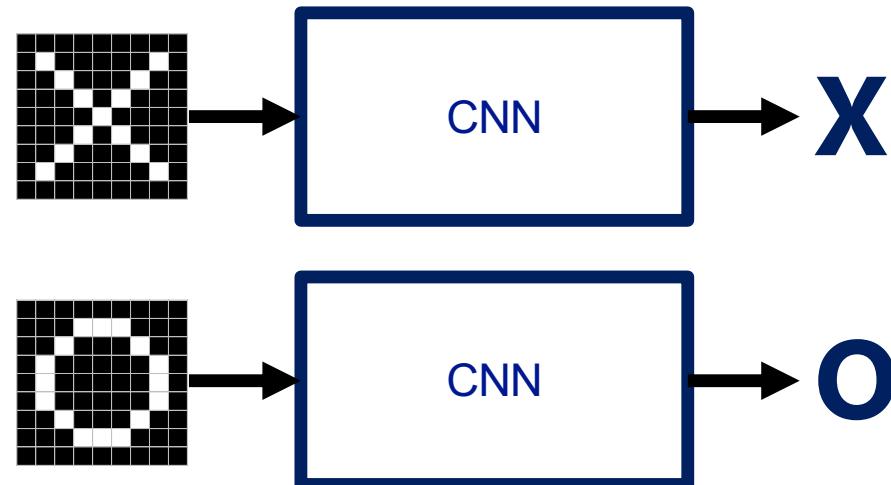


Prediction



Classify these images:  
Ignores the spatial relations

# Classify these images: take into account the spatial relations



Slides and YouTube Video by [Brandon Rohrer \[August 2016\]](#)

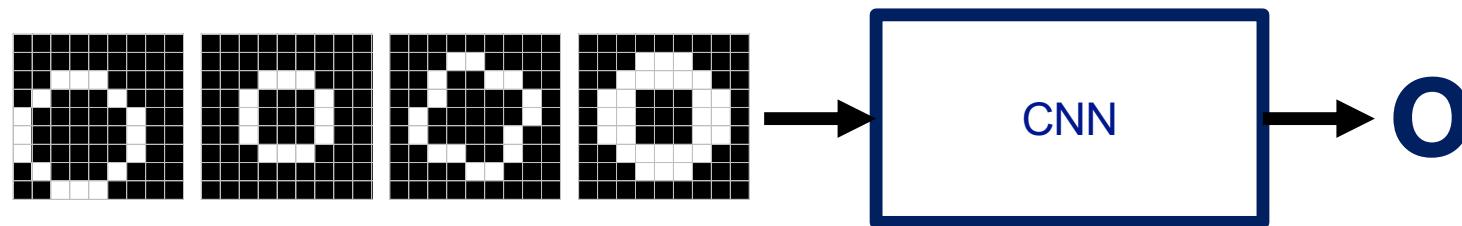
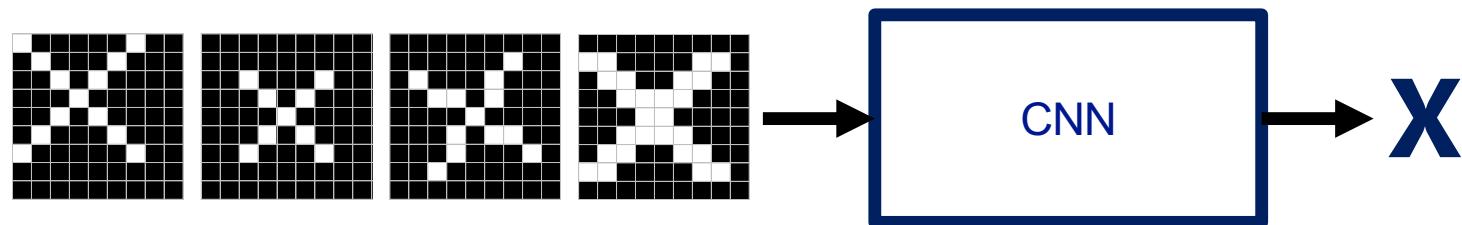
<https://www.youtube.com/watch?v=FmpDlaiMleA>

SLIDES

[https://github.com/brohrer/public-hosting/blob/master/how\\_CNNs\\_work.pptx?raw=true](https://github.com/brohrer/public-hosting/blob/master/how_CNNs_work.pptx?raw=true)

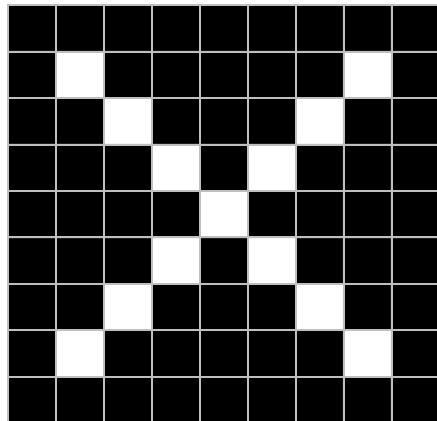
# Trickier cases

---



# Deciding is hard

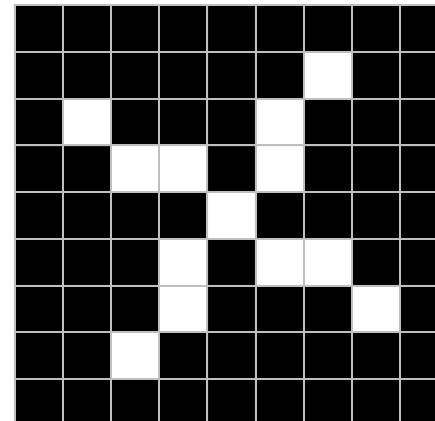
---



?

—

—



# What computers see

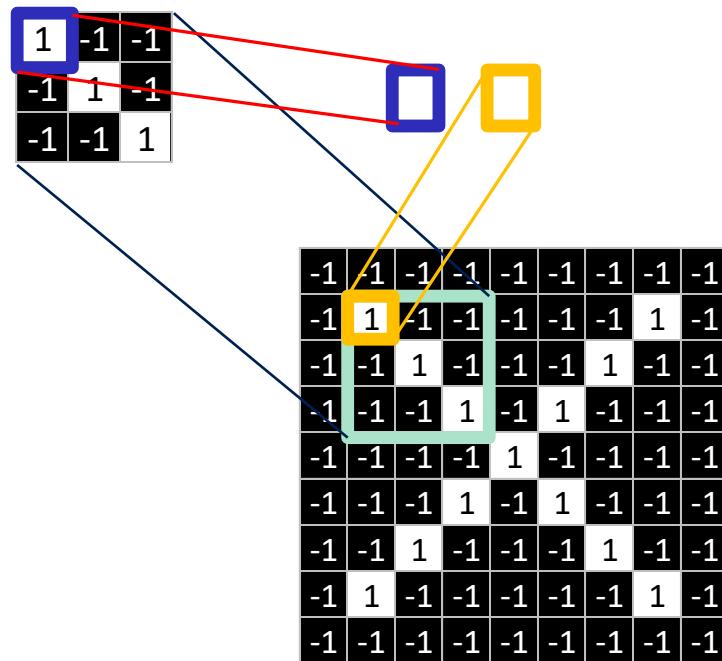
---

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1

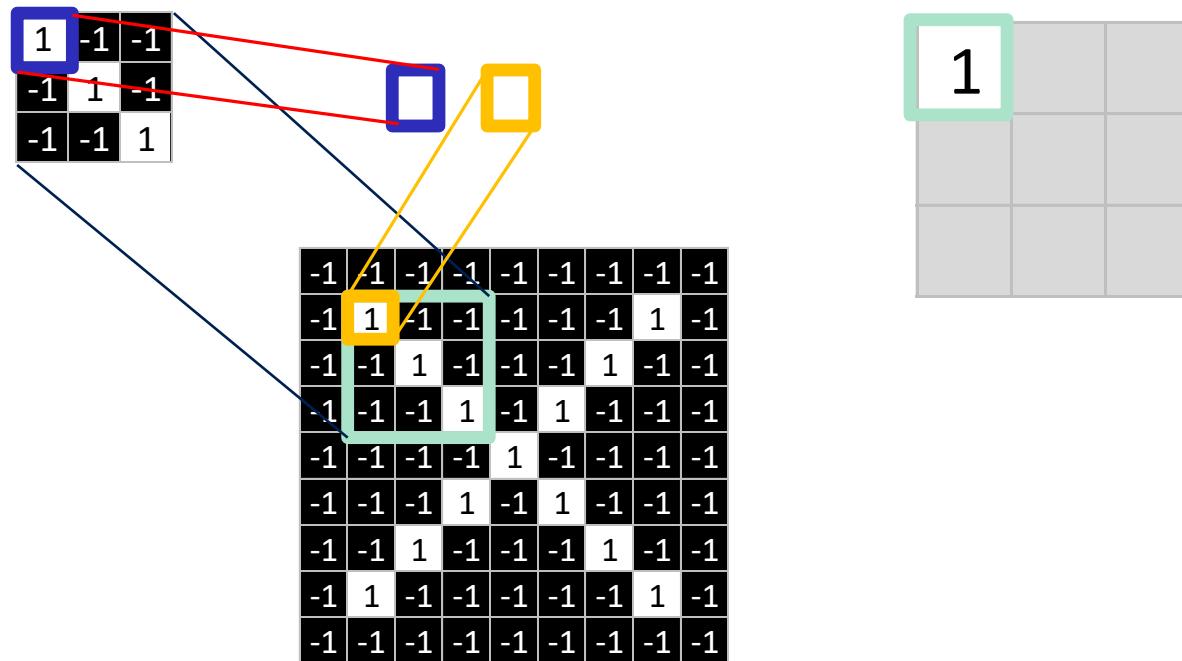


-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	1	-1	-1	-1
-1	-1	1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	1	-1	-1
-1	-1	-1	1	-1	-1	1	1	-1
-1	-1	-1	1	-1	-1	-1	-1	1
-1	-1	1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1

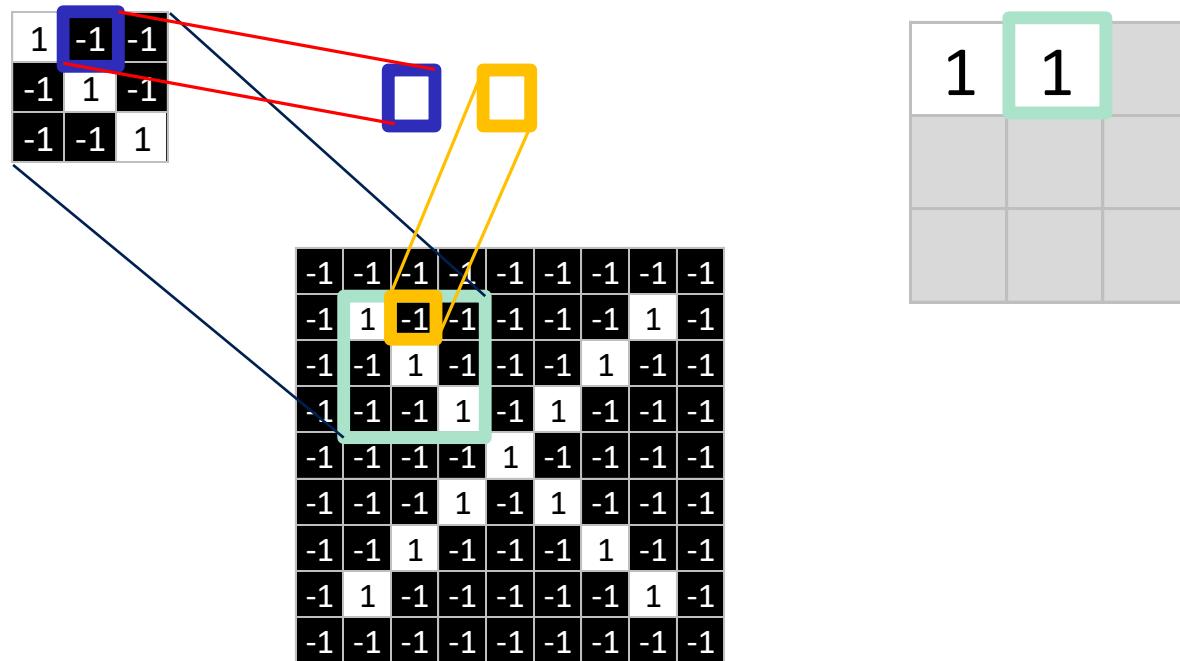
# Filtering: The math behind the match



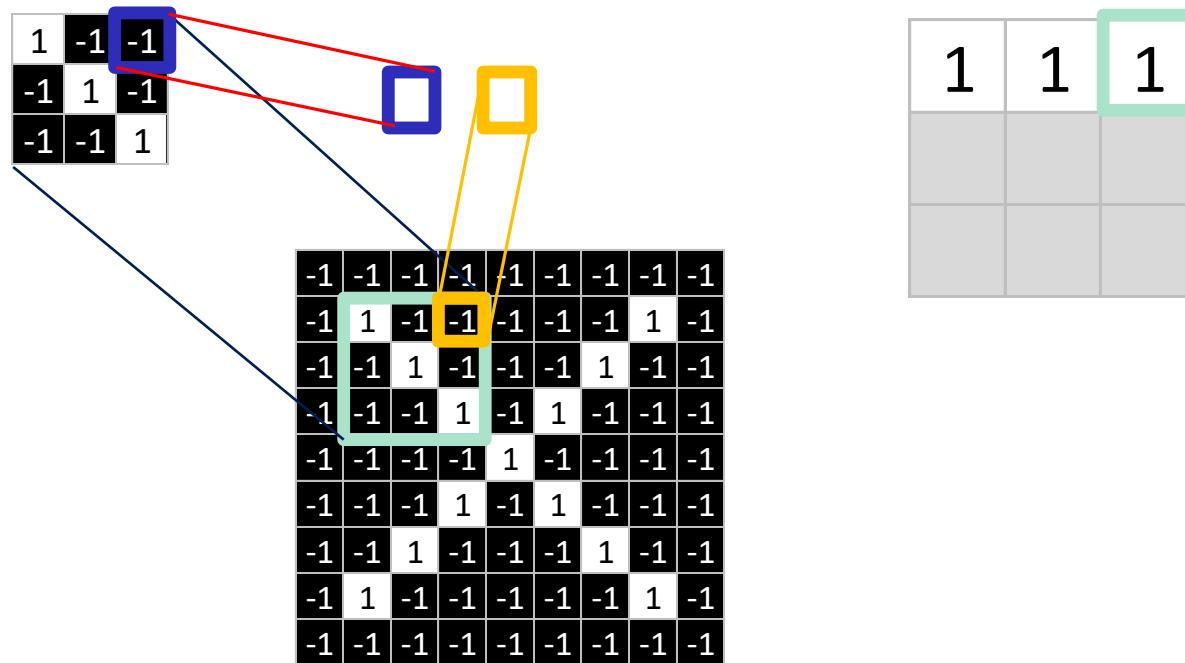
# Filtering: The math behind the match



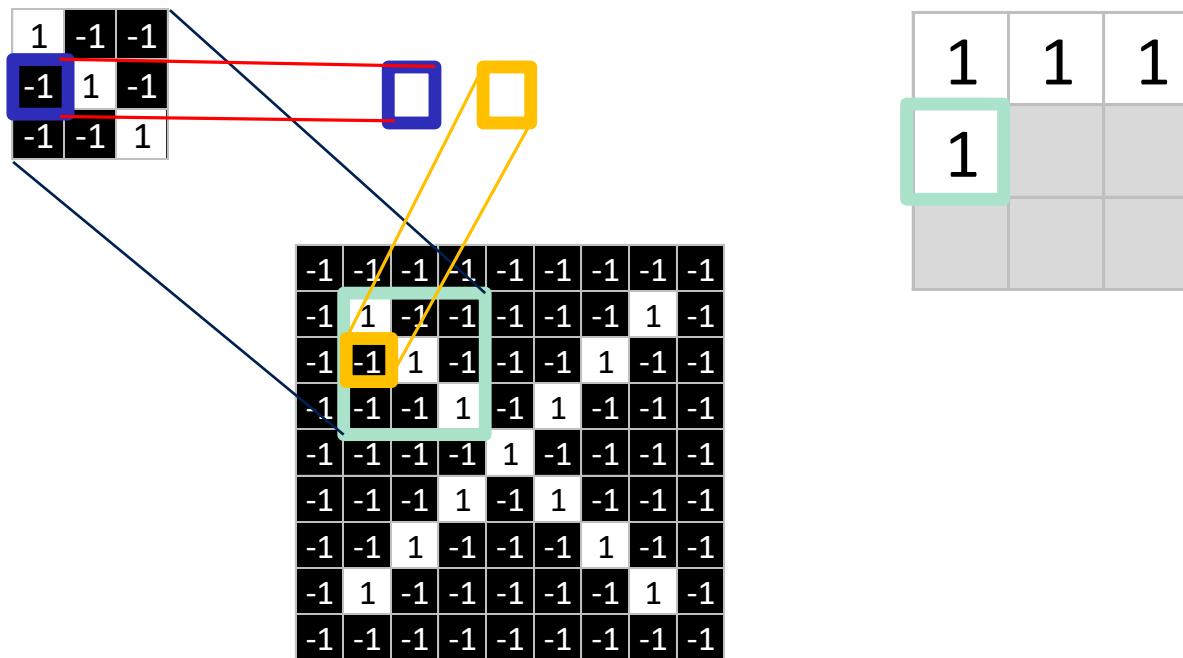
# Filtering: The math behind the match



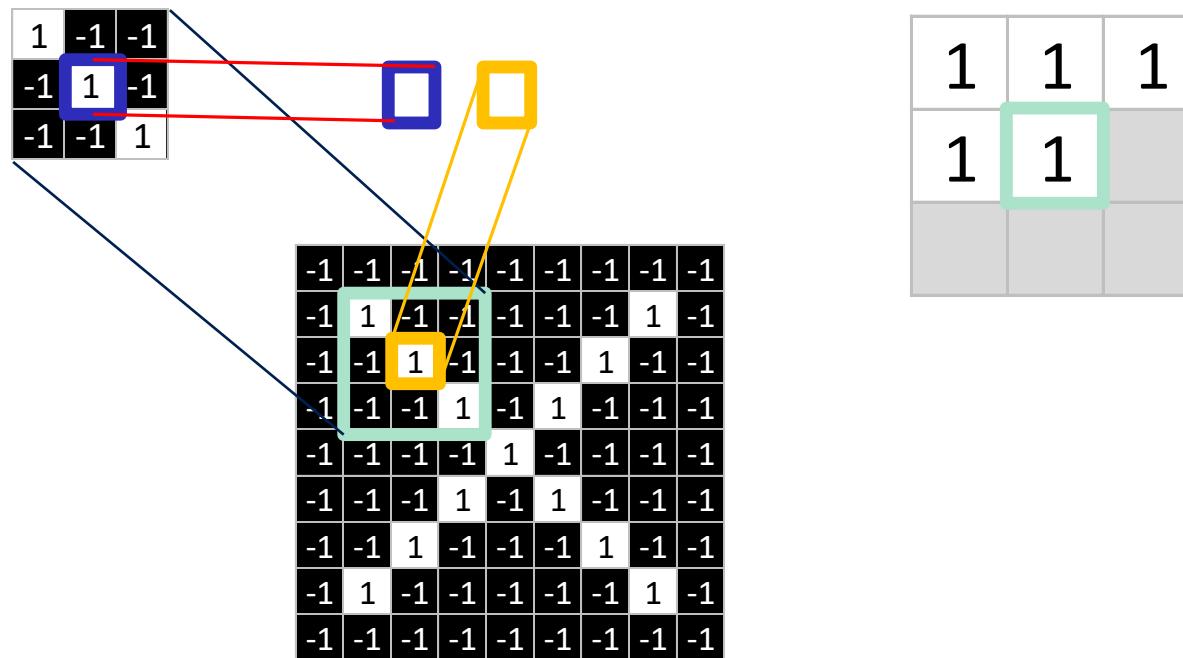
# Filtering: The math behind the match



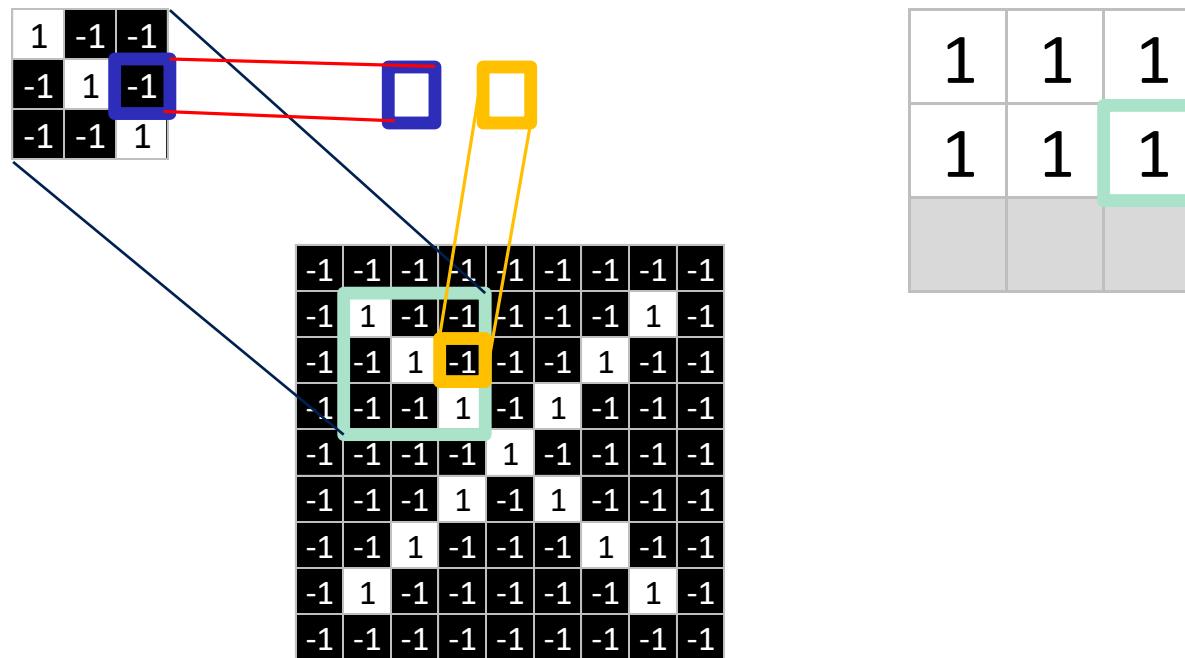
# Filtering: The math behind the match



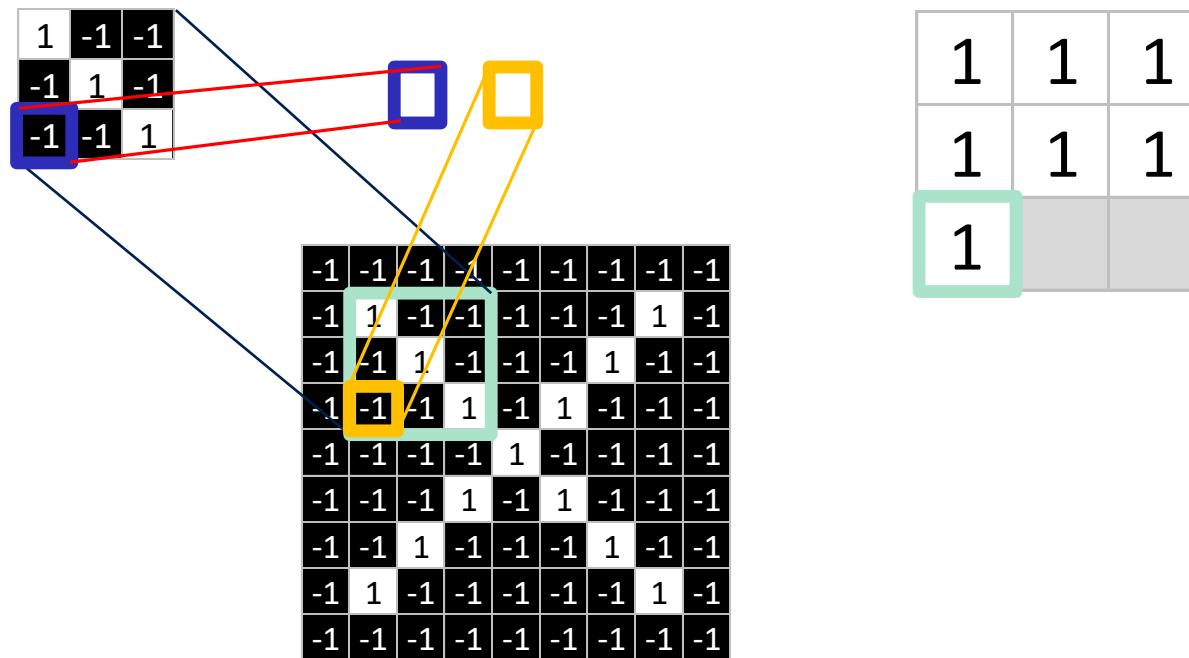
# Filtering: The math behind the match



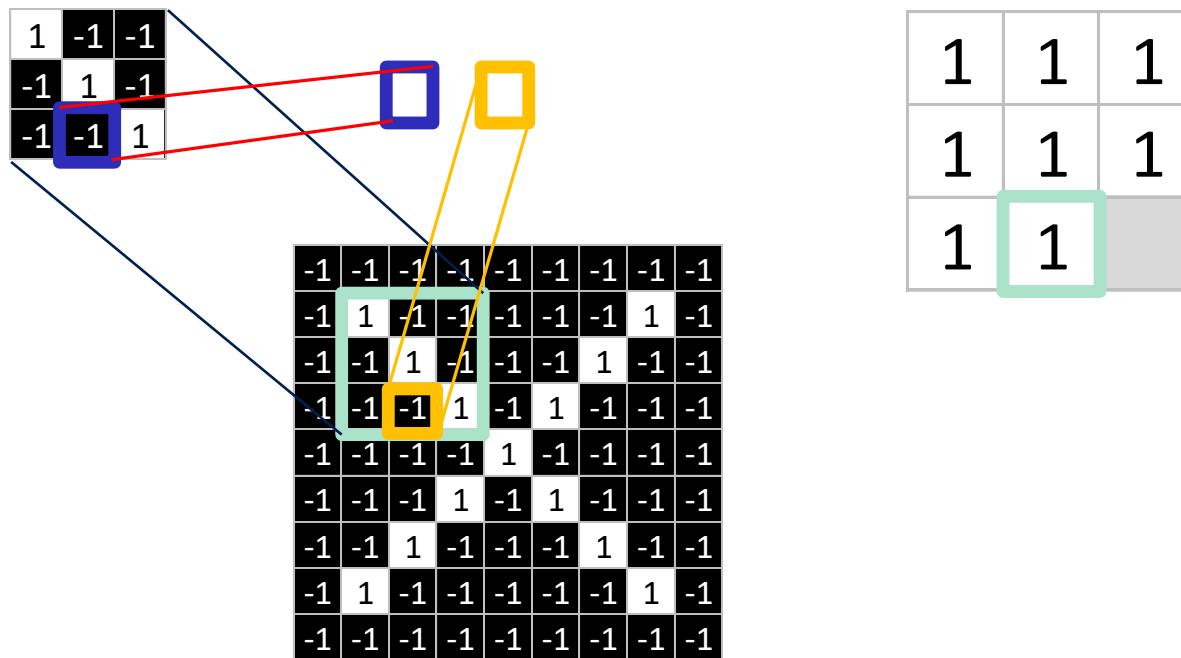
# Filtering: The math behind the match



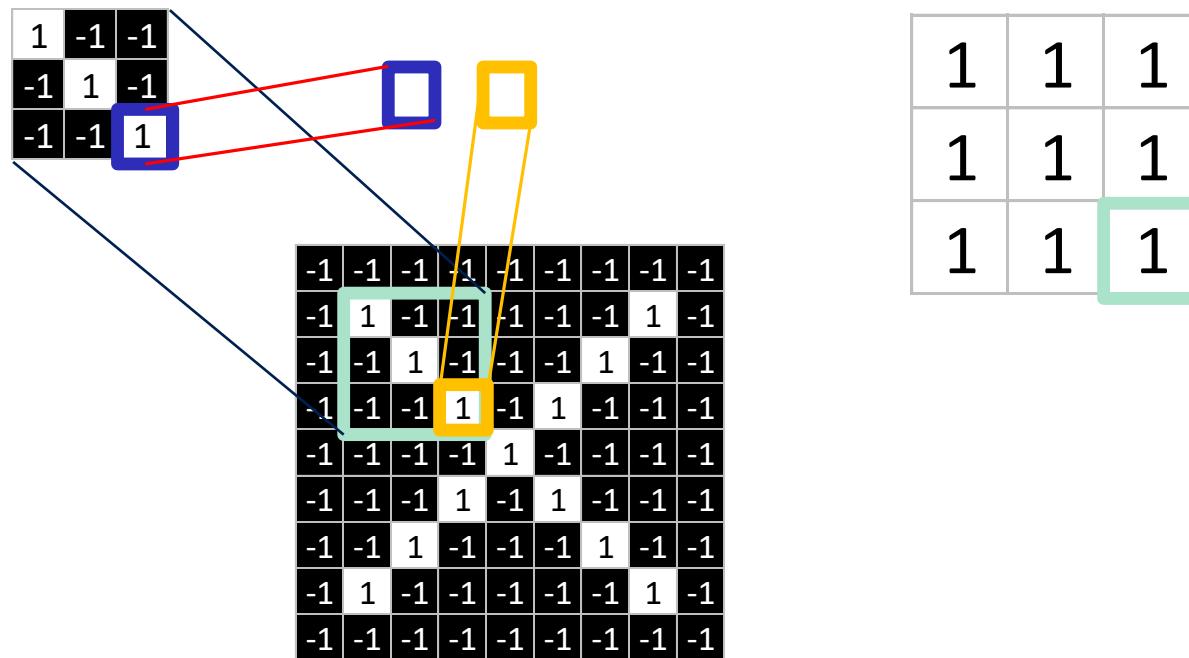
# Filtering: The math behind the match



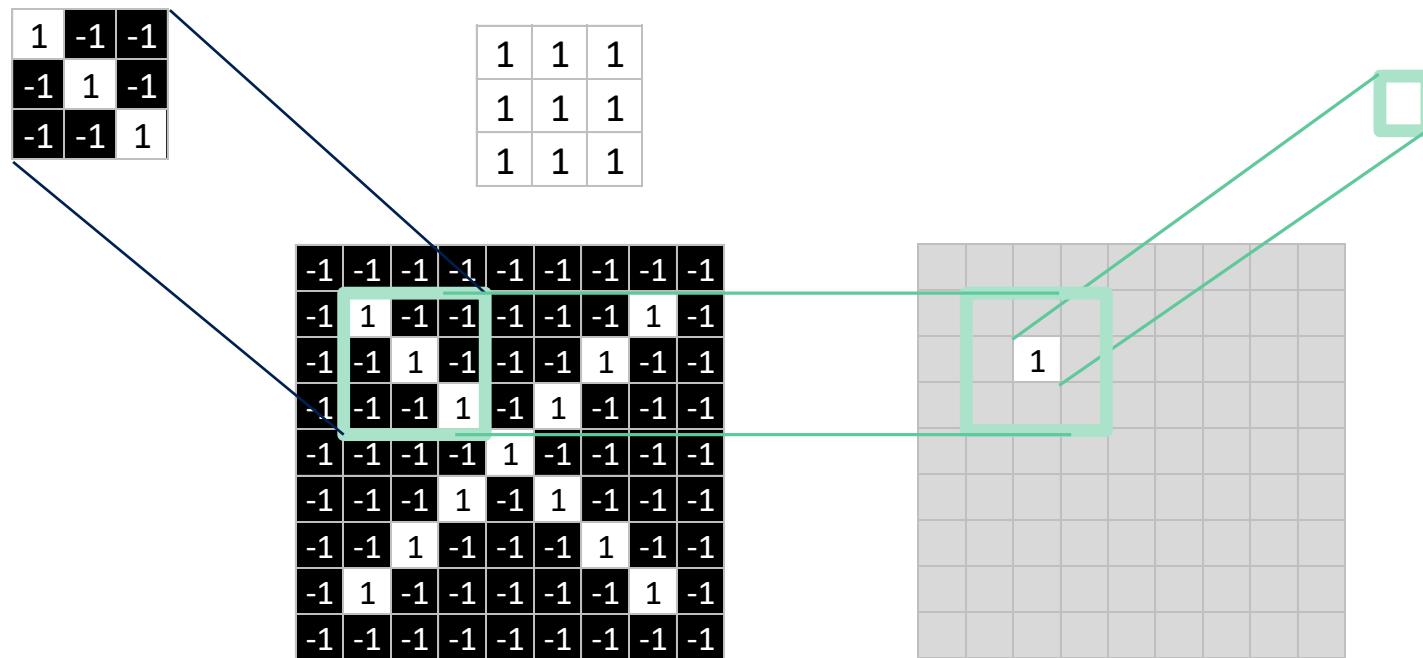
# Filtering: The math behind the match



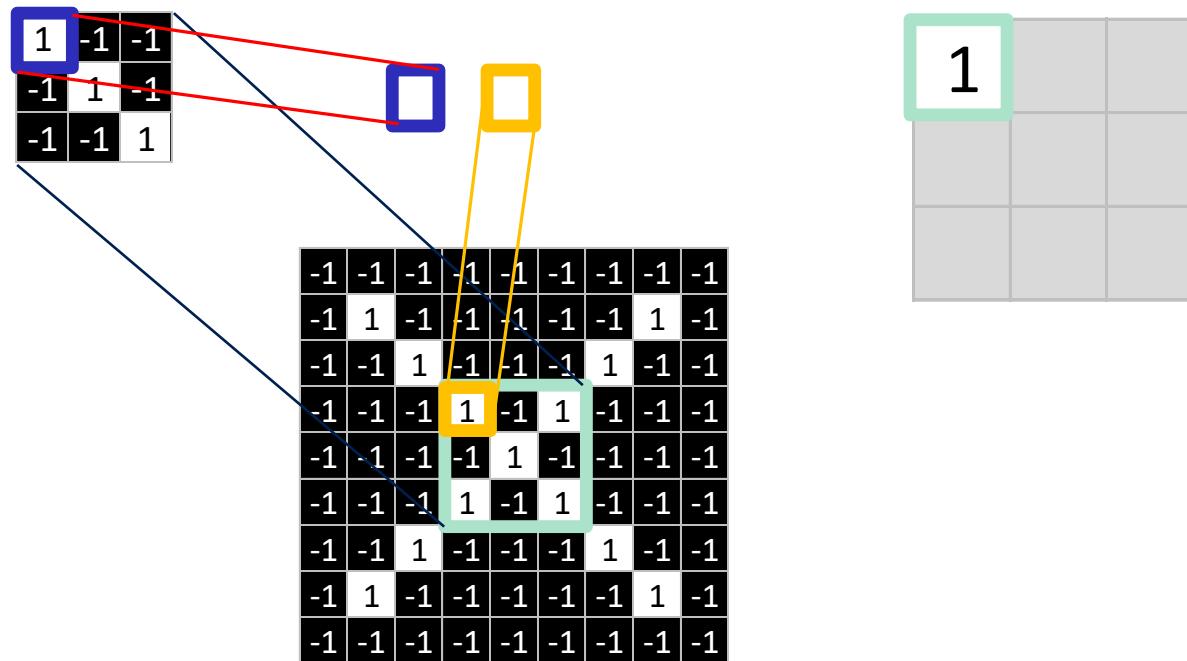
# Filtering: The math behind the match



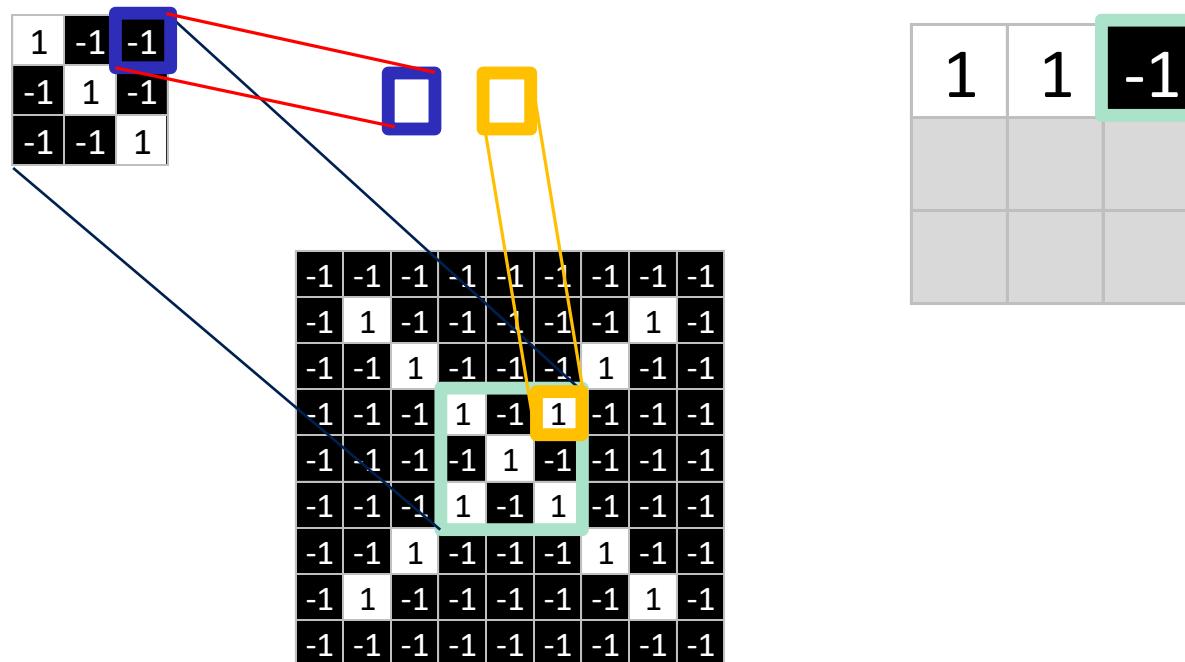
# Filtering: The math behind the match



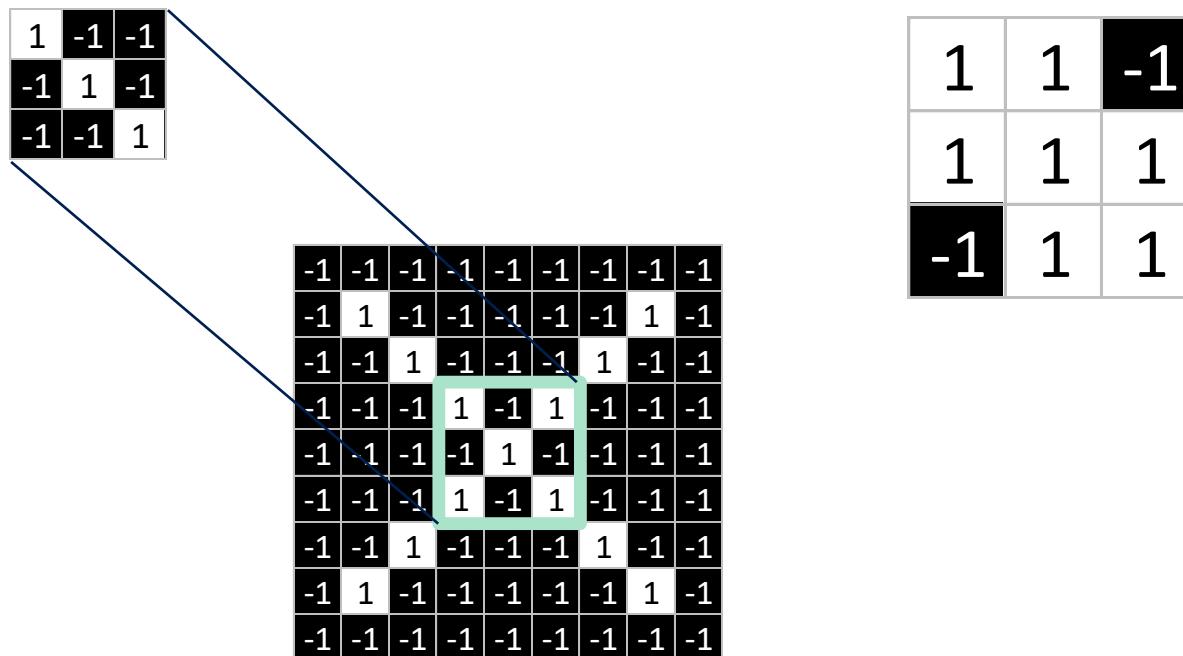
# Filtering: The math behind the match



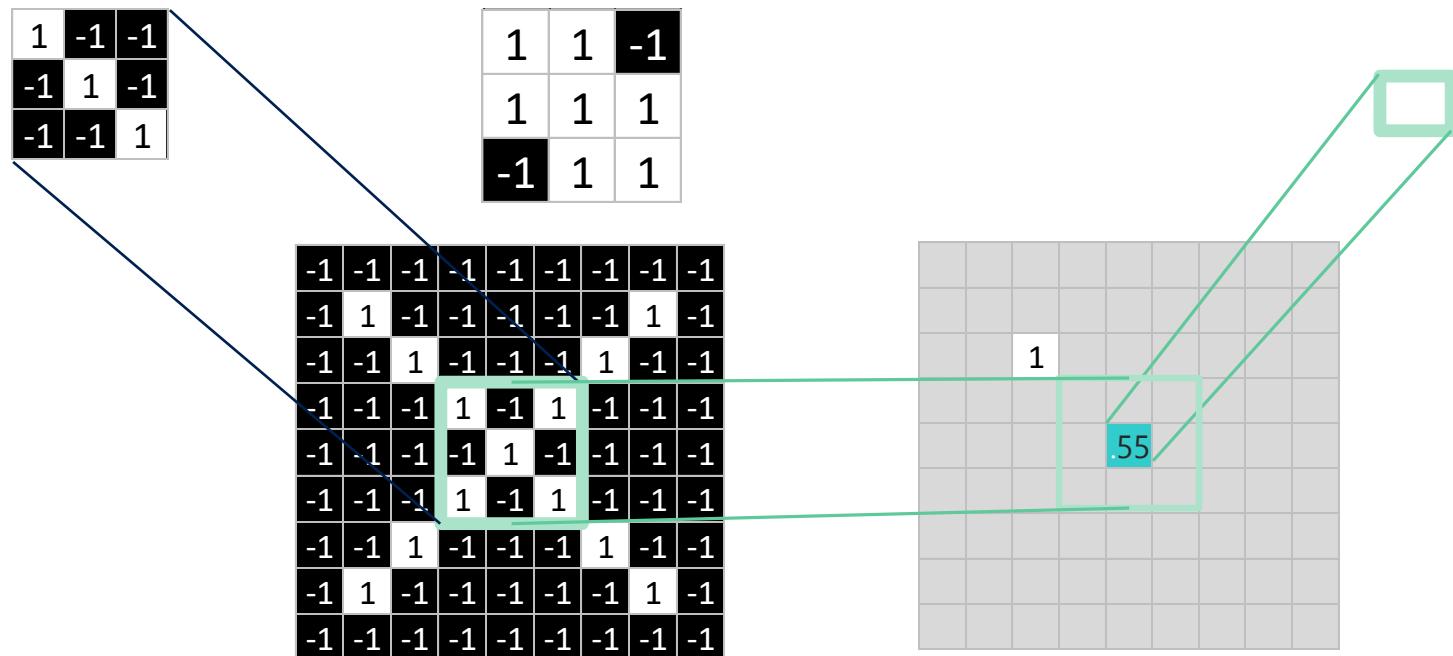
# Filtering: The math behind the match



# Filtering: The math behind the match



# Filtering: The math behind the match



# Convolution: Trying every possible match

1	-1	-1
-1	1	-1
-1	-1	1

-1	-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1	-1



0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

# Convolution: Trying every possible match

---

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1

⊗

1	-1	-1
-1	1	-1
-1	-1	1

=

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	1	1	-1	-1
-1	-1	-1	1	1	1	-1	-1	-1
-1	-1	-1	-1	1	1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	1	-1	-1



1	-1	-1
-1	1	-1
-1	-1	1

=

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	1	-1	-1
-1	-1	-1	-1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	1	-1	-1
-1	-1	-1	-1	-1	-1	-1	1	-1



1	-1	1
-1	1	-1
1	-1	1

=

0.33	-0.55	0.11	0.11	0.11	-0.55	0.33
-0.55	0.55	-0.55	0.33	-0.55	0.55	-0.55
0.11	-0.55	0.55	-0.77	0.55	-0.55	0.11
-0.11	0.33	-0.77	1.00	-0.77	0.33	-0.11
0.11	-0.55	0.55	-0.77	0.55	-0.55	0.11
-0.55	0.55	-0.55	0.33	-0.55	0.55	-0.55
0.33	-0.55	0.11	0.11	0.11	-0.55	0.33

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	1	-1	-1
-1	-1	-1	-1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	1	-1	-1
-1	-1	-1	-1	-1	-1	-1	1	-1



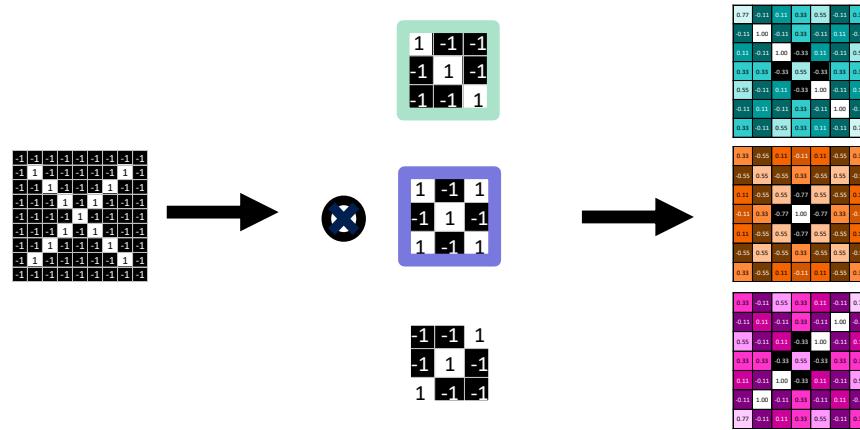
-1	-1	1
-1	1	-1
1	-1	-1

=

0.33	-0.11	0.55	0.33	0.11	-0.11	0.77
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.11	-0.11	1.00	0.33	0.11	-0.11	0.55
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.77	-0.11	0.11	0.33	0.55	-0.11	0.33

# Convolution layer

One image becomes a stack of filtered images



# Convolution layer

One image becomes a stack of filtered images

-1 -1 -1 -1 -1 -1  
-1 1 -1 -1 -1 -1  
-1 -1 1 -1 -1 -1  
-1 -1 -1 1 -1 1  
-1 -1 -1 1 1 -1  
-1 -1 -1 -1 1 -1  
-1 -1 -1 1 -1 -1  
-1 -1 -1 1 1 -1  
-1 -1 -1 -1 -1 1  
-1 -1 -1 -1 -1 -1



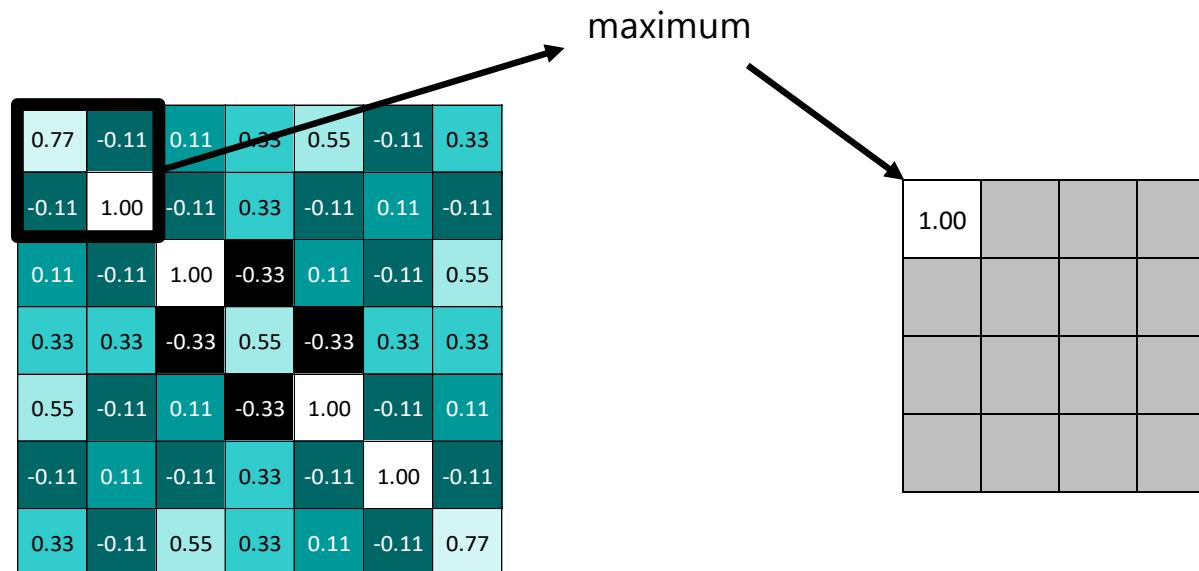
0.71	-0.11	0.11	0.09	0.55	-0.11	0.11
-0.11	1.00	-0.11	0.09	0.11	0.11	-0.11
0.11	-0.11	1.00	-0.09	-0.11	-0.11	0.11
0.09	0.11	-0.09	0.55	0.33	0.11	0.11
0.55	-0.11	0.11	-0.11	1.00	-0.11	0.11
-0.11	0.11	0.11	-0.11	1.00	0.11	0.11
0.11	-0.11	0.55	0.33	-0.11	0.11	0.11

# Pooling: Shrinking the image stack

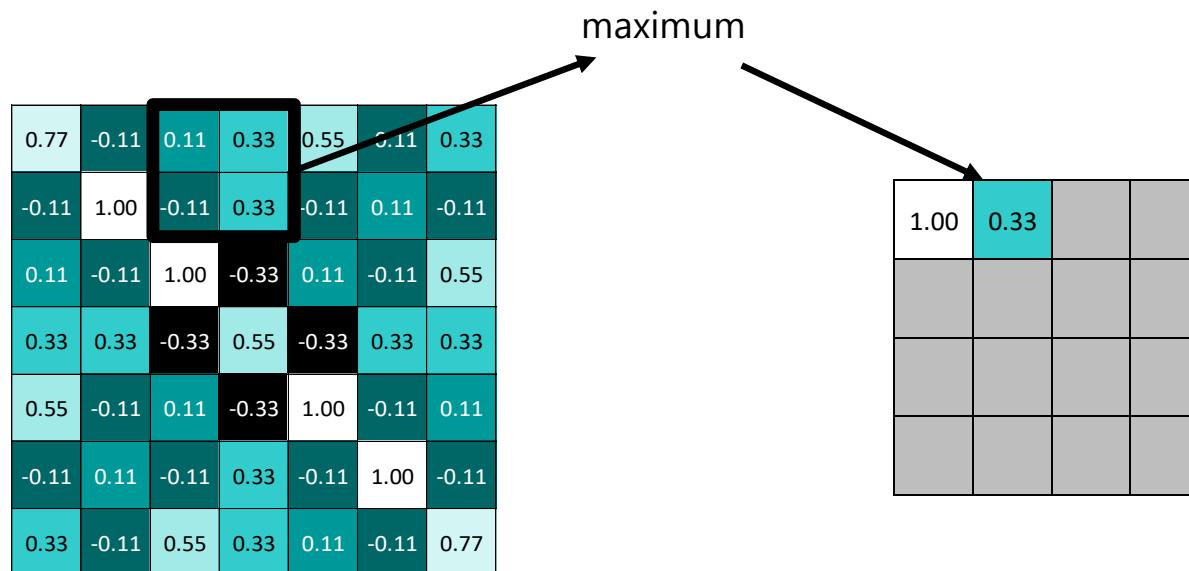
---

1. **Pick a window size (usually 2 or 3).**
2. **Pick a stride (usually 2).**
3. **Walk your window across your filtered images.**
4. **From each window, take the maximum value.**

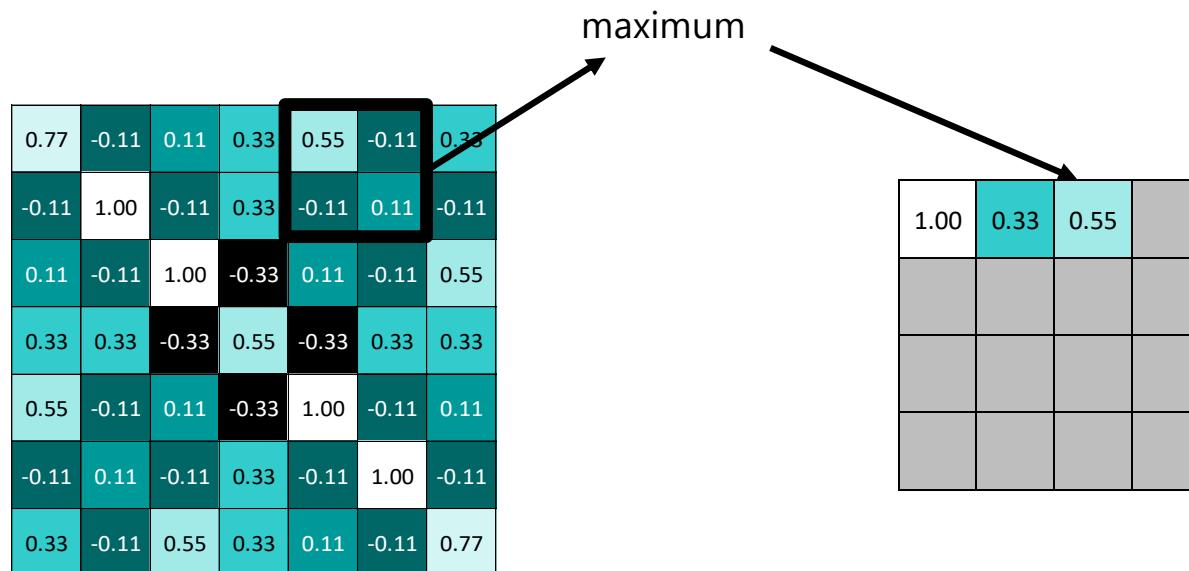
# Pooling



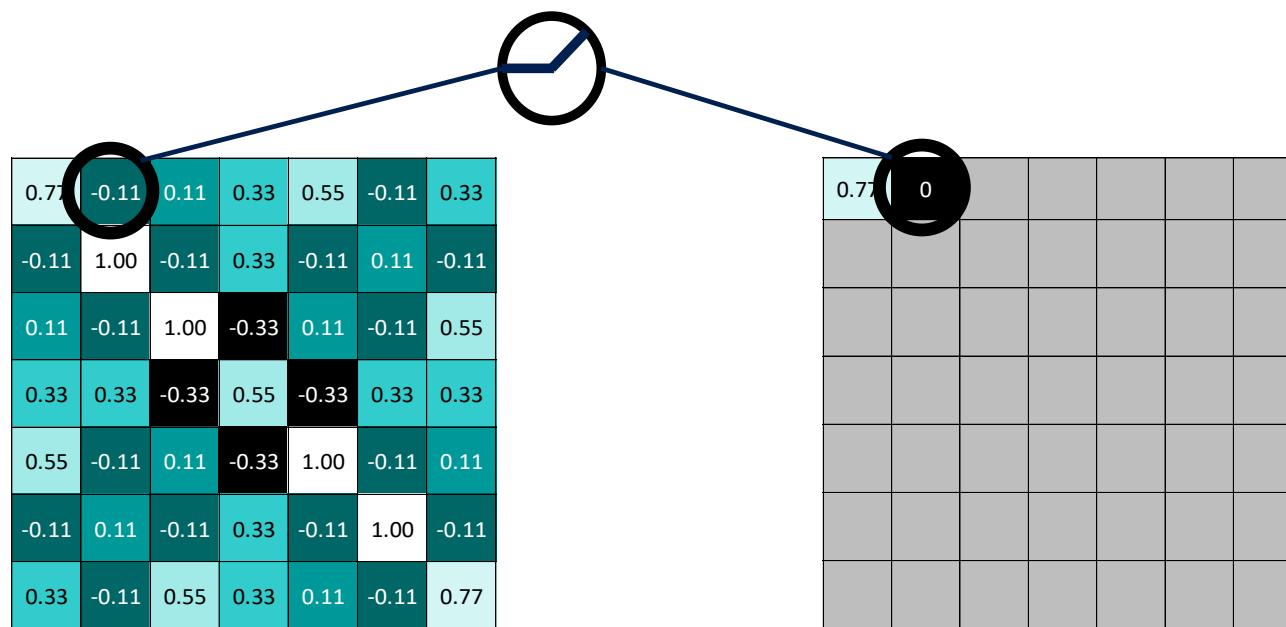
# Pooling



# Pooling

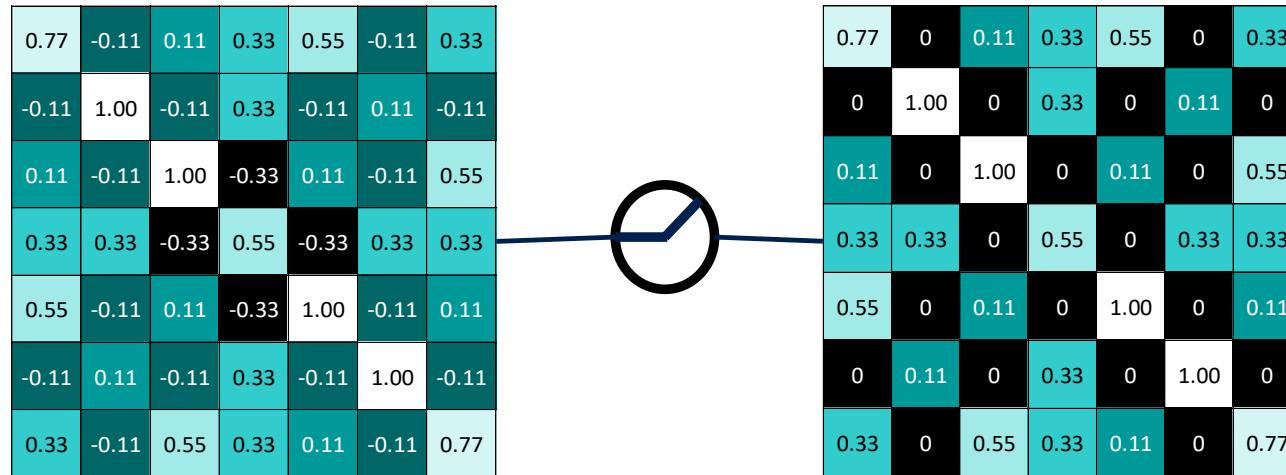


# Nonlinear Activation function: Rectified Linear Units (ReLUs)



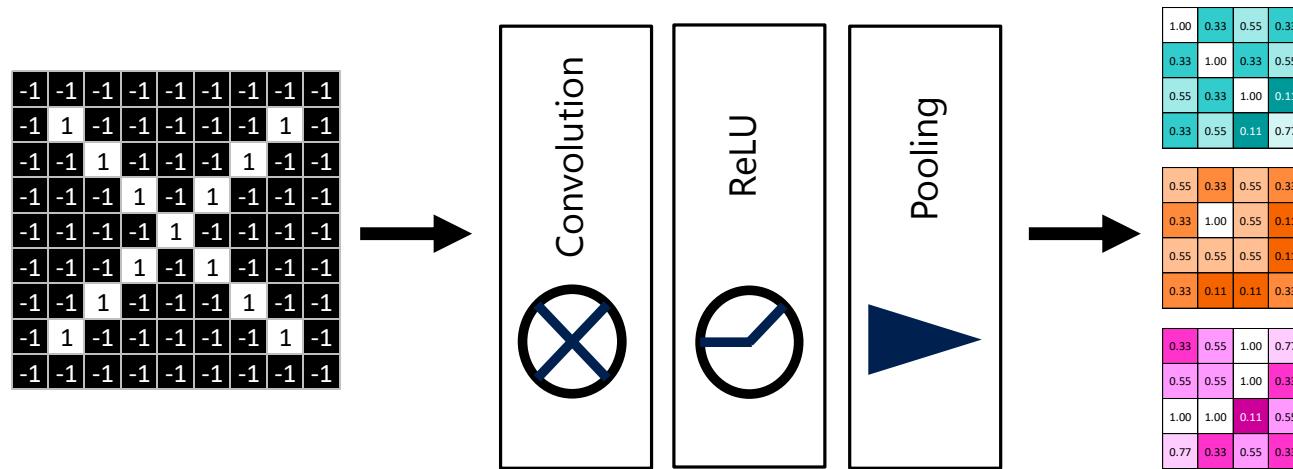
# Rectified Linear Units (ReLUs)

---



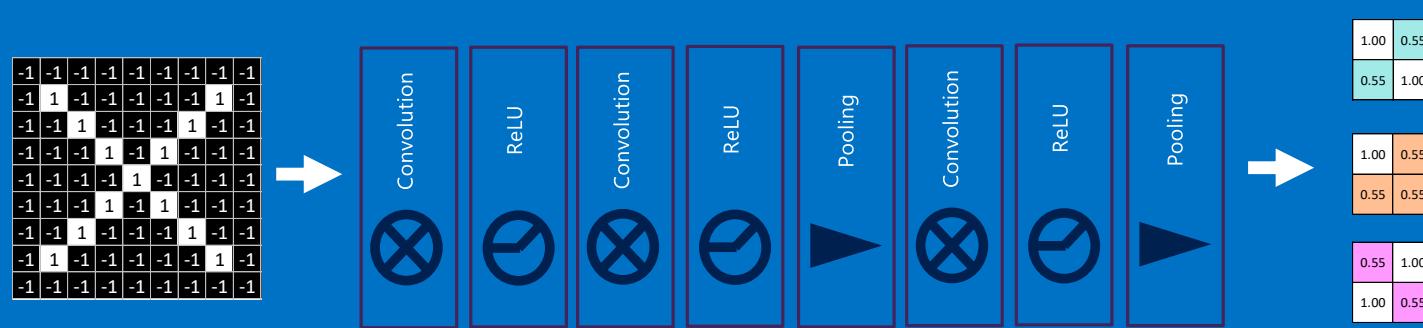
# Layers get stacked

The output of one becomes the input of the next.



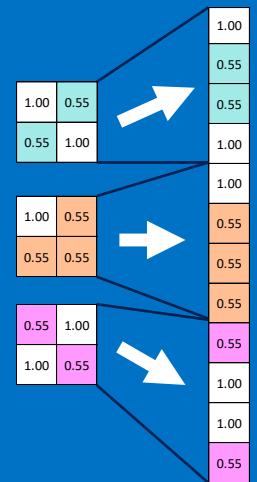
# Deep stacking

Layers can be repeated several (or many) times.



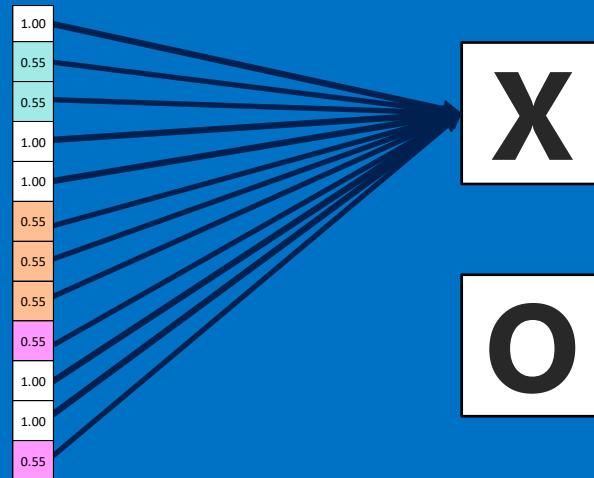
# Fully connected layer

Every value gets a vote



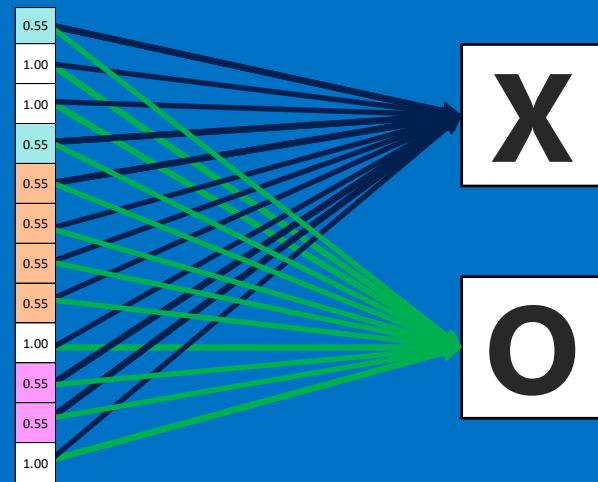
# Fully connected layer

Vote depends on how strongly a value predicts X or O



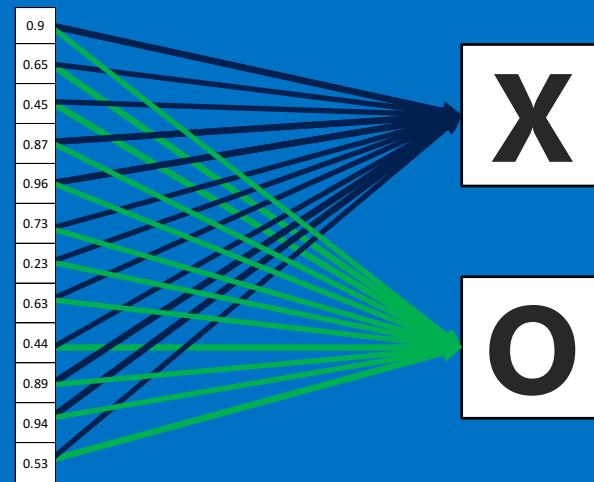
# Fully connected layer

Vote depends on how strongly a value predicts X or O



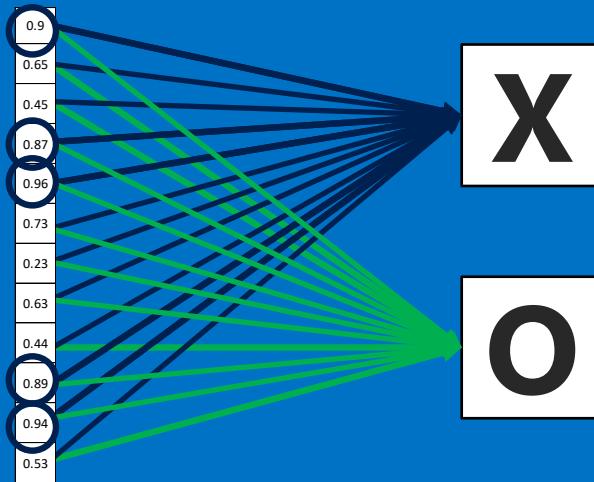
# Fully connected layer

Future values vote on X or O



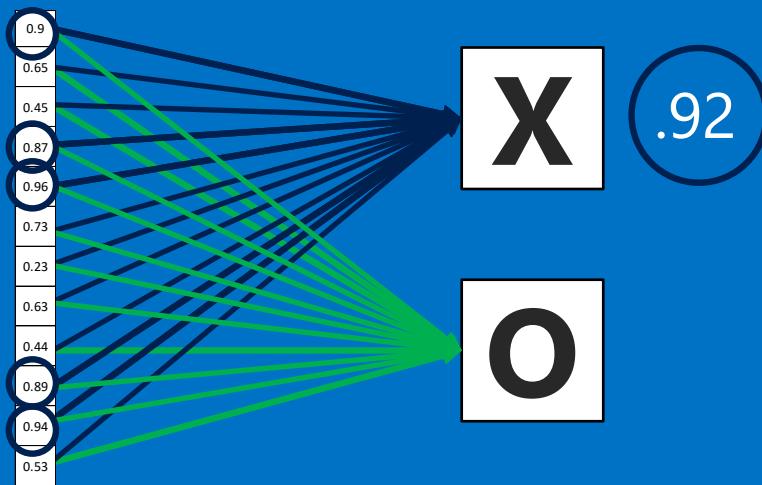
# Fully connected layer

Future values vote on X or O



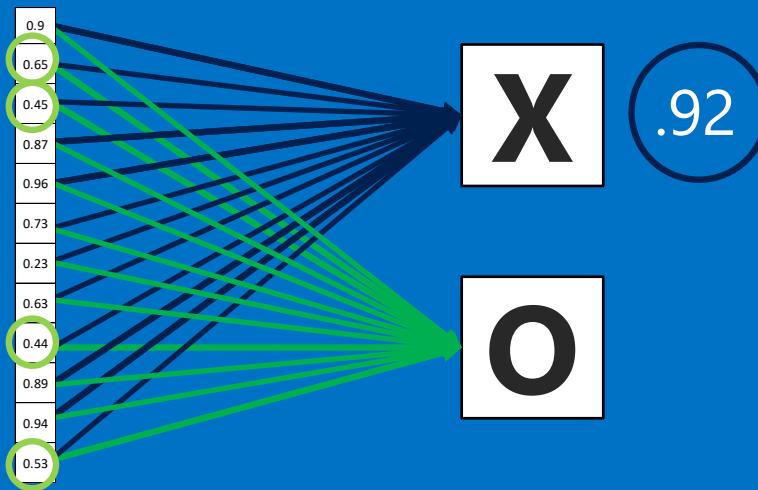
# Fully connected layer

Future values vote on X or O



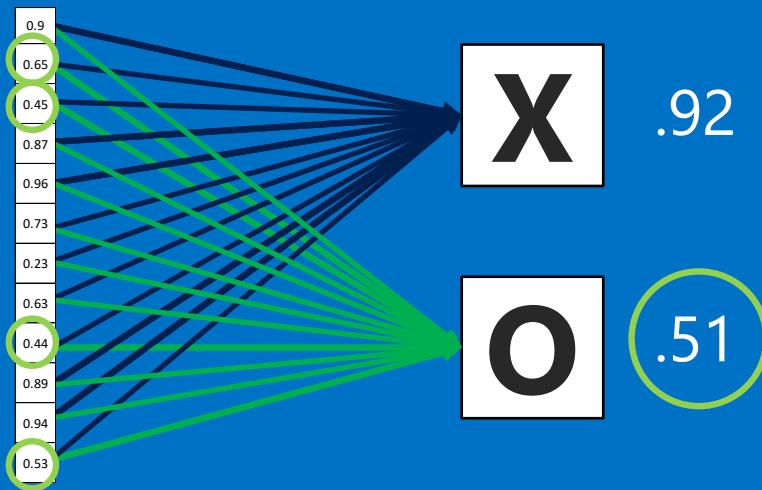
# Fully connected layer

Future values vote on X or O



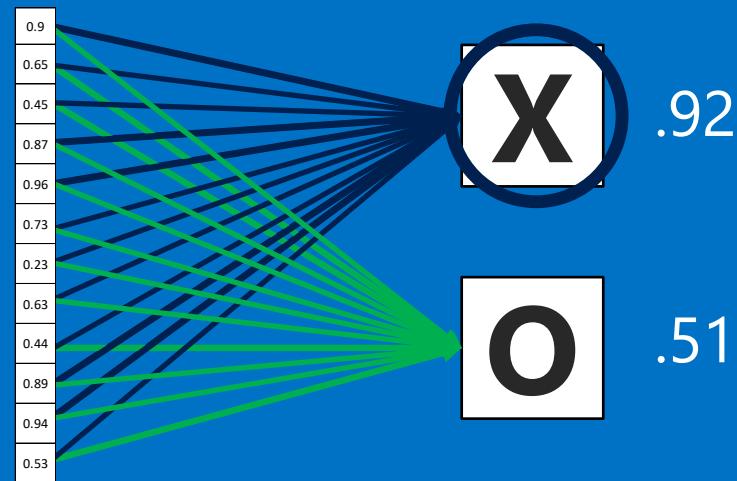
# Fully connected layer

Future values vote on X or O



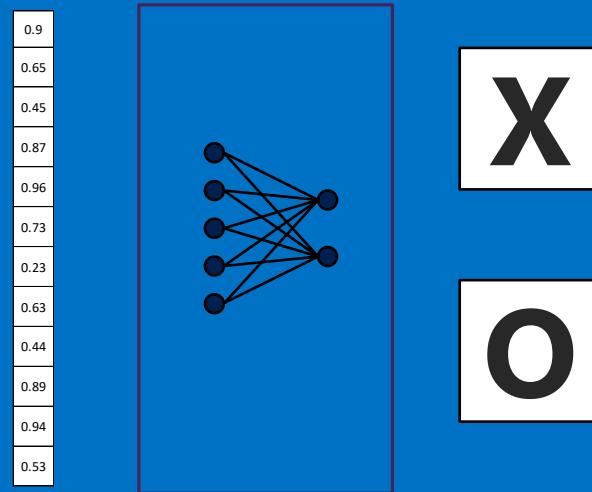
# Fully connected layer

Future values vote on X or O



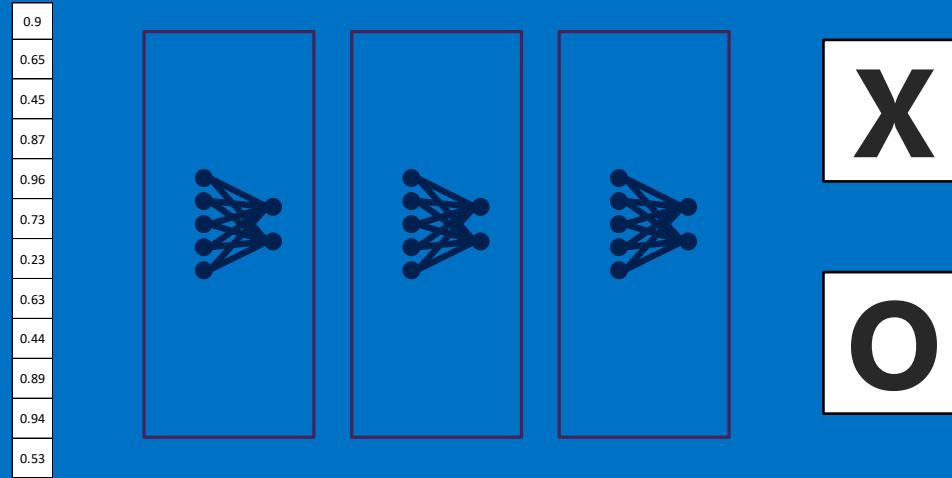
# Fully connected layer

A list of feature values becomes a list of votes.



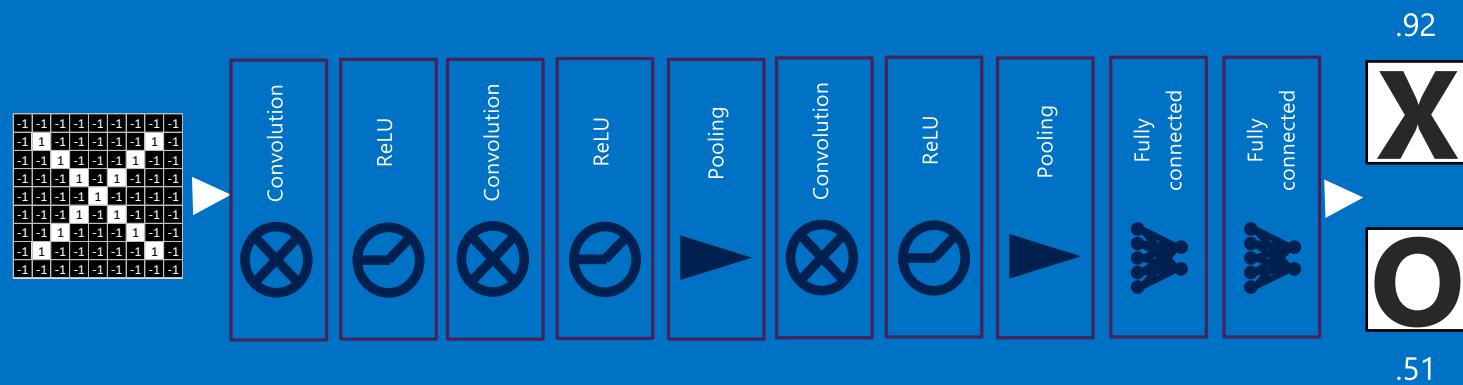
# Fully connected layer

These can also be stacked.



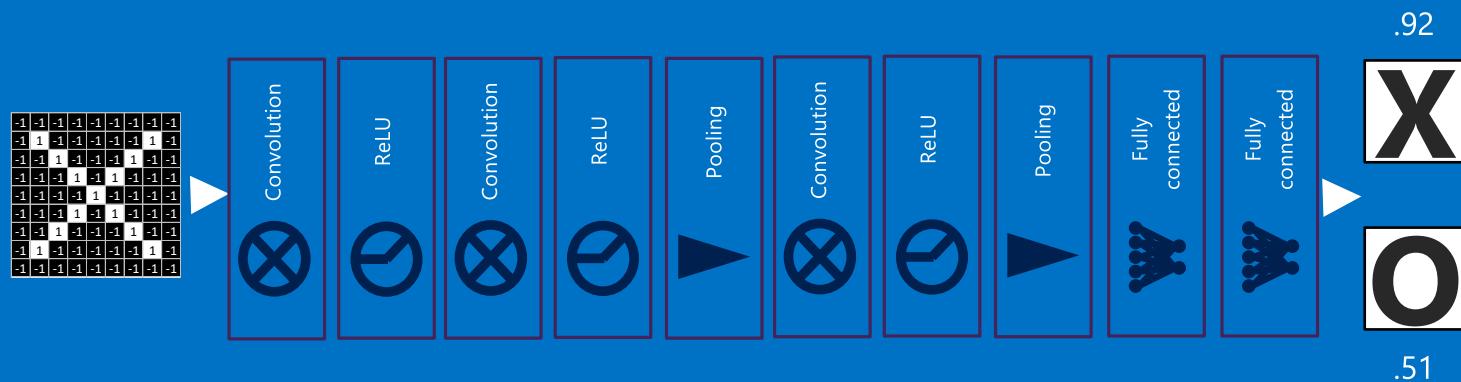
# Putting it all together

A set of pixels becomes a set of votes.



# Backprop (gradient descent) to learn weights/kernels

	Right answer	Actual answer	Error
X	1	0.92	0.08
O	0	0.51	0.49
Total		0.57	



Operation	Kernel	Image result
Identity	$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$	
Edge detection	$\begin{bmatrix} 1 & 0 & -1 \\ 0 & 0 & 0 \\ -1 & 0 & 1 \end{bmatrix}$	
	$\begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix}$	
	$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$	

---

Center element of the kernel is placed over the source pixel. The source pixel is then replaced with a weighted sum of itself and nearby pixels.

---

